How Long Do Markets Need to Fully React to Monetary Policy Announcements?

Paul L. Tran*

This version: 29 October, 2025

[Click here for the latest version]

Abstract

This paper shows that financial markets need more time to fully react to Federal Open Market Committee (FOMC) policy announcements than typically assumed. Using finance literature techniques and neural network methods for text analysis, I systematically estimate that on average, markets fully react to the information content of FOMC statements within an event window ending at least 30 minutes after release. This optimal window increases with the underlying maturity of an asset, reaching 50–60 minutes in length for maturities at least two quarters ahead. Additionally, statements with greater complexity, less similarity, and dissents are associated with longer event windows on average. I find that the correlation between monetary policy surprises measured within optimal versus 30-minute windows decreases with asset underlying maturity. These differences alter the forward guidance component of monetary policy shocks and magnify their estimated impact on interest rates, break-even inflation, and equity prices. Furthermore, constructing monetary policy shocks within optimal windows results in the responses of macroeconomic variables to become more precise.

Keywords: Event window studies, FOMC statements, monetary policy shocks, neural networks, natural language processing, text analysis

JEL Codes: C45, E52, E58, G14

*Department of Economics, The University of Texas at Austin. 2225 Speedway, BRB 2.128, C3100, Austin, TX 78712, USA. Email: pltran@utexas.edu. Website: https://paulletran.com/. I am indebted to my committee—Oli Coibion, Chris Boehm, Saroj Bhattarai, and Amy Handlan—for their invaluable guidance and support. For their helpful comments and suggestions, I also wish to thank Andi Mueller, Andres Drenik, Stefano Eusepi, Oliver Pfauti, Andrew Chang, and seminar participants at UT Austin. I am grateful to Chris Boehm and Niklas Kroner for sharing data with me. Finally, I thank Trieu Tran for her research assistance.

1 Introduction

> To establish the right price for a stock the market must have adequate information, but it by no means follows that if the market has this information it will thereupon establish the right price. The market's evaluation of the same data can vary over a wide range, dependent on bullish enthusiasm, concentrated speculative interest and similar influence, or bearish disillusionment. Knowledge is only one ingredient on arriving at a stock's proper price. The other ingredient, fully as important is sound judgement.
>
> Graham (1974)

Following the August 1999 Federal Open Market Committee (FOMC) announcement, the S&P 500 Index did not instantaneously settle at a new value. Instead, as shown in Figure 1, its price swung in both positive and negative directions for over an hour. This example highlights a key question: how much time do financial markets need to fully react to monetary policy announcements? Applied macroeconomists often measure these price changes within narrow windows around news events such as monetary policy announcements. Small windows avoid contamination from unrelated news, but strongly assume that the market's first response is its final response. On the other hand, wide windows allow more time for the market to react, but risk mixing the impact of monetary policy announcements with other information. Despite these trade-offs, the monetary policy literature has often defaulted to an "ad-hoc" 30-minute window. This "one-size-fits-all" approach is problematic because it assumes all markets take the same amount of time to react to monetary policy news, even though markets for longer-maturity assets might need more time to understand its more "abstract" information. If the chosen window is wrong, measurements of monetary policy surprises and shocks can become attenuated and less relevant.

I propose a systematic method that pins down these "optimal" event windows. First, I train a neural network for text analysis to approximate the underlying relationship between the text of monetary policy announcements and the market price reaction. This process creates a "text-based signal", which represents the network's predictions of the market price response based only on the FOMC statements, word-for-word. The key is that I have the network learn this relationship and generate this text-based signal for different event window lengths (e.g., a 10-minute window, 20-minute window, and so on). The "optimal" window is defined as the length where the neural network can make the most accurate predictions. This "Goldilocks" window is the time where the market has fully reacted to the monetary policy announcement, but not so much time that the price becomes overwhelmed by noise.[1]

My method yields four key findings. First, financial markets require more time to fully react to monetary policy announcements than the 30 minutes typically assumed. Specifically, I find that regardless of contract maturity, financial markets for futures and equities fully react within an event window starting 10 minutes before the announcement and *ending at least 30 minutes* after.

Second, the optimal window length is not "one-size-fits-all". It varies systematically with underlying asset maturity. Figure 2 summarises this result, showing the estimated optimal window length increases with the underlying maturity of the asset, rising from 40 minutes for short-horizon assets to 50–60 minutes for assets with underlying maturities at least two quarters ahead. Relatedly, I find that statements with greater complexity, less similarity to previous texts, and the presence of dissents also require longer event windows on average.

Third, the correlation between monetary policy surprises measured with my optimal windows versus the standard 30-minute window decreases as asset maturity increases. While interest-rate surprises for current and next FOMC meetings are similar, the correlation for long-maturity assets (e.g., fifteen years out) can decline by as much as 10%. These differences primarily alter the forward guidance component of monetary policy shocks, leading to larger estimated peaks and troughs or "shifts in importance" in the shock's composition.

Fourth, using the optimal window increases the estimated impact of monetary policy shocks about forward guidance on interest rates, break-even inflation, and equity prices. Furthermore, the estimated responses of macroeconomic variables also become more precise. This suggests that previously documented effects of monetary policy may be attenuated and imprecise due to the use of suboptimal window lengths.

[1] Additionally, this interpretation about the optimal event window is related to the findings of Casini and McCloskey (2025) about the relative dominance of the monetary policy surprise.

This paper is not the first to investigate the appropriate event window length. For example, Hillmer and Yu (1979) find that windows should last several hours, while Chang and Chen (1989) find they should span several days. Das and King (2021) find that asymmetric, narrower one-day event windows exhibit greater information content for earnings announcements. More recently, Casini and McCloskey (2025) argues that the window should be selected based on when the monetary policy surprise "dominates" other noise.[2] Using a heteroscedasticity-based statistical method, Boehm and Kroner (2025) select a fourteen-hour window. My paper contributes to this literature by introducing a new source of information. While the aforementioned studies infer the optimal window using only observed price movements, my method combines these dynamics with a "text-based signal" from a neural network that reads the FOMC statement itself. This allows my method to systematically estimate the window length where the "text-driven" price reaction is fully captured and noise is minimised on average. I also contribute to the high-frequency identification literature by detailing the effects of window choice on measured monetary policy shocks and surprises, particularly regarding forward guidance.

This paper joins a growing literature that uses text analysis to study the effects of central bank communication on expectations and the macroeconomy. Traditional methods in economics often relied on word counts or sentiment dictionaries to classify text (e.g., Gentzkow, B Kelly, and Taddy, 2019; Husted, Rogers, and Sun, 2020; Acosta, 2023; and others). A key limitation of these methods is that they can miss the nuances of context. In contrast, modern neural networks are designed to understand the full context and interdependencies of words. Recent papers have used these advanced tools to read FOMC statements and speeches, creating new and improved measures of monetary policy shocks (e.g., Handlan, 2022b; Doh, Song, and SK Yang, 2023; Piller, Schranz, and Schwaller, 2025; and others).

I contribute to this new literature from a different angle. While the studies mentioned above use text-analysis neural networks to directly create policy shocks, I use these methods

[2] The authors call this condition relative exogeneity. This is a technical condition which means the policy news is so large that it effectively drowns out all other market noise within that window.

to first address the event window timing problem. I apply the neural network to determine the optimal window length where the network can best understand the relationship between monetary policy communication and market price reactions. By first identifying the proper window, my method allows for the construction of more precise and relevant monetary policy surprises and shocks.

The remainder of this paper is structured as follows. Section 2 provides a framework of the dynamics of asset market prices in response to news. Despite strong and simplifying assumptions imposed, the framework motivates the necessity of my text-analysis neural network approach. Section 3 describes the input and output data of the text analysis method as well as the data used to investigate the effects of event window choice on monetary policy shock impacts. Section 4 details the process of systematically estimating the optimal event window lengths and the neural network approach behind it all. Section 5 presents the estimated length of time that financial markets need to fully react to monetary policy announcements. Section 6 details the effects of event window choice on monetary policy surprises and shocks. Section 7 details how the complexity, similar, and presence of dissents within the FOMC statements can affect the optimal window lengths. Finally, Section 8 concludes.

2 Motivating Framework

This section presents a motivating framework to show why the shortest or longest event windows are not always optimal. The framework illustrates how conflicting noise components can distort asset price dynamics following a news release. For simplicity, this framework imposes strong linearity and orthogonality assumptions on the noise components. These assumptions are purely illustrative, serving only to show how noise affects asset prices and, in turn, the optimal window length. My systematic estimation method does not impose these assumptions because the true composition of market price responses is unknown. Therefore, this framework and its related simulations serve to motivate the natural language processing approach used to extract a signal from the text of the news releases.[3]

Consider the price of one asset market at time t, P_t. Assume that the news is released at the beginning of the time period, $t = 0$. As shown in Equation 1, the asset price is the sum of three components:

$$\{P_t = P_t^f + \varepsilon_t^c + \varepsilon_t^n | t \geq 0\}, \tag{1}$$

where P_t^f is the fundamental price component, ε_t^c is cognitive noise, and ε_t^n is unrelated news.

P_t^f represents what asset price should be set by the market due to *only* the news release, meaning it is assumed that $P_t^f = P^f \in \mathbb{R}$.

Cognitive noise represents the collective "market noise" from participants' limited capacity to process information. Immediately after a news release, factors like algorithmic trading, over-reactions, under-reactions, or liquidity constraints mean market participants may have incomplete information or differing beliefs about the news. As time progresses, processing costs fall and investors learn from each other's price responses, leading to a convergence towards the fundamental price component. In other words, cognitive noise and its error "decay" over time, eventually converging to zero. The framework reflects this behaviour by representing cognitive noise as the following modified AR(1) process:

$$\varepsilon_t^c = \rho_c \varepsilon_{t-1}^c + e^{-\mathcal{D}t} \nu_t^c,$$

where coefficient $|\rho_c| < 1$, decay term $\mathcal{D} \in \mathbb{R}^+$, and random noise ε_t^c is normally distributed with mean zero and variance σ_c^2. I ensure the cognitive noise process exhibits the decaying behaviour over time by assuming that $|\frac{\rho_c}{\mathcal{D}}| < 1$ and that the variance of cognitive noise at time $t = 0$ is equal to σ_c^2.

The effect of news unrelated to the event of interest on the asset price is modelled as a

[3] Details behind the simulation process can be found in Appendix A.

random walk:

$$\varepsilon_t^n = \varepsilon_{t-1}^n + \nu_t^n,$$

where random noise ν_t^n is normally distributed with mean zero and variance σ_n^2. Similar to cognitive noise, the variance of the unrelated news process at time $t = 0$ is equal to zero.

2.1 The Effects of Noise on the Variance of Asset Prices

The expression of P_t as Equation 1 allows for the derivation of the variance of the asset price for all time $t \geq 0$ through iterative substitution.

At time $t = 0$ when news is released, one can express the variance of the asset price as $\text{Var}(P_0) = \text{Var}(\varepsilon_0^c) + \text{Var}(\varepsilon_0^n) = \sigma_c^2$. Similarly, the expression at time $t = 1$ is $\text{Var}(P_1) = \text{Var}(\varepsilon_1^c) + \text{Var}(\varepsilon_1^n) = \sigma_c^2(\rho_c^2 + e^{-2\mathcal{D}}) + \sigma_n^2$. Continuing this iterative process yields the following expression of

$$\begin{aligned}\text{Var}(P_t | t \geq 0) &= \left[\sum_{i=0}^{t} \rho_c^{2(t-i)} e^{-2\mathcal{D}i}\right] \sigma_c^2 + t\sigma_n^2 \\ &= \left[\frac{\rho_c^{2(t+1)} - e^{-2(t+1)\mathcal{D}}}{\rho_c^2 - e^{-2\mathcal{D}}}\right] \sigma_c^2 + t\sigma_n^2, \end{aligned} \tag{2}$$

where $\lim_{t\to\infty} \left[\frac{\rho_c^{2(t+1)} - e^{-2(t+1)\mathcal{D}}}{\rho_c^2 - e^{-2\mathcal{D}}}\right] = 0$ by the aforementioned framework assumptions.

I define t^{one} as the time where the variance of P_t is minimised. Solving for this time horizon yields the following indirect expression that provides important insight into the factors influencing the appropriate event window length:

$$t^{one} : \mathcal{D}\left[e^{-2(t+1)\mathcal{D}}\right] + ln(\rho_c)\rho_c^{2(t+1)} = \left[\frac{(e^{-2\mathcal{D}} - \rho_c^2)}{2}\right] \frac{\sigma_n^2}{\sigma_c^2}. \tag{3}$$

Two important findings from Equation 3 are that $\frac{\partial t^{one}}{\partial \sigma_n^2} < 0$ and $\frac{\partial t^{one}}{\partial \sigma_c^2} > 0$. In other words, an increased presence of unrelated news (cognitive noise) results in the time where the variance of the asset price is minimised to decrease (increase).[4] The dynamics caused by these two components give formal insight into why strictly using narrow event windows isn't appropriate.

2.2 Estimator Form

The dynamics described by Equation 3 provide insight into the trade-offs for a *single* news announcement. To apply this framework empirically such that the methodology of this paper is motivated, this concept must be extended to the *average* price reaction across many news releases. Therefore, the goal of this paper is not to find t^{one} for any single event, but to systematically estimate t^*, the time at which the market fully reacts to news announcements *on average.* Formally, consider N news announcements and the price of one asset market. For each news i that is released at time t, the asset price $P_{i,t}$ and its components will respond to the announcement. Therefore, I define t^* as the optimal time horizon that minimises the mean squared error (MSE) between $P_{i,t}$ and $P_{i,t}^f$ for all N news:

$$t^* : \min_t \frac{1}{N} \sum_{i=1}^{N} \left(P_{i,t} - P_i^f\right)^2 = \min_t \frac{1}{N} \sum_{i=1}^{N} \left(\varepsilon_{i,t}^c + \varepsilon_{i,t}^n\right)^2 \tag{4}$$

However, assume that the fundamental component is unobservable to the econometrician. Instead, suppose that the econometrician observes a noisy signal of the fundamental component, $s_i = P_i^f + \xi_i$, where $\xi_i \sim \mathcal{N}(0, \sigma_s^2)$. By making the core assumption that the sum of the asset price noise components, $(\varepsilon_t^c + \varepsilon_t^n)$, and the signal noise are *independent*, I

[4] Because Equation 3 is unable to have t isolated on one side, I numerically verify the dynamics of the t^{one} for various values of σ_c^2 and σ_n^2 in the indirect expression while holding the other parameters constant.

am able to derive the following expression for the minimisation problem of the MSE:

$$
\begin{aligned}
t^* : \min_t \frac{1}{N} \sum_{i=1}^{N} (P_{i,t} - s_i)^2 &= \min_t \frac{1}{N} \sum_{i=1}^{N} \left(P_i^f + \varepsilon_{i,t}^c + \varepsilon_{i,t}^n - P_i^f - \xi_i \right)^2 \\
&= \min_t \frac{1}{N} \sum_{i=1}^{N} \left(\varepsilon_{i,t}^c + \varepsilon_{i,t}^n - \xi_i \right)^2 \\
&= \min_t \frac{1}{N} \sum_{i=1}^{N} \left[\left(\varepsilon_{i,t}^c + \varepsilon_{i,t}^n \right)^2 + \xi_i^2 - 2\xi_i \left(\varepsilon_{i,t}^c + \varepsilon_{i,t}^n \right) \right] \\
&= \min_t \left\{ \mathbb{E}\left[\left(\varepsilon_{i,t}^c + \varepsilon_{i,t}^n \right)^2 \right] + \mathbb{E}\left[\xi_i^2 \right] - 2\,\mathbb{E}\left[\xi_i \right] \mathbb{E}\left[\left(\varepsilon_{i,t}^c + \varepsilon_{i,t}^n \right) \right] \right\} \\
\implies t^* : \min_t \frac{1}{N} \sum_{i=1}^{N} (P_{i,t} - s_i)^2 &= \min_t \left[\frac{1}{N} \sum_{i=1}^{N} \left(\varepsilon_{i,t}^c + \varepsilon_{i,t}^n \right)^2 + \sigma_s^2 \right] \qquad (5)
\end{aligned}
$$

This derivation implies that the econometrician can still solve for the optimal time horizon, t^*, using only the noisy signal s_i. This result is asymptotic: given infinite news events, the signal's precision has no effect on finding t^*. However, in finite samples like the monetary policy literature,[5] the signal's precision *does* determine the feasibility of estimating t^*. An equivalent and important interpretation is that t^* is the time where the average impact of noise is minimised. At this point, the fundamental component P_i^f (and its signal s_i) has the largest "share" in the observed asset price.

2.3 Framework Takeaways

I simulate the asset price process over time for multiple news announcements to illustrate the effects of cognitive and unrelated noise on the optimal time horizon, and to demonstrate why a good signal is needed for this MSE minimisation. The main implications are discussed in the main text, while Appendix A provides full details on the process.

I consider three scenarios for the asset price response: one with cognitive noise, one with unrelated news, and one with both. The simulations show that in all scenarios, the optimal time horizon (t^*) estimated from the noisy signal s_i is essentially equal to the true horizon

[5] For example, the FOMC only release eight scheduled monetary policy announcements in a year.

estimated from the fundamental price P_i^f. A "good" signal, therefore, makes it possible to estimate the time horizon that reflects a full market reaction. However, recall that this framework and its simulations impose strong linearity and orthogonality assumptions purely to demonstrate the effects of noise. In reality, the true compositions and processes of the fundamental and noise components are unknown, yet still influence the overall asset price.

Abstracting from this simple framework, I argue that the text-analysis neural network approach used in this paper does not need to know the true noise processes to approximate the underlying mapping between monetary policy communications and asset price changes. Therefore, the method is still able to extract a "good" signal.

3 Data

The sample period for my analysis runs from May 1999 to October 2019. As the principal entity determining U.S. monetary policy, the Federal Open Market Committee (FOMC) is closely watched by markets. Of particular importance are the released FOMC statements, which are the initial and primary source of monetary policy news for market participants.[6] Because markets are known to dissect and react to these statements word-by-word,[7] understanding the time it takes for them to fully react is critical for monetary policy event studies. However, the mapping from FOMC statements to asset price reactions is complex, nonparametric, and difficult to capture with traditional text analysis methods. Approximating this underlying relationship motivates the text-analysis neural network approach in this paper. Before detailing the systematic estimation, I provide descriptive information on the FOMC statements (i.e., the inputs) and the financial market asset prices (i.e., the outputs) used in this paper, as summarised in Appendix Table F2.

[6] Appendix Figure E1 shows U.S. interest spikes on Google Trends for phrases such as "FOMC meeting" and "FOMC statement" during scheduled meeting dates. Further support for financial market reactions coming primarily from the FOMC statements is the fact that the 1^{st}–3^{rd} query results on Google Search direct to the Board of Governors of the Federal Reserve System website.

[7] E.g., CNBC coverage of the January 2025 FOMC statement text.

3.1 Inputs: FOMC Statements

FOMC decisions are announced in a press release after its eight scheduled annual meetings. The FOMC occasionally holds unscheduled meetings, but I drop these statements from my sample. This exclusion is to ensure that measured asset price changes are driven by the statement's content, not by the surprise of an unscheduled meeting. I source all scheduled statements from the Board of Governors of the Federal Reserve System website, resulting in a sample size of 165 statements.[8]

Text Pre-processing To prepare the FOMC statements for the neural network, I pre-process all 165 statements by standardising them to plain UTF-8 text with uniform spacing. I also remove URLs, release timestamps, and the list of regional bank request approvals. Most importantly, I remove the FOMC member voting record from the end of each statement. This removal is prioritised for two main reasons. First, it ensures the neural network, which has a hard constraint on input length, focuses entirely on the main economic discussion (see Subsection 4.3). Second, while voting records are known to affect stock markets (C Madeira and J Madeira, 2019), they have less impact on interest-rate assets. Given my paper's focus on interest-rate futures, this exclusion should not significantly affect the estimation of the optimal event windows.[9]

Content and Evolution of Statements A typical FOMC statement discusses current macroeconomic conditions, communicates the Committee's economic expectations, and concludes with the new Federal Funds and discount rates. Following the Great Recession of 2008–2009, these statements began to include discussions of unconventional monetary policy, such as quantitative easing. This post-Crisis period saw an increase in statement complexity and length, particularly as the Committee relied more on communication to influence market expectations when it could no longer use rate changes. As seen in Figure 3, statement length grew rapidly until 2014. This trend partially reverses after 2014, but statements remain

[8] https://www.federalreserve.gov/monetarypolicy/fomc.htm

[9] Nonetheless, future versions of this paper will explore including the voting records as inputs.

longer on average than before the Great Financial Crisis.

The following example is an excerpt of the FOMC statement released on 01 August, 2018:

> Information received since the Federal Open Market Committee met in June indicates that the labor market has continued to strengthen and that economic activity has been rising at a strong rate. [...]
>
> The Committee expects that further gradual increases in the target range for the federal funds rate will be consistent with sustained expansion of economic activity, strong labor market conditions, and inflation near the Committee's symmetric 2 percent objective over the medium term. Risks to the economic outlook appear roughly balanced.
>
> In view of realized and expected labor market conditions and inflation, the Committee decided to maintain the target range for the federal funds rate at 1-3/4 to 2 percent. The stance of monetary policy remains accommodative, thereby supporting strong labor market conditions and a sustained return to 2 percent inflation.

3.2 Outputs: Interest-rate and Equity Futures

The outputs for the neural network are interest-rate and equity futures prices. I use intraday (tick) data for all assets, sourced from the Thomson Reuters Tick History database (LSEG). This high-frequency data is necessary to analyse the short event windows (e.g., 30 minutes) common in the high-frequency identification literature. Within these windows, the variation in interest-rate futures prices measures the change in expected interest rate paths in response to FOMC announcements.

I select a range of futures contracts commonly used to construct high-frequency monetary policy surprises (Kuttner, 2001; Gürkaynak, Sack, and Swanson, 2005; Gertler and Karadi, 2015; Nakamura and Steinsson, 2018; Jarociński and Karadi, 2020; Handlan, 2022b; and others). For short-term expectations, I use Federal Funds futures, which are liquid for measuring expectations up to three months out ($FF1$–$FF4$). To capture expectations from four months to one year ahead, I use Eurodollar futures ($EDcm2, EDcm3, EDcm4$), following the methodology of Acosta, Brennan, and Jacobson (2024). For longer horizons, I use Treasury futures with two ($TUc1, TUc2$), five ($FVc1, FVc2$), ten ($TYc1, TYc2$), and thirty-year ($USc1, USc2$) contracts, which correspond to underlying maturities of approximately two, four, seven, and fifteen years, respectively (Gürkaynak, Kisacikoğlu, and Wright, 2020). Finally, I include the S&P 500 Index and its E-mini futures ($SPX, ESc1, ESc2$). I use equities for two reasons: they are a central asset in the monetary policy literature, and their futures trade outside regular hours, providing the high-quality data necessary to test a wide variety of event window lengths.

Dependent Variable Construction I collect futures contract price levels at 10-minute intervals, from 10 minutes before to 18 hours after an FOMC statement release. The output of interest for the neural network is the log-price difference for each contract, constructed for event windows starting 10 minutes before and ending n minutes after the announcement:

$$DP_{t+n} = ln\left(\frac{P_{t+n}}{P_{t-10}}\right), \tag{6}$$

The neural network approach considers price log-differences for event windows up to 70 minutes in length.[10]

3.3 Data to Investigate Impacts of Monetary Policy Shocks

To investigate the effects of event window choice on monetary policy shocks, I use several additional datasets. Daily data for nominal Treasury yields are from Gürkaynak, Sack, and Swanson (2005).[11] Daily data for Treasury-inflation-protected-security (TIPS) yields and break-even inflation come from Gürkaynak, Sack, and Wright (2010).[12] Monthly data for industrial production (IP) and the consumer price index (CPI) are from Federal Reserve Economic Data. The monthly excess bond premium (EBP) measure is from Gilchrist and Zakrajšek (2012) and sourced from the Federal Reserve Board.[13]

[10]Longer event windows are not considered due to computational and financial constraints, but remain an interest for future versions of this paper.

[11]Treasuries: https://www.federalreserve.gov/data/nominal-yield-curve.htm

[12]TIPS and break-even inflation: https://www.federalreserve.gov/data/tips-yield-curve-and-inflation-compensation.htm

[13]https://www.federalreserve.gov/econresdata/notes/feds-notes/2016/files/ebp_csv.csv

4 Systematically Estimating Optimal Event Windows

Systematically estimating the optimal event window length hinges on approximating the underlying, nonparametric relationship between monetary policy communications and financial market price reactions. Obtaining the text-based signal needed for this estimation requires a precise approximation of this mapping. This precision, in turn, demands a method that can quantify the *full* information content of the FOMC statement, capturing all dimensions of communication that drive market reactions.

Traditional text analysis methods, such as "fitting predictive models on simple counts of text features" (Gentzkow, B Kelly, and Taddy, 2019), typically fail to capture the full information content of text. These methods often miss context and word interdependencies. For example, a "bag-of-words" approach would treat the phrases "*employment* went up, but *inflation* did not" and "*inflation* went up, but *employment* did not" as identical. Similarly, "n-gram" methods, which measure neighbouring words, miss information spread across a sentence. A trigram might capture "economic growth slowed" but miss the crucial reversal in the phrase "economic growth slowed, but is expected to pick up pace later this year". Therefore, these popular methods cannot realistically capture the "full" information content of FOMC statements.[14]

In contrast, text-analysis neural networks *can* capture complex relationships like context and interdependencies. For example, the words "slowed" and "pick up pace" could both be associated with "economic growth" for prediction, even if they are not adjacent. The power of these networks is mathematically justified by the Universal Approximation Theorem (Hornik, Stinchcombe, and White, 1989; and others), which states that any continuous function (such as the underlying mapping from FOMC statements to price reactions) can be approximated by a network with at least one hidden layer to any desired accuracy.[15]

[14] Papers like Babur and Cleophas (2017) have shown that model performance does not increase monotonically with the size of "n-gram" methods.

[15] The Universal Approximation Theorem is an existence theorem; it guarantees *existence* of a network with at least one hidden layer, but does not specify its structure. In practice, deep networks (multiple layers) are often used to reduce parameters and computational requirements for complex approximations.

Given their success in other social sciences, these methods offer similar benefits for studying monetary policy communication in economics (Gentzkow, B Kelly, and Taddy, 2019).

4.1 Approach: XLNet-Base, a Neural Network for Text Analysis

I employ the XLNet-Base language model from Z Yang et al. (2019), which builds upon the transformer architecture of Vaswani et al. (2017).[16] This model abandons the common "masking" technique, which can create training discrepancies and incorrectly assumes word independence.[17] Instead, XLNet-Base uses a permutation-based pre-training phase, allowing it to learn context bi-directionally without the weaknesses of masking. This makes XLNet-Base better suited for semantic tasks than unidirectional generative models (e.g., ChatGPT).[18] The model's versatility makes it ideal for fine-tuning on new regression tasks, such as approximating the underlying relationship between FOMC statements and price reactions. I start with the pre-trained XLNet-Base model and fine-tune it to "transfer" its general language knowledge to this specific task, allowing me to obtain the text-based signal.[19]

4.2 The Inputs, Outputs, Approach, and Optimal Window Length

Recall from the motivating framework that the optimal event window is where noise is minimised on average, allowing the fundamental price component to have the largest "share" of the price change. The text-based signal is the neural network's *predicted* price change for a given window, using only the FOMC statement text. It follows that the optimal window length is where the network's approximation simultaneously has the highest predictive performance, because only within this window is noise minimised on average, and produces the

[16]The open-source, pre-trained neural network has 12 layers, 12 attention heads, 117 million parameters, and a 32,000 token vocabulary. See Hugging Face for details.

[17]E.g., if "target" and "rate" are masked in "...increase its target rate", their implicit relation is ignored.

[18]Generative models are "left-to-right" and predict a word using only preceding words.

[19]Transfer learning, as shown by Z Yang et al. (2019) and others, reduces data requirements for new tasks while maintaining high accuracy.

most precise signal. This simultaneity drives my systematic estimation procedure:

1. For a given event window length, XLNet-Base regresses futures price changes on the FOMC statement text to assess its predictive performance and signal precision.

2. I repeat this regression for all event windows up to 70 minutes in length.

3. The optimal event window for that asset is the length that yields the network's best predictive performance.

4.3 Stratified-sampling Cross Validation

To prepare the neural network for systematic estimation, I split the sample five times, with 20 per cent of FOMC statements in each testing subsample. This splitting is stratified by the FOMC's rate decision, the FOMC Chair, pre-/post-2007, and statement word count to ensure equal distributions. This process yields five distinct training-testing folds such that no statement is shared across testing subsamples.[20]

These splits consist of input-output pairs: the FOMC statements (inputs) and the futures price log-differences (outputs). For each event window length and futures contract, I train XLNet-Base on a training subsample to approximate the underlying relationship. I then test its performance on the held-out subsample. This process is repeated five times, once for each split. The final accuracy for each event window length is the average performance across all five test splits. This method ensures a robust measure of out-of-sample performance and accounts for prediction variance across the splits.[21]

I restrict the context window of XLNet-Base (i.e., the number of word tokens it considers per input) to 512 tokens for each FOMC statement.[22] Given the average statement length is roughly 327 words (peaking at 800), the 512-token window provides adequate headroom to

[20] This process is known as stratified 5-fold cross-validation in the machine learning literature.

[21] Stratified 5-fold cross-validation minimises differences between the full-sample and subsample distributions, which is crucial for transfer learning on finite samples.

[22] Tokens are converted text inputs used to reduce the network's vocabulary size (e.g., "decreasing" and "increasing" might become {"de", "in", "creas", "ing"}).

understand most statements. A list of the main hyperparameters for fine-tuning is available in Appendix Table F3.

4.4 Accuracy Metrics

The primary accuracy metric for neural network fine-tuning is a *generalised* R^2 statistic from Hawinkel, Waegeman, and Maere (2024), denoted R^2_{OOS}. This metric is chosen for two main reasons. First, the conventional R^2 (i.e., the proportion of variance explained by a model) breaks down for non-linear methods like neural networks. For such models, the total variance no longer neatly decomposes into model and residual components, meaning the squared Pearson correlation coefficient does not equal R^2. Second, the R^2_{OOS} formula is specifically designed to assess out-of-sample performance, not in-sample fit.[23] Mathematical details of this statistic are in Appendix B.

As is standard in machine learning, model performance is compared to a baseline. This baseline assumes *no* relationship between the FOMC statement text and futures price changes, naively predicting with the in-sample average. Therefore, the R^2_{OOS} measures the reduction in predictive error achieved by the neural network compared to this naive baseline.

During fine-tuning, XLNet-Base also tracks other metrics: the out-of-sample Pearson correlation between predicted and actual values, the out-of-sample mean absolute error, and the in-sample mean squared error. This last metric is important for verifying that the network is genuinely learning the underlying relationship.

The decisive criterion for systematic estimation is the out-of-sample R^2 *averaged across sample splits* for each event window length and futures contract:

$$\overline{R^2_{OOS}} = \frac{\sum_{i=1}^{K} R^2_{OOS}}{K}, \tag{7}$$

[23] Using a conventional R^2 yields similar optimal window lengths and network quality, but R^2_{OOS} is a more appropriate metric.

where $K = 5$ is the number of sample splits. This $\overline{R^2_{OOS}}$ statistic measures the neural network's average performance improvement within a given window, relative to the naive baseline. The optimal event window is therefore the length that yields the largest $\overline{R^2_{OOS}}$. This is the point where the network's predictive power and generalisability is highest, which only occurs when the impact of noise is minimised and the fundamental reaction is fully captured.

4.5 Fine-tuning the Neural Network for Approximation

Individual parameters in a neural network are not interpretable like coefficients in a parametric model (Athey and Imbens, 2019). Therefore, this paper does not describe the causal effect of specific words within FOMC communications on asset price reactions. Instead, the network's training process approximates this underlying, complex relationship. I then use its predictions to pin down the event window size that best reflects the market's full reaction to monetary policy communications.

Fine-tuning a 110-million-parameter model like XLNet-Base makes overfitting an inevitable issue, which appears as a gap between R^2_{OOS} and in-sample R^2. I mitigate this difference by tuning two key hyperparameters: the number of training iterations and the learning rate. Too many iterations cause the network to over-fit the training data, deteriorating out-of-sample accuracy. Too few iterations will cause it to under-fit, failing to learn the underlying relationship. The learning rate governs how quickly the network updates its parameters. A rate that is too high causes the neural network weights to update too dramatically, forgetting prior information and degrading out-of-sample predictions. Conversely, a rate that is too low will also result in underfitting, as the network never fully learns the mapping between FOMC communications and asset price changes. Because I use transfer learning, the network's initial parameters are already optimised for general English. My goal is therefore to limit parameter updates within and across training iterations in order to preserve this generalisability while adapting the network to the specific task.

Ultimately, training the neural network is a balancing act: XLNet-Base has to meaningfully approximate the underlying relationship without overfitting, which would make the predictions non-generalisable and create a bad signal. I perform this balancing act using a two-step process. First, I set the maximum training iterations to 2040 steps (120 epochs). Second, I perform a "hyperparameter sweep" on the learning rate using Bayesian optimisation.[24] During this sweep, accuracy metrics from Subsection 4.4 are tracked. Training stops via early stopping: When out-of-sample accuracy begins to degrade. After a 24-hour sweep for each split, I select the learning rate and training iteration that yield the highest R^2_{OOS}. Given the potential for volatile out-of-sample predictions on finite samples, I select the model from this iteration only if subsequent iterations show a permanent decline in out-of-sample accuracy. I use the Simple Transformers library from Rajapakse, Yates, and Rijke (2024), which simplifies fine-tuning Transformer models like XLNet-Base.

4.6 Addressing Look-ahead Bias

A valid concern for text-analysis neural networks is look-ahead bias. As explained in Sarkar and Vafa (2024), this bias occurs when a network is pre-trained on future data to predict past values. The risk stems from large, unknown pre-training corpora, which may include data from the future.[25] I argue that this bias is mitigated in my method for two reasons: the known pre-training corpora of XLNet-Base and the specific language composition of FOMC statements.

First, the authors of XLNet-Base mitigate this concern by restricting the pre-training data to *only* BookCorpus and the entire English Wikipedia. This restriction makes it highly unlikely that the network learned any specific information about the intersection of FOMC communications and futures market price reactions during its pre-training. Importantly, Z Yang et al. (2019) showed that this limited pre-training data does not negatively impact the

[24]The set of considered learning rates is $[1e-5, 9e-5]$, assuming a uniform distribution.

[25]Look-ahead bias is especially prevalent for closed-source large language models (e.g., ChatGPT) whose pre-training data are not publicly accessible.

model's generalisability compared to other models.

Second, look-ahead bias is mitigated because the information content in FOMC statements are temporally anonymised. Similar to Glasserman and Lin (2023), the pre-processed text contains no references to relative times (e.g., t and $t+1$). Therefore, even if XLNet-Base's pre-training data included future FOMC statements, the network has no chronological information to exploit. It cannot "figure out" the temporal order of the statements from the text alone, preventing it from learning from future information.

5 How Long Until Markets Fully React to MP News?

In short, the systematic estimation is searching for the event window length where the neural network's predictive power is highest. As argued in the motivating framework, this is the point where noise has a minimal average impact on the asset price change. This optimal window length allows the network to best approximate the underlying relationship, yielding the most precise and generalisable text-based signal.

Figure 4 presents the results for selected federal funds, Eurodollar, and Treasury futures, as well as the S&P 500 Index. Appendix Figures E2–E8 show the results for all futures contracts. In each sub-figure, the horizontal axis shows the end time of the event window lengths (e.g., $t+20$ implies a 30-minute window ending 20 minutes after release), and the vertical axis is the R^2_{OOS} metric. The cross points show the $\overline{R^2_{OOS}}$ statistic for each window, with the solid yellow point marking the maximum. Box-and-whisker plots show the distribution of R^2_{OOS} across the five sample splits for each window length.[26]

The results show that regardless of futures maturity and asset type, the event window length yielding the largest $\overline{R^2_{OOS}}$ is always at least 40 minutes, ending 30 minutes after FOMC statement release. Examining other distributional metrics from the box-and-whisker plots, such as the 25^{th} or 75^{th} percentiles, leads to the same conclusion. Financial markets

[26]Standard errors are not calculated due to the small number of sample splits ($K=5$). While K could be increased (e.g., via leave-one-out cross-validation), the computational costs are prohibitive for this version of the paper.

appear to fully react within a 40-minute event window *at minimum*, implying the popular 30-minute window is too short.[27]

The optimal event window length increases with the underlying maturity of the futures, a trend summarised in Figure 2. For Federal Funds futures (measuring policy expectations from the current FOMC meeting), the average window is 40 minutes for front-month contracts, rising to 45 minutes for two- and three-month-ahead contracts (measuring policy expectations for the upcoming FOMC meeting). For Eurodollar futures (two to four quarters ahead), the window is 50 minutes. This 50-minute length also applies to 2- and 5-year Treasury futures, but increases to 60 minutes for 10- and 30-year futures. One possible explanation is that traders of longer-horizon assets are more exposed to "soft" information (e.g., forward guidance), which is harder to interpret (Indriawan, Jiao, and Tse, 2021). This may increase belief and information uncertainty and result in higher risk premia (Piazzesi and Swanson, 2008). Relatedly, Okada (2025) document that longer-duration Treasury bonds are more affected by the FOMC announcement premium. Observed liquidity differences may be a symptom of these optimal event window lengths. For example, Fleming and Piazzesi (2005) find abnormal trading volumes persist longer in longer-maturity Treasury markets after FOMC announcements.

For equities, the S&P 500 Index and its front-month E-mini futures market ($ESc1$) fully react within 50 minutes, while the market for second-month E-mini futures ($ESc2$) takes 60 minutes. The neural network estimates that the average optimal window length is 53.3 minutes. These longer-than-30-minute windows may be because equities are less directly impacted by policy rates, leading to more investor disagreement and misinterpretations about interest rate decisions (Zhang, Kappou, and Urquhart, n.d.). The difference between $ESc1$ and $ESc2$, despite both sharing the S&P 500 Index's infinite underlying maturity, may stem from their differing contract expirations. $ESc2$ has a more distant expiration and lower liquidity than $ESc1$. This could cause $ESc2$ traders to weigh "soft" information more heavily,

[27] Although not shown to avoid visual clutter, this conclusion also holds when using the median R^2_{OOS}. The optimal event window is never 30 minutes in length.

requiring more time to fully process the announcement.

Overall, the $\overline{R^2_{OOS}}$ broadly exhibits a "hump" shape. The neural network's predictive performance increases with event window length, peaks at 30–50 minutes after statement release, then declines. This pattern is consistent with the motivating framework. As the window initially lengthens, predictive performance improves as cognitive noise decays and the market's full reaction is captured. After the peak, performance declines as unrelated news begins to dominate the observed price reaction, reducing the network's predictive accuracy. The optimal event window length is therefore the point where the average impact from both noise components is minimised, allowing the neural network's approximation to be the most generalisable. This window captures the full reaction with the least possible noise relative to the baseline of naively predicting with the in-sample average.

Importantly, the neural network achieves a positive $\overline{R^2_{OOS}}$ for all event window lengths and assets, confirming a systematic relationship between FOMC statement text and asset price changes. Relative to the baseline of naively predicting with the in-sample average, the network's generalisability within the optimal window improves by an additional 2–17 percentage points (7.4 p.p. on average) compared to its performance in a conventional 30-minute window.[28] This improvement in predictive performance occurs at the optimal window length precisely because it is the point where the average impact of noise is minimised, allowing the network to best approximate the underlying relationship.

5.1 Different Event Windows, Different Market Responses

I now compare the market price responses measured within the estimated optimal windows to those from the conventional 30-minute window. Figure 5 depicts this comparison for selected futures and the S&P 500 Index. Appendix Figures E9–E15 show the comparison for all interest-rate and equity futures. In each sub-figure, the horizontal axis plots the price log-difference within the optimal window, while the vertical axis plots the 30-minute window

[28] Appendix Table F4 summarises these results for each futures contract.

response. The blue dots represent market reactions on scheduled FOMC meeting dates, and the 45-degree line is shown in grey and dashed.

To characterise whether markets under- or over-react to FOMC statements, I regress the price log-differences from the optimal windows on those from a 30-minute window. If the slope coefficient is statistically less than one for at most $\alpha = 0.05$, I define the market as under-reacting to the FOMC statement. Otherwise, the reactions are not considered statistically different. In each sub-figure, the regression line is coloured red for under-reaction and grey otherwise.

Financial markets for all interest-rate futures appear to under-react to FOMC statements, ex post. Furthermore, the slope of the regression decreases as the underlying maturity increases. This is consistent with the post-FOMC announcement drift described in Indriawan, Jiao, and Tse (2021) and Brooks, Katz, and Lustig (2023), who both point to information processing limitations as a source of bond market under-reactions. Indriawan, Jiao, and Tse (2021) argue that interest-rate markets are more exposed to "soft" information (e.g., forward guidance), which is costlier to process. Brooks, Katz, and Lustig (2023) find that slow adjustments by mutual fund investors also play a role. The 16 December, 2008 FOMC meeting provides a notable example. In response to the Great Recession, the FOMC cut the federal funds rate to a 0–0.25% range. This decision was a surprise, as markets expected a smaller reduction to 0.25–0.5%. This surprise, combined with the statement's commitment to "employ all available tools to promote the resumption of sustainable economic growth and to preserve price stability", likely contributed to markets needing more time to fully react. The top sub-figure of Figure 6 illustrates this: for all Eurodollar futures, the price change becomes increasingly positive as the event window expands beyond the initial 30-minute reaction.

The S&P 500 Index also appears to under-react, ex post, corroborating findings from Neuhierl and Weber (2024) and Golez, P Kelly, and Matthies (2025), who suggest market participants struggle to react to the full information in FOMC announcements. However, the S&P 500 Index also exhibited the most price reversals, often moving from negative to positive as the event window expanded. A key example is the 16 September, 2008 meeting, shown in the bottom sub-figure of Figure 6. The market initially responded negatively, but this trend reversed within a 40-minute window and continued to increase until market close. One explanation is that market beliefs reversed as participants processed the FOMC's decision to maintain interest rates. With more time, cognitive noise may have died out, allowing a negative "knee-jerk" reaction to reverse as the market fully processed the positive information. This aligns with Bordalo et al. (2024), who document that stock market overreactions are common and tend to reverse as beliefs correct with time.

Overall, these findings support the premise of my paper: the optimal event window length is not a parameter to be simply assumed. Instead, I document how the window length that best reflects the market's full reaction to FOMC statements varies systematically across asset types and underlying maturities. Conventional, shorter windows fail to capture the market's full processing of complex information, revealing systematic under-reactions in interest-rate markets and significant price reversals in equities that are otherwise missed.

6 What Happens to Monetary Surprises and Shocks?

This paper has shown that regardless of asset type and underlying maturity, the optimal event window length is always longer than the standard literature assumption of 30 minutes. However, how are monetary policy surprises and shocks affected when changing to the optimal window lengths?

Proper investigation into this question requires that I use a common event window length when constructing the surprises because the "invertibility", or "spanning" (Duffee, 2013), of information about macroeconomic variables from the yield curve is well-defined for a temporal cross-section of yields. In contrast, it is not currently known if these macro-finance term structure models are well-defined for different points of time on the yield curve. As a result, I construct monetary policy surprises for each of the window lengths found to be optimal by the neural network: *40-, 50-, and 60-minute windows.* The rest of this section uses the same data and considered FOMC meeting dates as described in Section 3.

6.1 Effects of Window Choice on Monetary Policy Surprises

For each FOMC meeting date, I construct interest-rate surprises within 30 minutes and the three found optimal event window lengths. Following Kuttner (2001), Gürkaynak, Sack, and Swanson (2005), Nakamura and Steinsson (2018), and others, I use federal funds rate to cover interest rate expectations up to three months out, Eurodollar futures to capture expectations from about four months to one year ahead, and Treasury futures for interest rate expectations out to fifteen years. Appendix C provides details behind the conversion from changes in futures contract prices to interest-rate surprises.

Figure 7 plots the correlation between the monetary policy surprises calculated within the optimal event window lengths v. 30 minutes. The horizontal axis depicts the underlying maturity of the surprises (e.g., "0–1M" represents the interest-rate surprise for the FOMC meeting happening in the current or next month). The red-dotted, yellow-dashed, and green-mixed lines represent the correlations calculated for surprises measured within the optimally-found 40, 50, and 60 minutes, respectively. The blue-solid line represents the correlations averaged across the three optimal window lengths. For all underlying maturities of all surprises, we see that the correlations decline when increasing the optimal event window length. Increasing the underlying maturity of the surprises also sees the correlation fall by a larger magnitude under longer optimal window lengths. For example, while monetary policy expectations about current and next FOMC meetings are more or less identical regardless of optimal window choice, the differences between 30- and 50-minute windows result in a 10% decline in the coefficient when looking at monetary policy surprises out to fifteen years ahead. In other words, as expectations about monetary policy are further into the future, calculating interest-rate surprises become more sensitive to the choice of event window lengths imposed around FOMC statement releases.

6.2 Effects of Window Choice on Monetary Policy Shock Impacts

In this subsection, I use the interest-rate surprises calculated within 30 minutes and the optimal event window length as financial instruments to identify exogenous variation in monetary policy through several popular methods in the literature. I choose to include interest-rate surprises from Treasury futures in my instrument set for two reasons. First, recent studies (e.g., Brennan et al., 2024; An, Stedman, and Lusompa, 2025) have shown that using instruments with longer underlying maturities can help prevent the understatement of monetary policy stimulus during the effective lower bound period.[29] Second, my findings from Figure 7 imply that only changing the event window length causes larger effects to Treasury futures markets than shorter maturity assets.[30] For conciseness, I choose to use the *median* of the window lengths found to be optimal by the neural network: *10 minutes before and 40 minutes after statement release.*

Shock Construction I construct monetary policy shocks following the construction methods of Gürkaynak, Sack, and Swanson (2005), Nakamura and Steinsson (2018), and Jarociński and Karadi (2020) (henceforth GSS, NS, and JK, respectively).[31] Rather than keep track of the entire expectations path of interest rates implied by futures prices, the authors of these papers reduce this dimensionality using principal component analysis on multiple

[29] For my sample period, the effective lower bound of the federal funds rate is defined as defined as 16 December, 2008 through 16 December, 2015.

[30] As a robustness check, all exercises performed in this subsection are re-done using the original instrument set of the authors behind the monetary policy shock construction methods. The corresponding figures and tables can be found in Appendix D and show that the results from the main text are overall robust to the choices of interest rate-surprises as instruments.

[31] I summarise the names, notation, and construction of the monetary policy shock series in Appendix Table F5.

monetary policy surprises to extract the common variation in one or two dimensions.[32,33] GSS use the first and second principal components to identify two shocks (factors) due to the multi-dimensional nature of monetary policy. The authors rotate both components such that the first component drives changes in interest-rate surprises for the current federal funds rate, while the second component has no effect on interest-rate surprises for the current federal funds rate. This rotation yields the "target" and "path" factors of monetary policy, which I call GSS_T and GSS_P, respectively. Because principal components have no interpretable units, GSS re-scale both factors such that GSS_T is one-for-one with the interest-rate surprise for the current federal funds rate while GSS_P is re-scaled to have equal impact on the implied rate surprise from the four-quarter Eurodollar futures contract. NS and JK use only the first principal component because, as explained in Bauer and Swanson (2023), the first principal component is essentially a weighted average of the target and path factors. NS re-scale their first principal component to be one-for-one with the daily change in the zero-coupon, nominal one-year Treasury yield, which I denote as NS_{MP}. JK first differentiate from prior studies by re-scaling the first principal component to have equal standard deviation with that of the implied rate surprise in the four-quarter Eurodollar futures. To account for possible information effects from the central bank, JK then impose sign restrictions on the first principal component based on its co-movement with stock market changes.[34] The series that has negative (positive) co-movement with stock prices is considered to be the monetary policy (central bank information) shock, which I call JK_{MP} (JK_{CBI}).

For easy interpretation and comparison between all monetary policy shock series, I re-scale all shock series to be one-for-one with the daily change in the zero-coupon, nominal

[32]Principal component analysis uses eigenvalue decomposition of a given dataset's covariance matrix to project the data onto new dimensions based on its variance. The first principal component, representing the first coordinates on the new dimensions, captures the largest common variation of the original variables. The second principal components captures the second largest common variation and so on.

[33]The main results of JK use the implied interest-rate surprises from $FF4$. However, the first principal component of multiple surprises are used in the modern update of the paper's shock: https://github.com/marekjarocinski/jkshocks_update_fed

[34]I calculate stock market changes within the same event window length used to measure the monetary surprises.

one-year Treasury yield.[35] Appendix Table F6 provides summary statistics of each of the shock series. Because these methods together account for the various characteristics of monetary policy, such as its impact beyond the immediate horizon, its multi-dimensional impact, and the possible existence of information effects, any differences found in the shocks and their effects from this exercise would be the impact of event window choice.

Visual Shock Differences To begin my investigation, I plot the shock series constructed in both event window lengths, displayed in Figure 8. For each sub-figure, the horizontal axis depicts the FOMC statement release dates. The vertical axis is in units of percentage points after re-scaling each shock to be one-for-one with the daily change in the nominal one-year Treasury yield. The black-solid and red-dotted lines represent the shock series derived from 30-minute and optimal window lengths, respectively. Regardless of construction method, we see that much of the peaks and troughs of shocks derived within 30 minutes see larger values in magnitude at many of those FOMC statement release dates for shocks constructed using optimal window lengths.[36] For example, both the December 2008 and March 2009 FOMC meetings see larger negative shocks by 3–6 basis points in magnitude when using the optimal event window length, possibly suggesting that the full impacts of the information content within the FOMC statements aren't yet captured within 30 minutes.

Another episode worth pointing out is the August 2011 FOMC meeting date, where the statement maintained the federal funds rate range between zero and 25 basis points, but was the first communication that included explicit dates on the expected path of monetary policy in its forward guidance language (Crump, Eusepi, and Moench, 2013). We can see that the optimal event window length is associated with a more negative monetary policy shock in both GSS_P and NS_{MP} on this release date. When imposing the sign restrictions of Jarociński and Karadi (2020), the negative central bank information shock increased in

[35]The robustness checks performed in Appendix D show that the choice of scaling has no overall effect on the findings of the main text.

[36]Plots for GSS_T are not shown because regardless of window choice, the shocks are essentially identical. This similarity makes sense since markets of all underlying maturities need little time to fully react to the immediate federal funds rate change in the FOMC statements.

magnitude by 3.6 basis points while the monetary policy shock stayed roughly the same. One possible interpretation of these differences is that because the federal funds rate was already at the zero lower bound, language indicating that the future path of the federal funds rate would remain at these low levels could be interpreted as an information surprise about weaker economic conditions than initially anticipated. The full extent of this negative information shock might only materialise when financial markets are given enough time to react to the statement.

On the other hand, changing event window lengths can result in "shifts in importance" in the composition of the monetary policy shock. When comparing JK_{MP} and JK_{CBI} derived with a 30-minute window on the September 2008 FOMC meeting, the monetary policy shock is the main effect from the statement release, while the central bank communication shock is close to zero. When using a 50-minute window however, this interpretation flips. Looking at the statement from this FOMC meeting, we can see that the forward guidance language from the central bank suggested moderate expected growth due to accommodative monetary policy and its efforts to promote liquidity in financial markets. Combined with its decision to hold the federal funds rate range steady, market participants would likely update upwards their views about economic prospects. In line with a positive central bank information shock, the S&P 500 increased in value during trading hours.

When comparing the identified shocks from JK overall, JK_{CBI} sees a higher frequency of increased amplitude when increasing the event window length. Although many FOMC meeting dates see increased amplitude for the JK_{MP} shock, the actual size of these differences due to event window choice is relatively smaller. One likely explanation for the larger differences in the central bank information shock is a combination of the changes to the S&P 500 Index from event window choice and the imposed sign restrictions of the construction method.

Responses of Interest Rates I now look at how the event window choice affects the impact of monetary policy shocks on nominal and real interest rates for different maturities. I use the daily change in Treasury yields and daily change in TIPS yields to represent nominal and real interest rates, respectively. For both sets of dependent variables, I calculate the daily changes using the end-of-day yields for the day of and before the FOMC announcement. Note that both the 30-minute and optimal event window lengths used to construct the monetary policy shocks are nested within the daily window for Treasury and TIPS yield changes. This timing restriction prevents the yield changes to affect the shocks. The regression specification is as follows:

$$y^{i,j} = \beta_0^{i,j,k,l} + \beta_1^{i,j,k,l}(Shock)^{k,l} + \varepsilon^{i,j,k,l} \tag{8}$$

where i indicates daily changes in Treasury or TIPS yields as the dependent variable, j indexes the horizon of the yield, k denotes the considered monetary policy shock used as the regressor, and l denotes whether the shock was constructed within the optimal event window length or 30 minutes.

Regression results for the responses of nominal and real interest rates to the monetary policy shocks within different event windows are presented in Table 1 and Table 2, respectively. In both tables, each estimate comes from a separate OLS regression of the form described in Equation 8. Starting from the left, columns 2–5 (2–4) of Table 1 (Table 2) contain estimates coming from regressions where the dependent variable is the daily change in nominal (real) yields and the independent variable is a monetary policy shock series constructed within 30 minutes around scheduled FOMC announcements. Estimates in columns 6–9 (5–7) have equivalent meaning, but the shocks are constructed within the median optimal event window of 50 minutes instead. Bolded columns 10–12 (8–10) are the differences between the corresponding coefficient estimates, representing the effect from event window choice.

When opening up the event window around scheduled FOMC announcements from 30 minutes to the optimal length, all monetary policy shocks that capture the effects of forward guidance have larger effects on both nominal and real interest rates. For example, increasing the event window length results in JK_{MP} to have an effect that is 26, 24, and 23 basis points larger on the 2-, 5-, and 10-year real yields measured using TIPS, respectively. The statistical significance of the coefficients at the 1% level remain when using the optimal event window length. Interestingly, the target factor under the methodology of Gürkaynak, Sack, and Swanson (2005), GSS_T, has smaller coefficient estimates for nominal interest rates at all horizons when using the optimal event window length, but an unclear direction for real yields. Only estimates for the 2-year nominal yield are statistically significant when the target factor is constructed under both window lengths. One likely explanation for these results is that the full sample period includes the effective lower bound, a time period where the stimulus of monetary policy through changes in the current federal funds rate was limited and greater emphasis was put into forward guidance.

The main takeaway from the regression results of both tables is that measuring the reaction of financial markets to monetary policy announcements within too small of an event window prevents market participants from fully processing the "soft" information content of the announcements, especially those pertaining to the future path of interest rates. As a result, the effects of monetary policy shocks about forward guidance on financial market variables can become hampered. Indeed, these regression differences support an interpretation suggested by Gürkaynak, Sack, and Swanson (2005), which is that financial markets can fully react to announced changes in the federal funds rate target within a short amount of time. However, information within the FOMC statements about the outlook of policy and economic conditions are more complex and require additional time to fully process and price in by market participants.

Responses of Break-even Inflation I also investigate the effect of window size on the impact of monetary policy shocks on break-even inflation as measured by the difference between nominal Treasury rates and TIPS rates. The regressions follow the same form as Equation 8, but the dependent variable is now the break-even inflation measure. Starting from the left in Table 3, the columns 2–4 contain estimates coming from regressions where the dependent variable is the daily change in break-even inflation at various maturities and the independent variable is the monetary policy shock, both constructed within 30 minutes. Columns 5–7 have equivalent meaning, but both variables are constructed within the median optimal event window length. Columns 8–10 contain the percentage change between the corresponding coefficient estimates.

Similar to the findings of Nakamura and Steinsson (2018), a majority of the estimates are not statistically significant. However, there are two interesting insights from the exercise. First, expanding the event window size to the median optimal length of 50 minutes results in an increase in the statistical significance for certain shock estimates. For example, the decline in break-even information at the furthest horizon from NS_{MP} is statistically significant at the 5% level when using the optimal window length. A similar conclusion is reached when looking at the effects to break-even inflation from GSS_T. Second, there are several instances where monetary policy shocks containing forward guidance negatively impact break-even inflation by several basis points. Like the takeaway from the regression results of interest rates, expanding the event window length to the optimal size results in the impact of monetary policy shocks to become stronger and more precise.

Responses of Equities I finish my analysis on the effects of event window size on financial variables by studying equity prices. Table 4 presents analogous results for the S&P 500 Index and its E-mini futures. The form of the regressions is the same as Equation 8, except the dependent variable is now 100 times the price log-difference of equities. From the left, the columns 2–4 contain estimates coming from regressions where the dependent variable is the price change in the S&P 500 Index and its E-mini futures and the independent variable is the monetary policy shock, both constructed within 30 minutes. Columns 5–7 have equivalent meaning, but both variables are constructed within the median optimal event window length. Columns 8–10 contain the percentage change between the corresponding coefficient estimates.

Increasing the event window around scheduled FOMC announcements from 30 minutes to the optimal length causes all monetary policy shocks containing forward guidance to impact stock prices more negatively. Opening up the event windows results in the decline of the S&P 500 Index from JK_{MP} to grow by 18.25%, while GSS_P and NS_{MP} negatively affect stock prices by a more severe 11.51% and 1.23% when using the optimal event window length, respectively. The differences in negative impacts to equity prices become larger in magnitude because of event window choice when examining the responses of the second-month S&P 500 E-mini futures contract. Interestingly, the larger effects of event window choice on the stock price changes have much larger variance across the monetary policy shocks when compared to the effects seen on interest rate changes.

Similar to the earlier investigation on interest rates, the response of stock prices to the target factor from GSS becomes weaker by 11.99% when using the median optimal window length. Although the estimates are still statistically significant at the 1% level and go in the correct direction, the weaker effects are possibly due to the event window lengthening causing a "shift in importance" towards the path factor. Using the optimal window length allows more time for the stock market to fully react to the FOMC statements. The longer window captures market participants quickly reacting to a negative shock to the current federal funds rate, then fully reacting to process the effects of tightening forward guidance and other "soft" information. With regards to the central bank information shock, one possible explanation for the difference in effects could have to do with the aforementioned combination of the larger idiosyncratic volatility of stock prices (relative to Treasury yields) and the imposed sign restrictions of JK. As discussed in Subsection 5.1, the S&P 500 Index under-reacts to monetary policy announcement ex post, but also experiences frequent instances of over-reaction through price reversals. By expanding the event window to the median optimal length of 50 minutes, these two factors can result in different decompositions of the central bank information shock, amplifying the series in both directions and weakening its overall effect on stock prices.

I note that the effects of event window choice on the Jarociński and Karadi (2020) shocks shouldn't be interpreted in isolation because of the sign restrictions method behind extracting JK_{MP} and JK_{CBI}. Because the two shocks are defined based on the co-movement with changes in equity prices, this restriction should naturally cause both shocks to have large impacts on the S&P 500 Index and its E-mini futures. Nonetheless, it is still seen that other monetary policy shocks containing forward guidance, like GSS_P, negatively impact equity prices and futures by similar magnitudes to those from JK_{MP}.

Responses of the Macroeconomy via Local Projections To study the effects of event window choice on the impact of monetary policy transmission into macroeconomic variables, I estimate impulse responses using a local projection approach. Following Gertler and Karadi (2015) and Bauer and Swanson (2023), I include the log of CPI, the log of IP, the nominal two-year Treasury yield, and the EBP. Measured by Gilchrist and Zakrajšek (2012), the EBP represents the effective risk-bearing capacity of the financial sector and is essentially the difference between private and public bonds. Including the variable in the model specification is what allows for the transmission of monetary policy shocks from financial markets into economic variables. In order to have all variables be the same frequency, I convert all shock series into monthly frequency such that months with FOMC meetings have a monetary policy shock of zero. Summary statistics for the monthly specification variables and shock series can be found in Appendix Tables F7 and F8.

I use the lag-augmented local projection method from Olea and Plagborg-Møller (2021) to estimate and graph the impulse response functions of each outcome variable from each monetary policy shock series. The primary reason why I choose this approach is because adding an additional lag to the local projection specification as controls results in correct inference with using only Eicker-Huber-White heteroscedasticity-robust standard errors, avoiding the need to correct for serial correlation in the local projections. Therefore, I run a separate regression of each outcome variable on each shock:

$$y_{t+h}^{i,l} = \theta^{i,k,l}(Shock)_t^{k,l} + controls + \eta^{i,k,l} \tag{9}$$

where i indexes one of the four outcome variables, k denotes the considered monetary policy shock, l denotes whether the shock was constructed within the optimal event window length or 30 minutes, and h is the number of months in the future. The impulse response functions are responses of macroeconomic variables to a 100 basis point increase in the monetary policy shock from the common re-scaling to be one-for-one with the zero-coupon, nominal one-year Treasury yield. For all the shock series, this represents a contractionary policy shock.

Figure 9 visually compares the impulse responses of CPI to the different monetary policy shock series constructed in either the optimal event window length or 30 minutes.[37] Initial impressions are that the point estimates of the impulse responses are qualitatively similar regardless of event window choice for all outcome variables. While there exists some differences (e.g., the decline in CPI is more severe in response to NS_{MP} at farther horizons when the shock is constructed within the optimal window length), the lengthening of the event window doesn't result in dramatic changes in the impulse responses. However, what is visually noticeable is that the confidence intervals from the optimal window length are smaller than those from the 30-minute window for most responses. To formally measure this difference, I first calculate the width of the confidence intervals for each time horizon of each macroeconomic variable response to each monetary policy shock series constructed within the optimal window length or 30 minutes. I then take the ratio of the confidence interval width using the optimal window length to the confidence interval width using the 30-minute window. This ratio being less (greater) than one implies that the impulse responses to monetary policy shocks constructed using the optimal window length are more (less) precise compared to those when using 30 minutes. Calculating the average and median

[37]Because the discussed results in the main text are robust to the response comparisons of the other macroeconomic variables, they are visually presented in Appendix Figures E16–E18.

confidence interval width ratio yields 0.9173 and 0.9410, respectively. In other words, constructing monetary policy shocks within the optimal median window length of 50 minutes results in *more precise* responses of macroeconomic variables.

Overall, the heterogeneity in the differences amongst the responses of interest rates, stock prices, and the broader macroeconomy support the findings of Brennan et al. (2024): the data on long-term rates and construction method used for monetary policy shocks are important contributors. I provide evidence that the dimension of event window length must also be considered. By choosing the popular 30-minute window in the literature, the effects of monetary policy shocks about forward guidance on financial variables can more or less be dampened because markets have not yet fully reacted to the policy announcements. Furthermore, constructing monetary policy shocks within the event window length where the average impact of noise is minimised results in the estimated responses of macroeconomic variables to be more precise.[38]

7 Statement Features and Their Effects on Windows

Recall that systematic estimation is the neural network regressing the changes in asset prices within different windows on the text of FOMC statements. In other words, systematic estimation could be thought of as "jointly" estimating the and optimal event window length and text-based signal with greatest precision. Unfortunately, the financial and mandatory hardware constraints required by the neural network to perform this "joint" estimation are why the current version of this paper is only able to consider event window lengths ending up to 60 minutes after FOMC statement release.

An alternative method that potentially circumvents these constraints and provides some initial insights is what I call the "one signal" approach. First, assume that the predicted price changes within the "jointly" estimated optimal window (i.e., the most precise signal) is *constant* for all time. As a result, event window lengths longer than an hour after FOMC

[38] These confidence interval width ratios are robust to the choice of monetary policy surprise instruments.

statement release can estimated by following the motivating framework more closely. Before detailing the analyses, I note that the core assumption of the "one signal" approach is a strong one. Because the weights and embeddings of XLNet-Base are sensitive to changes in both its input and output variables, it is entirely possible that the "jointly" estimated event window length is different from that which has the largest $\overline{R^2_{OOS}}$ under the "one signal" approach.[39]

The "one signal" approach uses the text-based signal from the "jointly" estimated optimal event window and calculates $\overline{R^2_{OOS}}$ for all event window lengths longer than the optimal window length. This shortcut provides two advantages. First, I can observe how the different dimensions of monetary policy communication potentially affect the optimal event window lengths. Second, I can check if there exists an event window beyond the optimal window length that has a larger $\overline{R^2_{OOS}}$, which I perform and discuss in the Appendix D.

7.1 How FOMC Statement Characteristics Affect Window Lengths

An interesting question to consider is if different aspects within the text of FOMC statements affect the optimal event window length. I use the "one signal" approach to provide some insight into whether the complexity of the information content in FOMC statements, the degree of similarity between sequential statements, and the mention of dissents within the statements change the optimal event window lengths. These heterogeneity exercises rely on the core assumptions that the text-based signal associated with the "jointly" estimated optimal event window is *constant* for all time *and* that this signal is still relevant when conditioning the FOMC statements into subsamples.[40]

Text Complexity I first observe how the optimal event window length changes for futures maturities when conditioned on the complexity of the FOMC statements. I assess the *semantic* complexity using the standard readability formula of the Flesch-Kincaid Grade

[39] A rough analogy is how iterative generalised method of moments does not always not converge.

[40] The ideal solution would be to acquire large subsamples of FOMC statements satisfying the these characteristics, then have XLNet-Base systematically estimate the optimal event window lengths for all characteristic subsamples.

Level index. These readability formulas determine the grade level required for a reader to properly read and comprehend a provided text based on sentence structure (i.e., average sentence length), word structure (i.e., percentage of difficult words in a sentence), and word phonology (i.e., average number of syllables per word).[41]

I first split my FOMC statements depending on whether their grade levels are below or above the median grade level (i.e., 16.5, which is the comprehensive level of some degree of a master's education). I then calculate the MSEs under the "one signal" approach for each subsample of FOMC statements for every futures maturity. A summary of this exercise is found in the first two columns of Table 5, which displays the minimised MSEs and associated event window lengths averaged across all asset types and maturities. Immediately, one can see that FOMC statements with a required readability level of at least some of a master's degree have a longer average optimal event window of 71 minutes. In contrast, "simpler" statements with a grade level up to some of a master's degree is associated with a shorter average event window length of 60 minutes. Interestingly, both subsamples of FOMC statements possess similar averaged MSEs. One potential explanation for this combination is that the price responses of financial markets to "complicated" statements are relatively affected more by cognitive noise. This association is consistent with the idea that more complex information results in differing interpretations and reactions to said information, causing increased market volatility and trading volume. Indeed, papers such as Smales and Apergis (2017) have shown that the increasing complexity of FOMC statement language over time is associated with increasing volatility across equity, bond, and currency markets. In other words, financial markets might need more time to fully react to FOMC statements with more complicated language.[42]

[41]Appendix Figure E19 displays the evolution of the Flesch-Kincaid Grade Level for my sample period of FOMC statements. Descriptive statistics for the text complexity measure can be found in right column of Appendix Table F9.

[42]Further research is needed into understanding how the overall complexity of monetary policy communications affects the ability of people, even technically trained professionals like financial market participants, to fully comprehend and respond to these announcements. As shown in McMahon (2023) for example, efforts by the Bank of England to simplify their communications have resulted in decreases to semantic complexity over time, but a simultaneous increase in motivating complexity over time.

Text Similarity Prior studies (e.g., Acosta and Meade, 2015; Handlan, 2022a) have documented that while FOMC statements are relatively concise, the semantics and language have varied more often than typically assumed over time. Using the "one signal" approach, I also analyse how the optimal event window length changes when conditioned on the degree of similarity between the FOMC statements. The method I use for this exercise is a bag-of-words model, which represents text as an unordered, weighted-frequency of words. Importantly, this method assumes that documents with similarly weighted frequencies for words are discussing similar topics. By transforming the FOMC statements into vectors of word frequencies, I can perform vector analysis and create a measure of similarity between any two document vectors.

The weighted frequency of words frequently used in text analysis is called Term Frequency-Inverse Document Frequency (TFIDF). Conversion of my sample of FOMC statements into the TFIDF matrix is done through the Python text analysis library of sklearn. Specifically, each word in a document is multiplied by the number of times it occurs within a particular document over the frequency of documents where the word appears. The purpose behind this ratio is to properly weigh down the importance of frequently occurring words that don't actually provide much meaning with regard to distinguishing FOMC statements from one another. For example, words such as "a" and "the" appear in all documents with high frequency, but don't signal anything about any particular statement. Therefore, the value for these words are weighed down by the high number of occurrences across documents, yielding a *weighted frequency*. Additionally, words like "unemployment", "inflation", and "federal funds rate" will also be weighed down because monetary policy discussed in every FOMC statements has these components. However, it should be noted that these words are weighed down by a smaller degree relatively speaking because they could still convey some importance when discerning the degree of similarity between two statements (e.g., if one statement discusses about inflation relatively more than another statement that focuses more on labour market conditions). Lastly, words that do not appear in every FOMC statement, such as "persists" or "tight", are given a higher weighted frequency because they potentially signal a unique economic environment when the FOMC meeting took place compared to other meeting dates. The differences in these unique and informative words are what allow me to calculate the degree of similarity between FOMC statements.

The calculation of the TFIDF matrix is as follows: Let D be the set of FOMC statements and T be the set of all terms that appear within and across all statements in D. Let a single document and term be indexed by $d \in D, t \in T$, respectively. Therefore, the term frequency of a particular term t in document d can be defined as:

$$tf_{d,t} = \ln\left(\frac{tc_{d,t}}{nt_d}\right) + 1, \tag{10}$$

where $tc_{t,d}$ is the number of times term t appears in document d, nt_d is the total number of words in d, and both the natural logarithm and addition term serve to smooth out the frequency measure. We can then calculate the inverse document frequency as follows:

$$idf_{d,t} = \ln\left(\frac{nd}{df_{d,t} + 1}\right) + 1, \tag{11}$$

where nd is the number of documents in set D and $df_{t,d}$ is the number of documents that term t appears. Multiplying the two terms gives us the weighted frequency of each term t in document d:

$$TFIDF_{d,t} = tf_{d,t} * idf_{d,t}. \tag{12}$$

Calculating $TFIDF_{d,t} \forall d \in D, t \in T$ and combining them yields a $D \times T$ matrix, where each element $TFIDF_{d,t}$ represents the number of times a term appears in a particular document, divided by the number of words in said document, then multiplied with the document frequency (ratio between the number of documents where the term appears and the total number of documents). Essentially, the higher the TFIDF value is for a term, the more informative that word is with regard to distinguishing the information content of those FOMC statements from others.

Although the TFIDF matrix can be produced on the current pre-processed FOMC statements, the results might not be accurate in representing the similarity of the intended or comprehended information content of the statements. For example, the method could potentially view the terms "Federal Funds Rate" and "federal funds rate" as different. In other words, this bag-of-words model is not as "sophisticated" at discerning context from the FOMC statements in their original grammatical structure compared to the neural network, causing the $TFIDF_{d,t}$ values to be inaccurate. I mitigate this issue by deploying additional preprocessing steps on my sample of FOMC statements. First, I make every word lowercase. Second, I remove common words that convey little semantic meaning, such as articles, pronouns, and conjunctions.[43] The last step involves converting all words into their "base" form (e.g., the words "increases", "increasing", and "increase" are all combined into the base form of "increas"). Calculating the TFIDF matrix on the additionally cleaned FOMC statements from May 1999 to October 2019 yields a matrix with 165 rows (statements) and 966 columns (terms).[44]

With the TFIDF matrix calculated, I am able to produce a corresponding matrix whose values represent a measure of similarity between any two pairs of FOMC statements. Specifically, this document similarity matrix is created by multiplying the TFIDF matrix with its transpose:

$$\text{Document Similarity Matrix} = TFIDF \cdot TFIDF^T. \tag{13}$$

Consider a document d row vector of the TFIDF matrix. For each term $t \in T$, the $TFIDF_{d,t}$ values are only positive if the terms appear in document d, otherwise the value is zero. Additionally, the vectors for every document all have the same magnitude due to normalisation by the number of documents and terms in the matrix. Because the product between the TFIDF

[43] I choose to retain numbers because the FOMC's targetted federal funds rate or discussions about forward guidance convey important semantic content.

[44] Appendix Table F10 presents a list of 30 base terms that have the highest TFIDF scores for my sample of FOMC statements. Recall that terms with higher TFIDF values are more informative terms with regard to differentiating the FOMC statements with them from statements without.

matrix and its transpose is essentially the dot product between every pair of document row vectors in the sample of FOMC statements, I can define the degree of similarity between any two FOMC statements as the cosine of the angle between these documents, also called the *cosine similarity*:

$$S^{A,B} := \cos\theta = \frac{\mathbf{A}\cdot\mathbf{B}}{\|\mathbf{A}\|\|\mathbf{B}\|}, \tag{14}$$

where $\mathbf{A}, \mathbf{B}$ are the vectors for FOMC statements A and B, respectively. The more similar two FOMC statements are, the angle between the two corresponding vectors approaches zero, resulting in the cosine of the angle to go to one. A cosine similarity value of one between two FOMC statements means they are identical, while a value of zero means the two documents have no common words. In other words, the more base terms shared between two FOMC statements, the higher cosine similarity value is assigned to the pair.

The Document Similarity Matrix for FOMC statements from May 1999 to October 2019 allows for me to compare the degree of similarity between any two documents for any points of time d and d'. The series of document-pairs $(d, d-1)$ represents the degree of similarity between sequential FOMC statements, which I call as S^1 for the rest of the paper and is plotted in Figure 10.[45] As time progresses, the cosine similarity between sequential FOMC statements increases. The increased persistence and standardisation of FOMC statement wording begins roughly between early-2009 and early-2010, which coincides with the period when the FOMC was in the process of reducing the federal funds rate towards the zero lower bound. As discussed in Handlan (2022a), one explanation behind this occurrence is that the FOMC wanted to avoid market surprises through reducing the variation in statement vocabulary after the financial crisis. Even when the zero lower bound ended with the Board raising the federal funds rate at the end of 2015, the cosine similarity measures remained high through 2019 due to the efforts by the FOMC to increase the transparency of its monetary policy under Chairs Yellen and Powell.

[45] Appendix Figure E20 represents the Document Similarity Matrix for my sample of FOMC statements as a heat map. Descriptive statistics for S^1, the cosine similarity between sequential FOMC statements, can be found in Appendix Table F9.

I split my FOMC statements depending on whether S^1 value is below or above the median (i.e., 0.88), labelling the subsample of statements as being "different" or "similar" relative to their immediately previous document, respectively. Calculating the MSEs for each subsample of FOMC statements for every futures maturity, I present the results of this exercise in the third and fourth columns from the left of Table 5. "Different" FOMC statements observe the smallest MSEs within an event window length of 62 minutes when averaged across all asset types and maturities, while "similar" statements see minimised MSEs within a 51-minute window on average. Similar to the condition of text complexity, both subsets of statements have similar minimised MSEs on average. Intuitively, if an FOMC statement contains less similar base terms and associated ideas with its predecessor, financial markets could require more time in order to discern and react to the new information of the current FOMC statement. Conversely, markets would presumably react to statements that share a lot of terms and concepts with past policy announcements in a short amount of time because most of said information has already been incorporated into the price.

Presence of Dissents The final characteristic of the FOMC statements I consider under the "one signal" approach is whether the presence of dissents in the monetary policy announcements affects the optimal event window lengths of asset types. Similar to prior studies, I find that roughly 40% of FOMC statements in my sample have recorded dissents. While this suggests dissents are not rare, they still provide important signals. According to Federal Reserve tradition, dissents are usually recorded only when Board members find the majority's opinion unacceptable (C Madeira and J Madeira, 2019). In other words, the very existence of a recorded dissent provides additional information for markets to process, which in turn could affect their reaction time. To investigate this characteristic, I split my sample of FOMC statements depending on whether the policy votes were unanimous or not.[46] The fifth and sixth columns of Table 5 present the calculated MSEs for each subset of statements. On average, statements with(out) dissents observe the smallest MSE under

[46] I do not record and distinguish between dissent votes for tighter policy, easier, policy, or other reasons.

the "one signal" approach within an event window length of 83 (53) minutes. The longer reaction time for statements with dissents may be because they introduce additional layers of information regarding policy transparency and internal disagreement, requiring more time for markets to process compared to the unanimous announcements. It should be noted that prior studies (e.g., C Madeira and J Madeira, 2019; Tsang and Z Yang, 2024) have shown that the presence of dissents have a nuanced impact on financial markets. While equity markets have a strong reaction to disclosed information contained in dissents, interest-rate markets exhibited muted impulse responses. My findings show that when averaged across all asset types (i.e., interest-rate based and equities), financial markets are impacted through needing more time to fully react to FOMC statements with the presence of dissents.

Overall, I show that monetary policy communication affects the response of financial markets along multiple dimensions beyond the "hard" information within the FOMC statements (i.e., the change to the federal funds rate). Because monetary policy communication is expressed with varying degrees of semantic complexity, similarity to prior policy announcements, and heterogeneous voting records, these characteristics overall can influence the time financial markets need to fully respond to the "soft" information within FOMC statements, such as forward guidance. Indeed, future research towards understanding the effects of these dimensions could provide deeper insight into how monetary policy communication exactly affects financial markets and in turn the broader macroeconomy.

8 Conclusion

The choice of event window length for measuring financial market responses to news has largely remained an "ad-hoc" decision in the empirical monetary policy literature. This paper challenges the conventions that have emerged from this practice by asking a crucial question: how does one choose the appropriate event window length for their study in monetary policy? By combining observed price dynamics with a text-based signal derived using neural networks methods for text analysis, I develop and implement a methodology to systematically estimate the event window lengths that best reflect the full market reaction to the information content of FOMC statements.

My systematic estimation yields two key findings that directly confront the common assumptions in the literature. First, the common 30-minute window is insufficient time for markets to fully react to FOMC announcements. My results show that, on average, markets fully react within an event window *ending at least 30 minutes after release*, corresponding to a total window length of at least 40 minutes. A 30-minute window is never found to be optimal for any asset type considered. Second, the optimal window length is not a universal parameter. In line with this paper's title, one window does not fit all; the optimal duration increases systematically with asset underlying maturity. This estimated length rises from 40 minutes for assets with the shortest underlying maturities to 50-60 minutes for futures with underlying maturities of two quarters or more. Relatedly, statements with greater complexity, less similarity, and the presence of dissents are associated with longer event windows on average.

The implications of these findings are not merely methodological, as the choice of event window has a tangible impact on the construction and interpretation of monetary policy surprises and shocks. I document that the correlation between surprises measured within the optimal windows and those from a conventional 30-minute window decreases with asset underlying maturity. These discrepancies are economically meaningful, resulting in observed differences in derived monetary policy shocks, particularly the forward guidance component. Furthermore, the responses of interest rates and break-even inflation to monetary policy shocks about forward guidance become larger when opening up the event window to the optimal length. This larger estimated impact suggests that shorter, conventional windows attenuate the measured effects of policy. Similar results are found in the response of stock prices to these shock series. With regards to the broader macroeconomy, I document that monetary policy shocks constructed within the optimal event window length result in the estimated responses of macroeconomic variables to become more precise.

Ultimately, this paper argues that determining the optimal event window length is an empirical question, not a parameter whose value should be simply assumed. The provided methodology offers a systematic approach for this basic but critical step in event window studies. The results indicate that by allowing for longer, asset-specific reaction times, we can construct monetary policy surprises that more accurately reflect the market's full response to central bank communication, leading to a more precise understanding of the effects of monetary policy. Indeed, as the literature considers broader measures of central bank communication in order to construct more holistic monetary policy shocks (e.g., Neuhierl and Weber, 2019; Swanson and Jayawickrema, 2023; Bauer and Swanson, 2023; and others), I argue that when the primary source of reaction is communicated through text, the methodology of this paper can be used to measure when markets fully react to these policy communications, resulting in more relevant and precise measurements of unanticipated monetary policy.

References

Acosta, Miguel (2023). "The Perceived Causes of Monetary Policy Surprises". Published manuscript.

Acosta, Miguel, Connor M. Brennan, and Margaret M. Jacobson (2024). *Constructing high-frequency monetary policy surprises from SOFR futures*. Tech. rep. Finance and Economics Discussion Series 2024-034. Board of Governors of the Federal Reserve System.

Acosta, Miguel and Ellen E. Meade (2015). *Hanging on Every Word: Semantic Analysis of the FOMC's Postmeeting Statement*. Tech. rep. FEDS Notes. Board of Governors of the Federal Reserve System.

An, Phillip, Karlye Dilts Stedman, and Amaze Lusompa (2025). *How High Does High Frequency Need to Be? A Comparison of Daily and Intradaily Monetary Policy Surprises*. Tech. rep. Research Working Paper no. 25-03. Federal Reserve Bank of Kansas City.

Athey, Susan and Guido Imbens (2019). *Machine Learning Methods Economists Should Know About*.

Babur, Önder and Loek Cleophas (2017). "Using n-grams for the Automated Clustering of Structural Models". In: *SOFSEM 2017: Theory and Practice of Computer Science*. Ed. by Bernhard Steffen, Christel Baier, Mark van den Brand, Johann Eder, Mike Hinchey, and Tiziana Margaria. Cham: Springer International Publishing, pp. 510–524.

Bauer, Michael D. and Eric T. Swanson (2023). "A Reassessment of Monetary Policy Surprises and High-Frequency Identification". In: *NBER Macroeconomics Annual* 37.1, pp. 87–155.

Boehm, Christoph E and T Niklas Kroner (2025). "Monetary Policy without Moving Interest Rates: The Fed Non-Yield Shock". SSRN Working Paper No 3812524.

Bordalo, Pedro, Nicola Gennaioli, Rafael La Porta, and Andrei Shleifer (2024). "Belief Overreaction and Stock Market Puzzles". In: *Journal of Political Economy* 132.5, pp. 1450–1484.

Brennan, Connor M., Margaret M. Jacobson, Christian Matthes, and Todd B. Walker (2024). *Monetary Policy Shocks: Data or Methods?* Tech. rep. Finance and Economics Discussion Series 2024-011r1. Board of Governors of the Federal Reserve System.

Brooks, Jordan, Michael Katz, and Hanno Lustig (2023). *Post-FOMC Announcement Drifts in U.S. Bond Markets*. Working Paper 25127. National Bureau of Economic Research.

Casini, Alessandro and Adam McCloskey (2025). *Identification and Estimation of Causal Effects in High-Frequency Event Studies*. arXiv: 2406.15667 [econ.EM].

Chang, Son J and Son-Nan Chen (1989). "Stock-price adjustment to earnings and dividend surprises". In: *Quarterly Review of Economics and Business* 29.1, pp. 68–81.

Crump, Richard K., Stefano Eusepi, and Emanuel Moench (2013). *Making a Statement: How Did Professional Forecasters React to the August 2011 FOMC Statement?* Liberty Street Economics 20130107. Federal Reserve Bank of New York.

Das, Somnath and Alexander Z. King (2021). "Measuring the informativeness of earnings announcements: The role of event windows". In: *The Quarterly Review of Economics and Finance* 82, pp. 350–367.

Doh, Taeyoung, Dongho Song, and Shu-Kuei Yang (2023). "Deciphering Federal Reserve Communication via Text Analysis of Alternative FOMC Statements". Federal Reserve Bank of Kansa City Working Paper.

Duffee, Gregory (2013). "Forecasting Interest Rates". In: *Handbook of Economic Forecasting*. Ed. by Graham Elliott and Allan Timmermann. Vol. 2. Handbook of Economic Forecasting. Elsevier. Chap. 7, pp. 385–426.

Fleming, Michael J. and Monika Piazzesi (2005). *Monetary Policy Tick-by-Tick*. Tech. rep. Working Paper. Federal Reserve Bank of New York.

Gentzkow, Matthew, Bryan Kelly, and Matt Taddy (2019). "Text as Data". In: *Journal of Economic Literature* 57.3, pp. 535–74.

Gertler, Mark and Peter Karadi (2015). "Monetary Policy Surprises, Credit Costs, and Economic Activity". In: *American Economic Journal: Macroeconomics* 1, pp. 44–76.

Gilchrist, Simon and Egon Zakrajšek (2012). "Credit Spreads and Business Cycle Fluctuations". In: *American Economic Review* 102.4, pp. 1692–1720.

Glasserman, Paul and Caden Lin (2023). *Assessing Look-Ahead Bias in Stock Return Predictions Generated By GPT Sentiment Analysis*.

Golez, Benjamin, Peter Kelly, and Ben Matthies (2025). "FOMC News and Segmented Markets". In: *Journal of Accounting and Economics* 79.2, p. 101767.

Gorodnichenko, Yuriy and Walker Ray (2017). *The Effects of Quantitative Easing: Taking a Cue from Treasury Auctions*. Working Paper 24122. National Bureau of Economic Research.

Graham, Benjamin (1974). *The Renaissance of Value: The Proceedings of a Seminar on the Economy, Interest Rates, Portfolio Management, and Bonds Vs Common Stocks, September 18, 1974*. Financial Analysts Research Foundation.

Gürkaynak, Refet S., Burçin Kisacikoğlu, and Jonathan H. Wright (2020). "Missing Events in Event Studies: Identifying the Effects of Partially Measured News Surprises". In: *American Economic Review* 110.12, pp. 3871–3912.

Gürkaynak, Refet S., Brian Sack, and Eric T. Swanson (2005). "Do Actions Speak Louder Than Words? The Response of Asset Prices to Monetary Policy Actions and Statements". In: *International Journal of Central Banking* 1.1.

Gürkaynak, Refet S., Brian Sack, and Jonathan H. Wright (2010). "The TIPS Yield Curve and Inflation Compensation". In: *American Economic Journal: Macroeconomics* 2.1, pp. 70–92.

Handlan, Amy (2022a). "FedSpeak Matters: Statement Similarity and Monetary Policy Expectations". Published manuscript.

Handlan, Amy (2022b). "Text Shocks and Monetary Surprises: Text Analysis of FOMC Statements with Machine Learning". Published manuscript.

Hawinkel, Stijn, Willem Waegeman, and Steven Maere (2024). "Out-of-Sample R2: Estimation and Inference". In: *The American Statistician* 78.1, pp. 15–25.

Hillmer, S.C. and P.L. Yu (1979). "The market speed of adjustment to new information". In: *Journal of Financial Economics* 7.4, pp. 321–345.

Hornik, Kurt, Maxwell Stinchcombe, and Halbert White (1989). "Multilayer feedforward networks are universal approximators". In: *Neural Networks* 2.5, pp. 359–366.

Husted, Lucas, John Rogers, and Bo Sun (2020). "Monetary policy uncertainty". In: *Journal of Monetary Economics* 115, pp. 20–36.

Indriawan, Ivan, Feng Jiao, and Yiuman Tse (2021). "The FOMC announcement returns on long-term US and German bond futures". In: *Journal of Banking & Finance* 123, p. 106027.

Jarociński, Marek and Peter Karadi (2020). "Deconstructing Monetary Policy Surprises—The Role of Information Shocks". In: *American Economic Journal: Macroeconomics* 12.2, pp. 1–43.

Kuttner, Kenneth N. (2001). "Monetary Policy Surprises and Interest Rates: Evidence from the Fed Funds Futures Market". In: *Journal of Monetary Economics* 47.3, pp. 523–544.

Madeira, Carlos and João Madeira (2019). "The Effect of FOMC Votes on Financial Markets". In: *The Review of Economics and Statistics* 101.5, pp. 921–932.

McMahon Naylor, Matthew (2023). *Getting Through: Communicating Complex Information.* Tech. rep. Staff Working Paper No. 1047. Bank of England.

Nakamura, Emi and Jón Steinsson (2018). "High-Frequency Identification of Monetary Non-Neutrality: The Information Effect". In: *The Quarterly Journal of Economics* 133.3, pp. 1283–1330.

Neuhierl, Andreas and Michael Weber (2019). "Monetary policy communication, policy slope, and the stock market". In: *Journal of Monetary Economics* 108. "Central Bank Communications:From Mystery to Transparency"May 23-24, 2019Annual Research Conference ofthe National Bank of UkraineOrganized in cooperation withNarodowy Bank Polski, pp. 140–155.

Neuhierl, Andreas and Michael Weber (2024). *Monetary Momentum.* Chicago Booth Research Paper, Fama-Miller Working Paper, University of Chicago, Becker Friedman Institute for Economics Working Paper, LawFin Working Paper 20-05, 2020-39, 37. Chicago Booth Research Paper Series, Chicago Booth: Fama-Miller Working Paper Series, Becker

Friedman Institute for Economics Working Paper Series, DFG Center for Advanced Studies - Foundations of Law and Finance.

Newey, Whitney K. and Kenneth D. West (1987). "A Simple, Positive Semi-Definite, Heteroskedasticity and Autocorrelation Consistent Covariance Matrix". In: *Econometrica* 55.3, pp. 703–708.

Okada, Masayuki (2025). "Interest Rate Risk around Monetary Policy Announcements and Asset Duration". Working paper.

Olea, José Luis Montiel and Mikkel Plagborg-Møller (2021). "Local Projection Inference is Simpler and More Robust Than You Think". In: *Econometrica* 89.4, pp. 1789–1823. eprint: https://onlinelibrary.wiley.com/doi/pdf/10.3982/ECTA18756.

Piazzesi, Monika and Eric T. Swanson (2008). "Futures prices as risk-adjusted forecasts of monetary policy". In: *Journal of Monetary Economics* 55.4, pp. 677–691.

Piller, Alexander, Marc Schranz, and Larissa Schwaller (2025). "Using Natural Language Processing to Identify Monetary Policy Shocks". Working Paper.

Rajapakse, Thilina C., Andrew Yates, and Maarten de Rijke (2024). "Simple Transformers: Open-source for All". In: *Proceedings of the 2024 Annual International ACM SIGIR Conference on Research and Development in Information Retrieval in the Asia Pacific Region*. SIGIR-AP 2024. Tokyo, Japan, pp. 209–215.

Sarkar, Suproteem and Keyon Vafa (2024). "Lookahead Bias in Pretrained Language Models". SSRN Working Paper No 4754678.

Smales, L.A. and N. Apergis (2017). "Does more complex language in FOMC decisions impact financial markets?" In: *Journal of International Financial Markets, Institutions and Money* 51, pp. 171–189.

Swanson, Eric T. and Vishuddhi Jayawickrema (2023). "Speeches by the Fed Chair Are More Important than FOMC Announcements: An Improved High-Frequency Measure of U.S. Monetary Policy Shocks". Unpublished manuscript.

Tsang, Kwok Ping and Zichao Yang (2024). *Agree to Disagree: Measuring Hidden Dissent in FOMC Meetings*. arXiv: 2308.10131 [econ.GN].

Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin (2017). "Attention Is All You Need". In: *CoRR* abs/1706.03762.

Yang, Zhilin, Zihang Dai, Yiming Yang, Jaime G. Carbonell, Ruslan Salakhutdinov, and Quoc V. Le (2019). "XLNet: Generalized Autoregressive Pretraining for Language Understanding". In: *CoRR* abs/1906.08237.

Zhang, Yu, Konstantina Kappou, and Andrew Urquhart (n.d.). "Post-FOMC Drift in the Equity Options Market". SSRN Working Paper No 5464630.

Figures and Tables

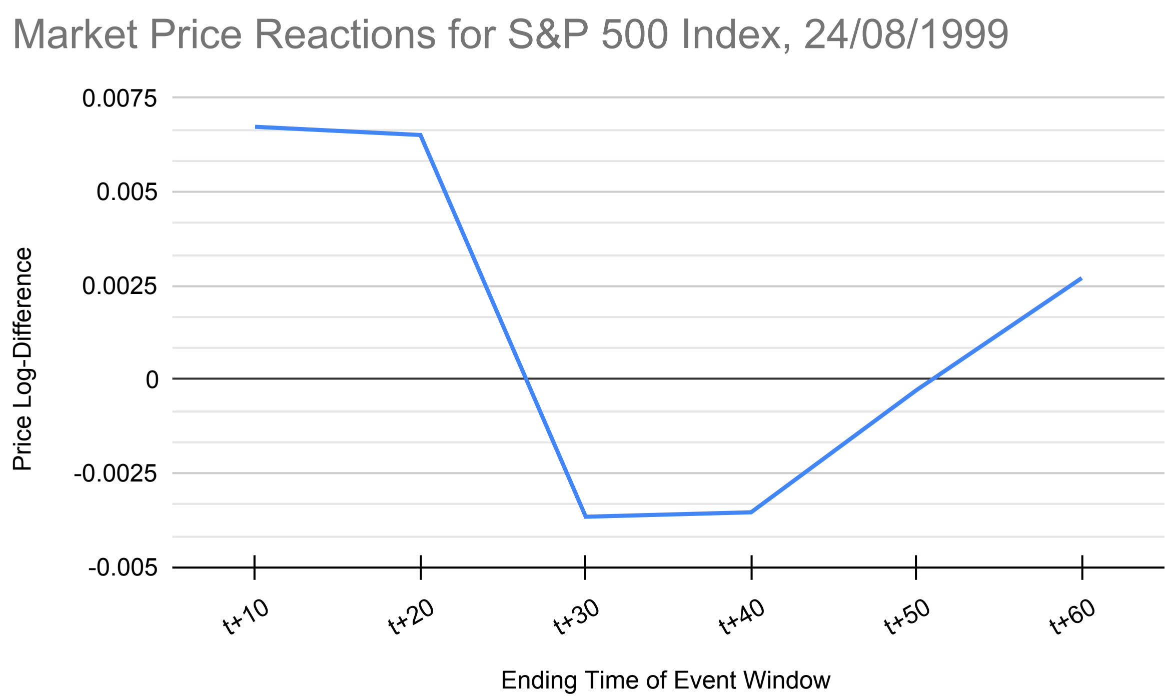

Figure 1: Responses of S&P 500 Index to FOMC Statements

Notes: The price log-differences of the S&P 500 Index for different event window lengths after the release of the 24 August, 1999 FOMC statement is depicted. The horizontal axis of each figure depicts the end time of the considered event window lengths and the vertical axis represents the price-log differences.

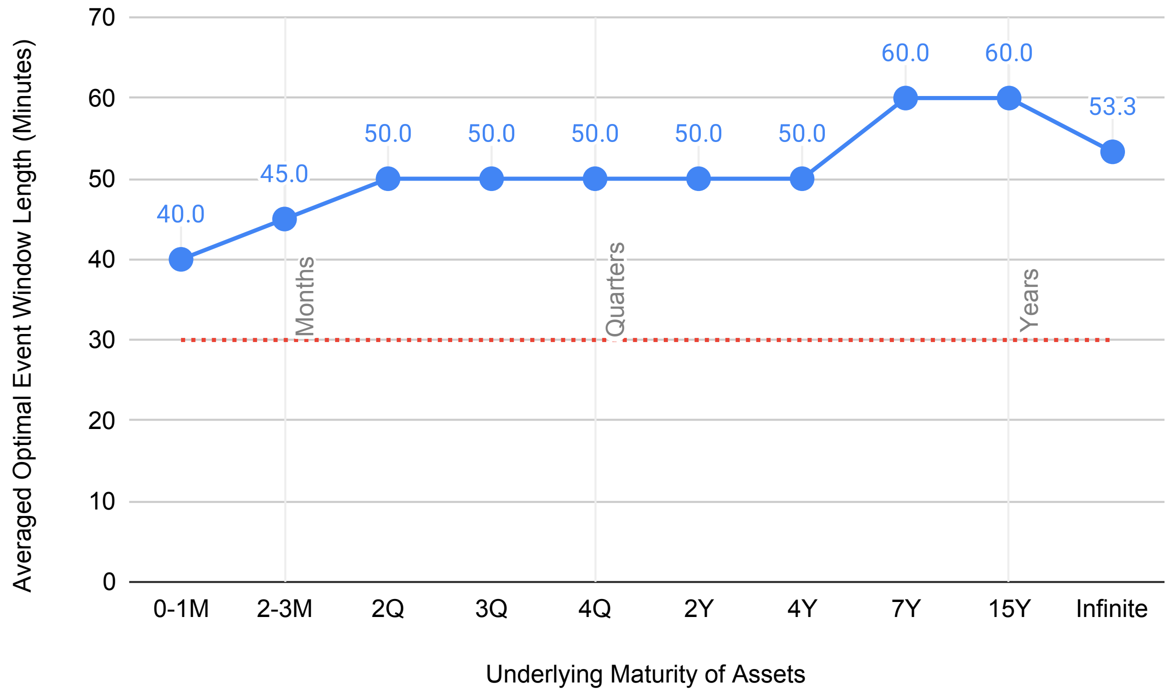

Figure 2: Averaged Optimal Event Window Lengths for Various Asset Underlying Maturities

Notes: The horizontal axis depicts the horizons of all considered futures contracts and equities. The vertical axis depicts the event window lengths in minutes, with all windows starting at 10 minutes before FOMC statement release. The red dotted line depicts the 30-minute window length, common in the literature. The systematically estimated event window lengths are averaged across futures maturities for each asset type. "0–1M" is the average of the event window lengths for front-month and 1-month-ahead Federal Funds futures. "2–3M" is the average of the event window lengths for 2-month and 3-month-ahead Federal Funds futures. "Infinite" is the underlying maturity of the S&P 500 Index and considered E-mini futures contracts. The underlying maturities for the Treasury futures contracts are approximated by Gürkaynak, Kisacikoğlu, and Wright (2020).

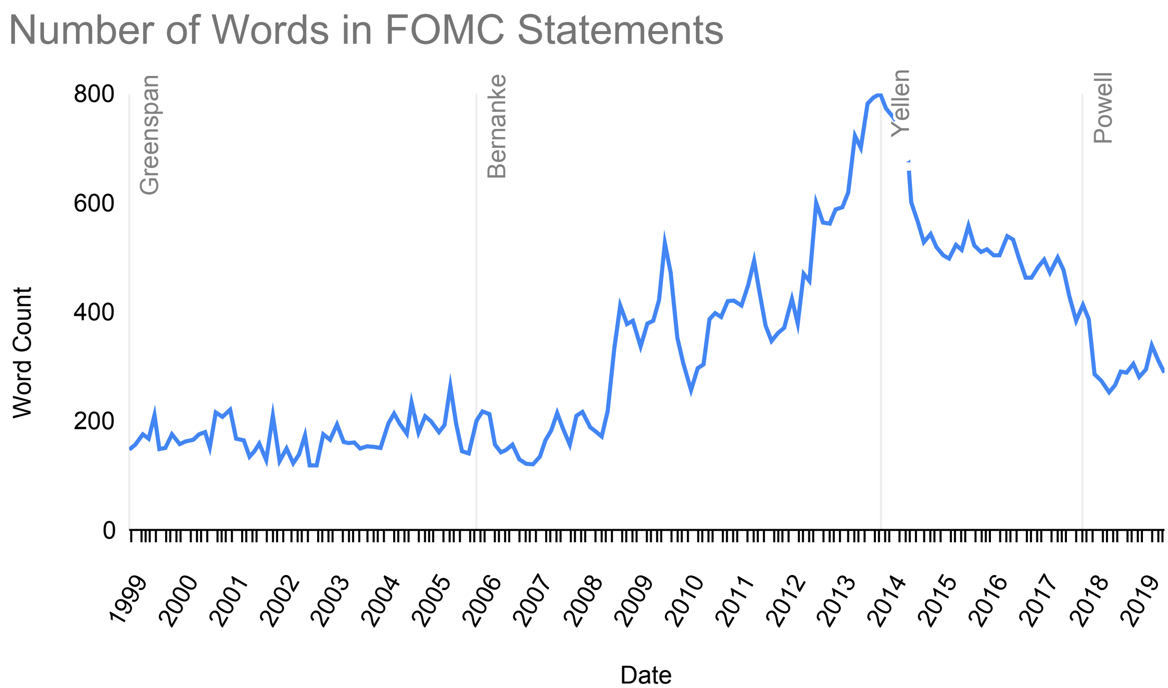

Figure 3: Number of Words in FOMC Statements, May 1999–October 2019

Notes: The above counts are for FOMC statements that have underwent pre-processing, which is explained in Subsection 3.1. From left to right, the vertical grey lines indicate the first FOMC meeting with Greenspan, Bernanke, Yellen, and Powell as Fed Chair.

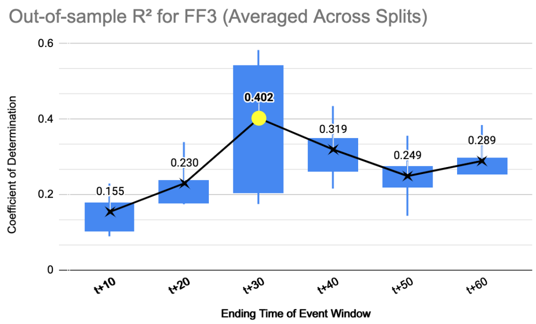

(a) $FF3$

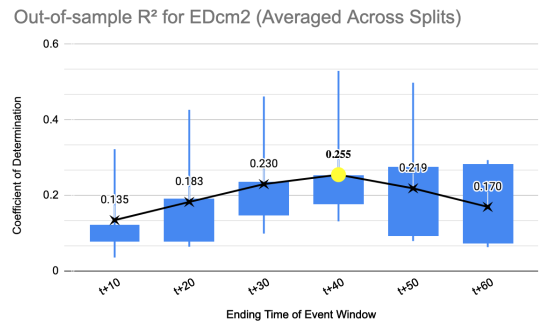

(b) $EDcm2$

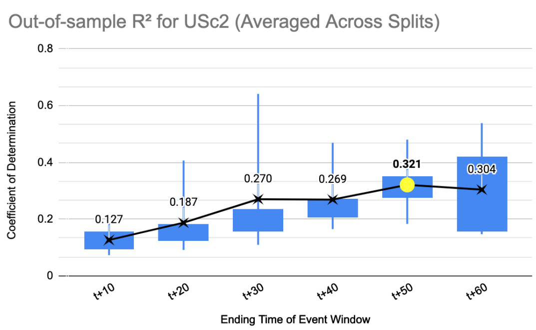

(c) $USc2$

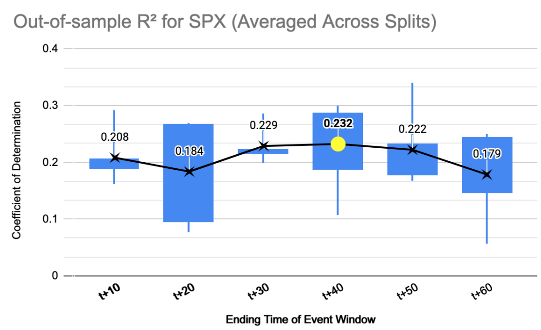

(d) SPX

Figure 4: Optimal Event Window Lengths for Interest-Rate Futures and S&P 500 Index
Notes: The horizontal axis of each figure depicts the end time of the considered event window lengths and the vertical axis represents the percentage points of price changes by the market directly and only due to the FOMC statement text. The cross points represent $\overline{R^2_{OOS}}$ for each event window size, where the solid yellow point represents the event window with the largest $\overline{R^2_{OOS}}$. For each event window, box-and-whisker plots are shown surrounding the corresponding averages. $0 \leq \overline{R^2_{OOS}} \leq 1$ represents the proportion of the mean squared error from predicting with the in-sample average explained by the superior out-of-sample performance of XLNet-Base.

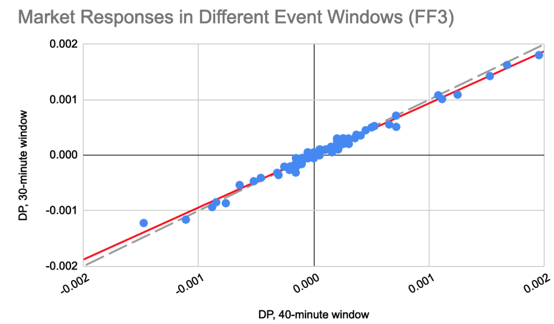

(a) $FF3^{**}$

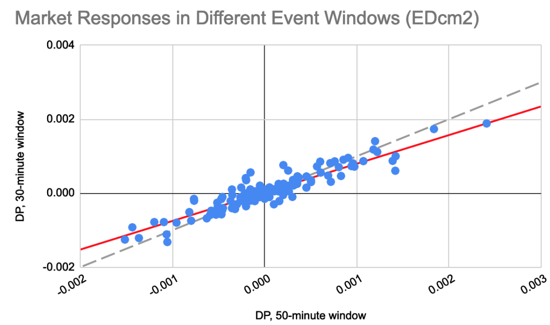

(b) $EDcm2^{***}$

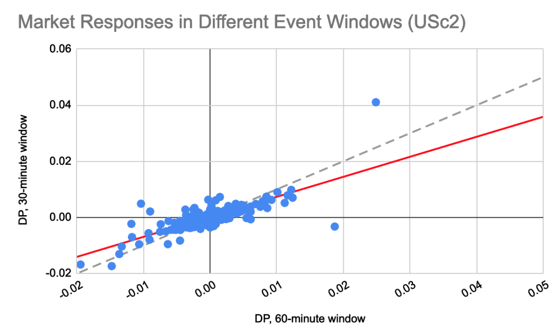

(c) $USc2^{***}$

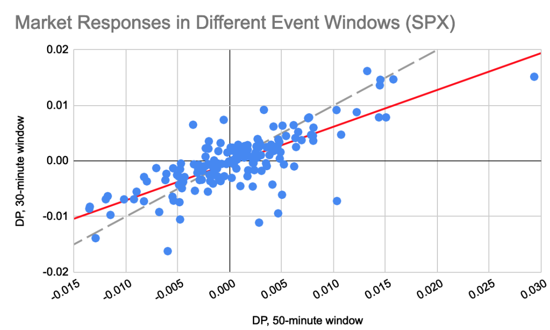

(d) SPX^{***}

Figure 5: Comparing Market Responses in Different Event Windows for Interest-Rate Futures and S&P 500 Index
Notes: The horizontal axis depicts the price log-difference within the systematically estimated event window length. The vertical axis represents DP_{t+20}. The 45-degree line is depicted as grey and dashed. The blue dots are market reactions on scheduled FOMC meetings. The price-log differences calculated within the optimal window lengths are regressed on DP_{t+20} through OLS. If the slope coefficient is greater (less) than one and its difference with one is statistically significant for at most $\alpha = 0.10$, then the financial market under-reacts to the information content of FOMC statements on release, ex post, and is depicted in red. ** sig. at the 5% level, *** sig. at the 1% level.

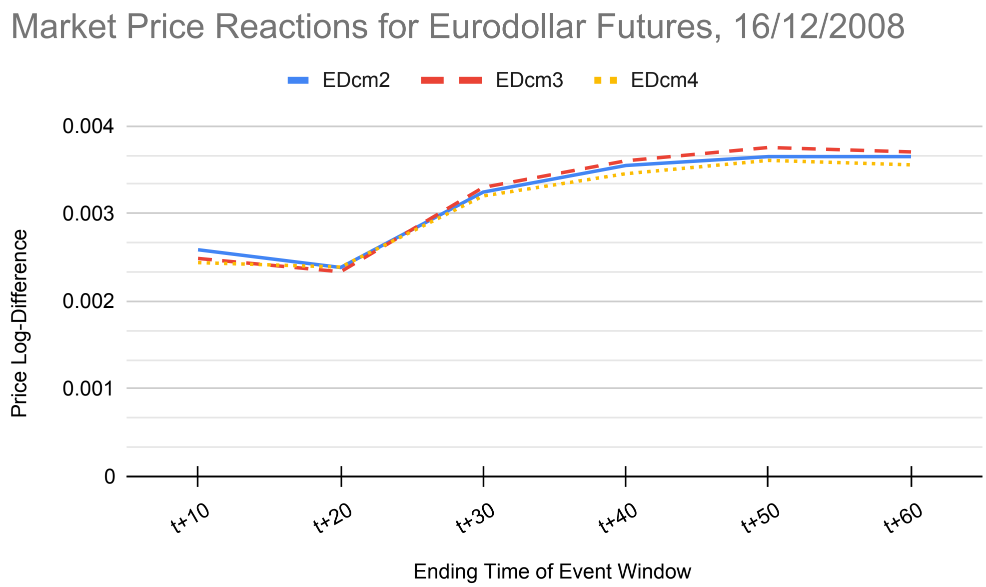

(a) Eurodollar Futures

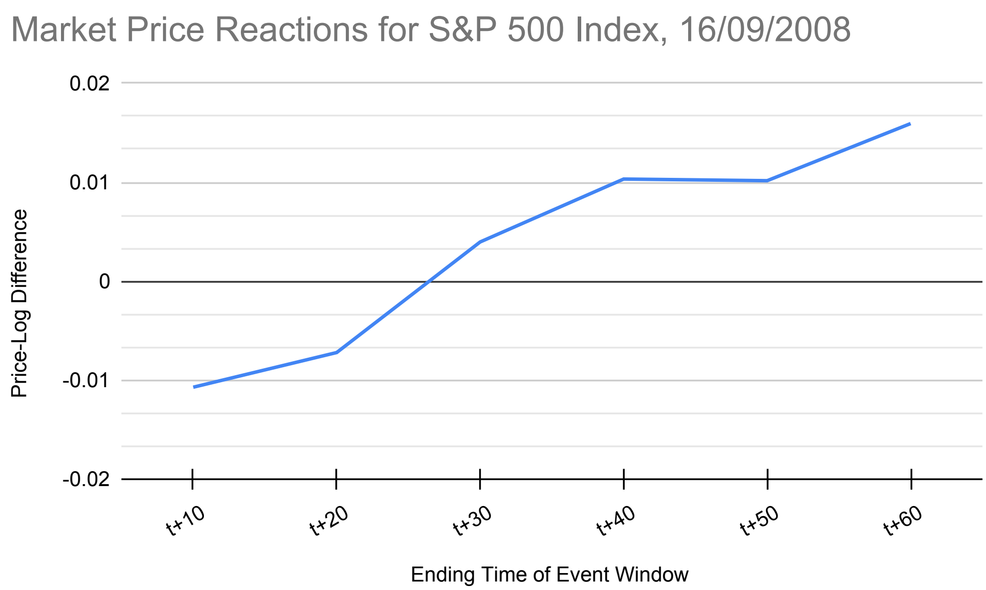

(b) SPX

Figure 6: Responses of Eurodollar Futures and S&P 500 Index to FOMC Statements
Notes: The top (bottom) sub-figure depicts the price log-differences of Eurodollar futures markets (the S&P 500 Index) for different event window lengths after the release of the 16 December (September), 2008 FOMC statement. The horizontal axis of each figure depicts the end time of the considered event window lengths and the vertical axis represents the price-log differences.

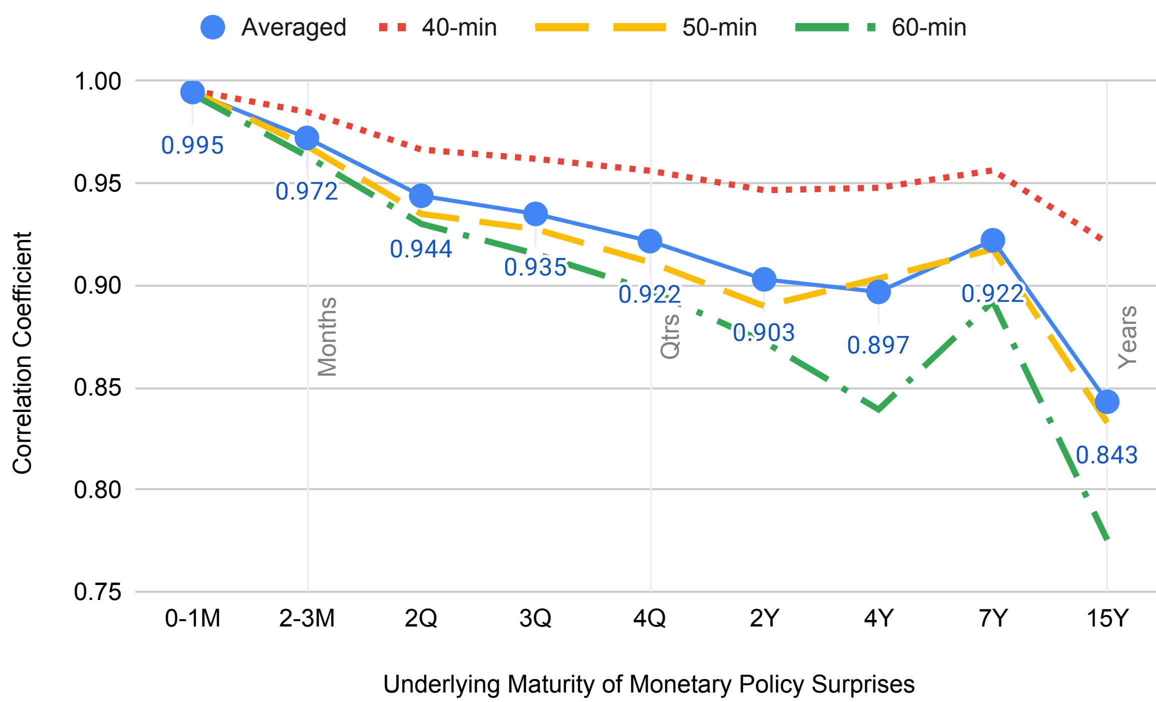

Figure 7: Correlation Between Monetary Policy Surprises Calculated in Optimal v. 30-minute Window Lengths

Notes: The horizontal axis depicts the underlying maturities of interest-rate surprises. The vertical axis represents the Pearson correlation coefficient between the surprises calculated within optimal event window lengths v. 30 minutes. The red-dotted, yellow-dashed, and green-mixed lines represent the correlations calculated for surprises measured within the optimally-found 40, 50, and 60 minutes, respectively. The blue-solid line represents the correlations averaged across the three optimal window lengths. "0–1M" is interest rate surprise calculated using information from front-month and 1-month-ahead Federal Funds futures. "2–3M" represent expectations calculated from 2-month and 3-month-ahead Federal Funds futures. The underlying maturities for the Treasury futures contracts and resulting surprises are approximated by Gürkaynak, Kisacikoğlu, and Wright (2020).

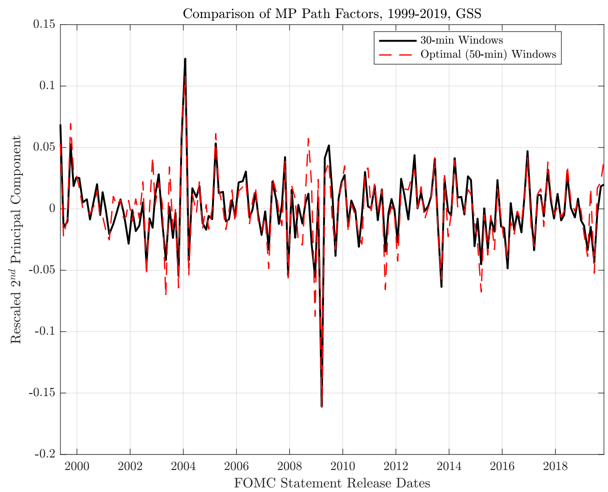

(a) GSS_P

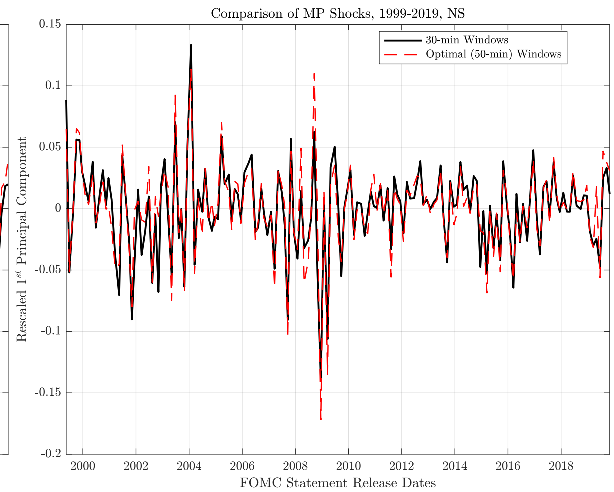

(b) NS_{MP}

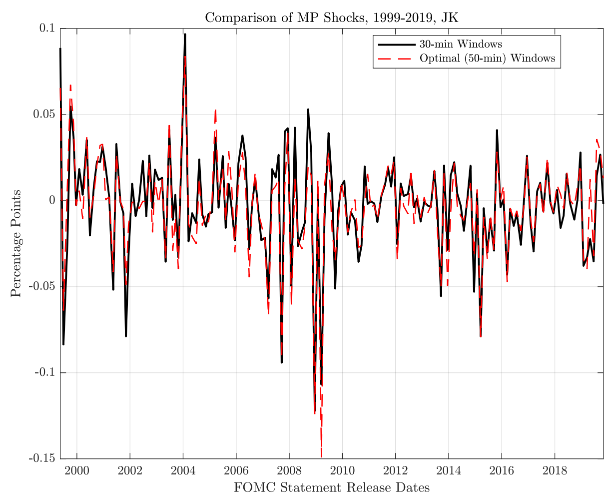

(c) JK_{MP}

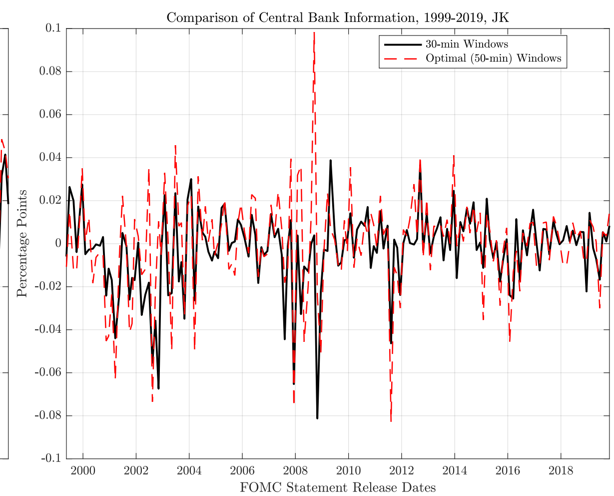

(d) JK_{CBI}

Figure 8: Comparing Monetary Policy Shock Series Derived from Optimal Window Length v. 30 Minutes

Notes: For each sub-figure, the horizontal axis represents FOMC statement release dates. The vertical axis depicts percentage points after rescaling each shock to be one-for-one with the daily change in the zero-coupon, nominal one-year Treasury yield. For all construction methods, the black-solid and red-dotted lines represent the shocks derived from surprises measured within 30 minutes and the median optimal event window length of 50 minutes, respectively.

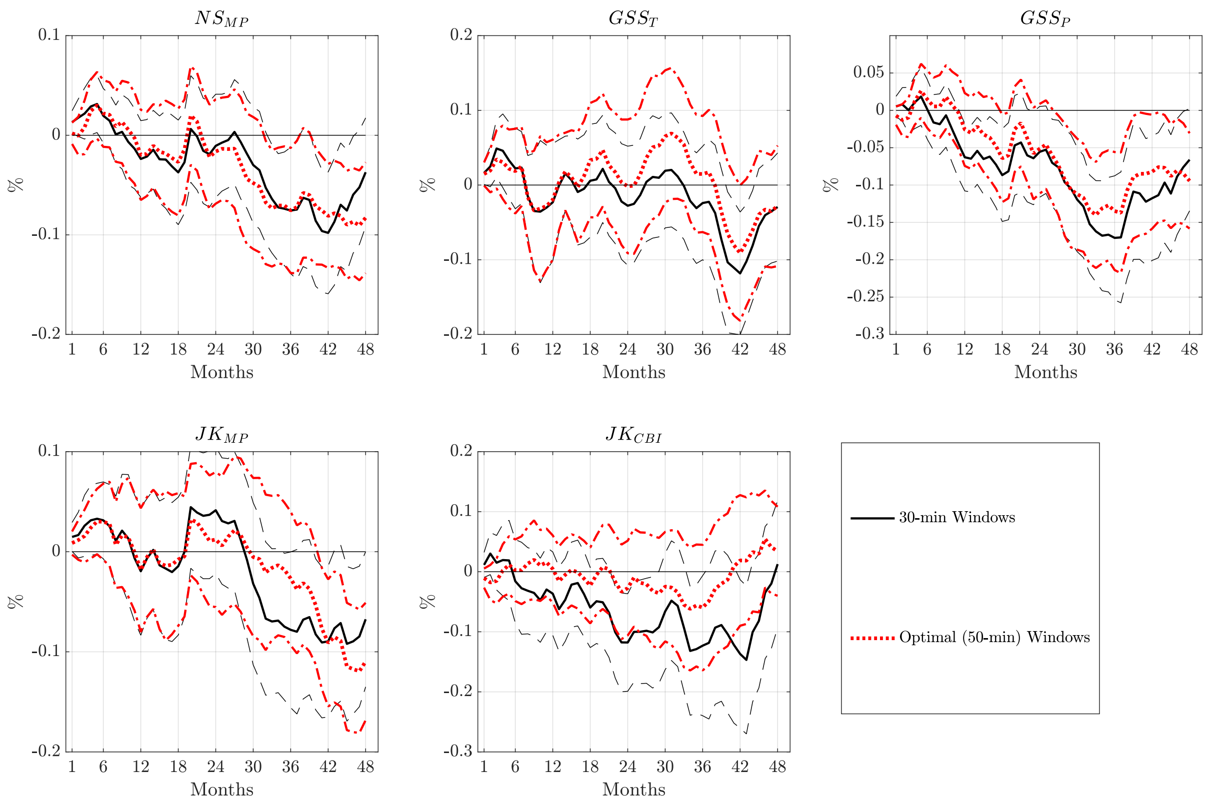

Figure 9: Effects of Event Window Choice on CPI Impulse Responses

Notes: Impulse responses are calculated using the lag-augmented local projection from Olea and Plagborg-Møller (2021). Confidence bands are at the 90% level. Standard errors are Eicker-Huber-White heteroscedasticity-robust standard errors. The responses are to a 100 basis point increase in the monetary policy shock series. The monetary policy shocks are constructed using the full set of monetary policy surprise instruments.

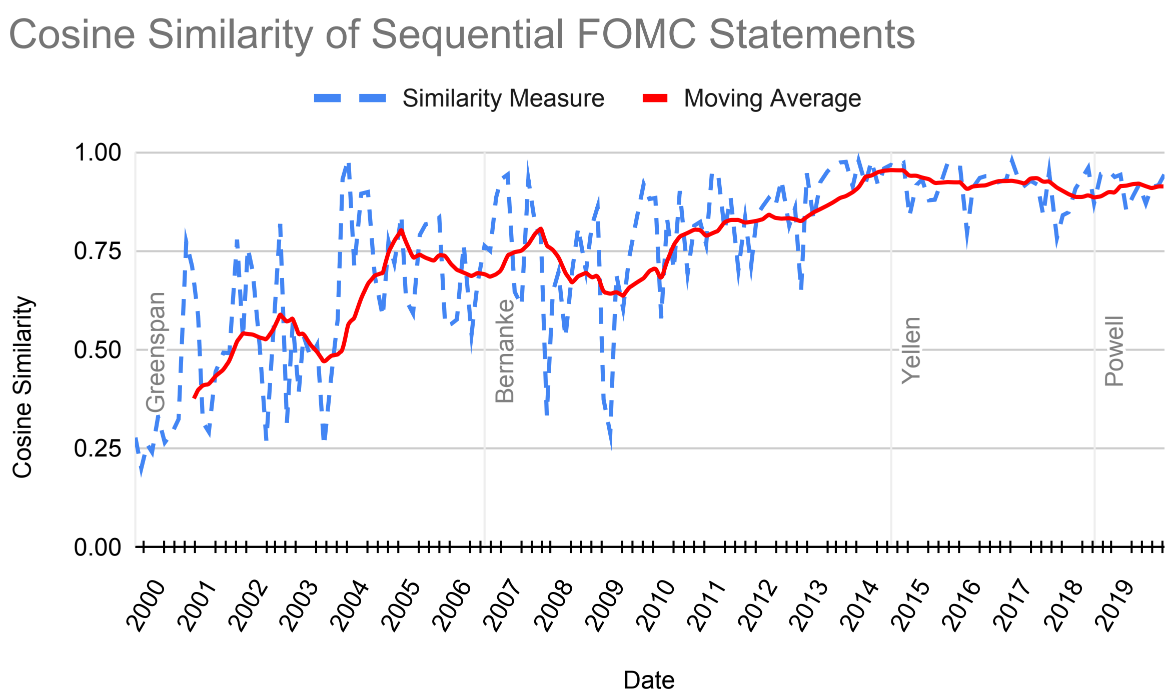

Figure 10: S^1 Cosine Similarity of Sequential FOMC Statements

Notes: S^1 is the cosine similarity between the TFIDF value of an FOMC statement and that of the immediately previous FOMC statement. Cosine similarity values closer to one (zero) mean the statements share more (less) common term usage. The moving average in solid red is calculated with a period of 10.

	30-minute Window				Optimal Window				**Difference**		
	ΔTY_1	ΔTY_2	ΔTY_5	ΔTY_{10}	ΔTY_1	ΔTY_2	ΔTY_5	ΔTY_{10}	ΔTY_2	ΔTY_5	ΔTY_{10}
GSS_T	1.00***	0.82***	0.15	-0.37	1.00***	0.78***	0.08	-0.42	**-0.04**	**-0.07**	**-0.05**
	(0.28)	(0.34)	(0.43)	(0.41)	(0.23)	(0.28)	(0.33)	(0.32)			
GSS_P	1.00***	1.46***	1.89***	1.64***	1.00***	1.51***	1.93***	1.66***	+**0.05**	+**0.04**	+**0.02**
	(0.09)	(0.09)	(0.25)	(0.36)	(0.11)	(0.10)	(0.21)	(0.30)			
NS_{MP}	1.00***	1.24***	1.29***	0.95***	1.00***	1.30***	1.39***	1.06***	+**0.06**	+**0.11**	+**0.11**
	(0.07)	(0.09)	(0.19)	(0.25)	(0.09)	(0.10)	(0.18)	(0.24)			
JK_{MP}	1.00***	1.30***	1.39***	0.99***	1.00***	1.35***	1.52***	1.16***	+**0.04**	+**0.13**	+**0.17**
	(0.11)	(0.15)	(0.28)	(0.36)	(0.12)	(0.16)	(0.32)	(0.43)			
JK_{CBI}	1.00***	1.04***	1.00***	0.82***	1.00***	1.20***	1.14***	0.85***	+**0.16**	+**0.14**	+**0.03**
	(0.25)	(0.30)	(0.31)	(0.29)	(0.23)	(0.25)	(0.27)	(0.26)			

Table 1: Differences in Responses of Nominal Interest Rates to Shocks from Event Window Choice

Notes: Each estimate comes from a separate OLS regression. For each regression, the dependent variable is the one-day change in end-of-day Treasury yields, represented by columns. The independent variable is the change in monetary policy shock series over a given event window around the time of scheduled FOMC announcements. Starting from the left, columns 2–5 represent the OLS regressions performed on shock series constructed within 30-minute windows. Columns 6–9 are equivalent, but for shocks constructed within the median optimal window length of 50 minutes. The sample period consists of all scheduled FOMC statement release dates from May 1999 through October 2019, resulting in 165 observations. Newey and West (1987) standard errors are reported in parentheses. Columns 10–12 represent the differences between the coefficients of the regressions using the optimal windows v. those using 30-minute windows. *** sig. at the 1% level.

	30-minute Window			Optimal Window			**Difference**		
	$\Delta TIPS_2$	$\Delta TIPS_5$	$\Delta TIPS_{10}$	$\Delta TIPS_2$	$\Delta TIPS_5$	$\Delta TIPS_{10}$	$\Delta TIPS_2$	$\Delta TIPS_5$	$\Delta TIPS_{10}$
GSS_T	-0.81	0.02	-0.19	-0.90	0.09	-0.16	**-0.09**	**+0.07**	**+0.03**
	(1.44)	(0.51)	(0.45)	(1.57)	(0.46)	(0.37)			
GSS_P	2.21***	1.96***	1.74***	2.20***	2.03***	1.75***	**-0.00**	**+0.06**	**+0.01**
	(0.49)	(0.46)	(0.44)	(0.37)	(0.38)	(0.37)			
NS_{MP}	1.17***	1.29***	1.08***	1.31***	1.47***	1.20***	**+0.14**	**+0.18**	**+0.13**
	(0.73)	(0.36)	(0.30)	(0.63)	(0.31)	(0.28)			
JK_{MP}	1.40***	1.40***	1.15***	1.66***	1.64***	1.38***	**+0.26**	**+0.24**	**+0.23**
	(0.83)	(0.47)	(0.41)	(0.63)	(0.49)	(0.46)			
JK_{CBI}	0.51	0.99***	0.85***	0.60	1.13***	0.84***	**+0.09**	**+0.14**	**-0.01**
	(0.85)	(0.33)	(0.25)	(0.92)	(0.33)	(0.25)			

Table 2: Differences in Responses of Real Interest Rates to Shocks from Event Window Choice

Notes: Each estimate comes from a separate OLS regression. For each regression, the dependent variable is the one-day change in end-of-day TIPS yields, represented by columns. The independent variable is the change in monetary policy shock series over a given event window around the time of scheduled FOMC announcements. Starting from the left, columns 2–4 represent the OLS regressions performed on shock series constructed within 30-minute windows. Columns 5–7 are equivalent, but for shocks constructed within the median optimal window length of 50 minutes. The sample period consists of all scheduled FOMC statement release dates from May 1999 to October 2019, resulting in 165 observations. Newey and West (1987) standard errors are reported in parentheses. Columns 8–10 represent the differences between the coefficients of the regressions using the optimal windows v. those using 30-minute windows. *** sig. at the 1% level.

	30-minute Window			Optimal Window			**Difference**		
	ΔBEI_2	ΔBEI_5	ΔBEI_{10}	ΔBEI_2	ΔBEI_5	ΔBEI_{10}	ΔBEI_2	ΔBEI_5	ΔBEI_{10}
GSS_T	1.63***	0.13	-0.18	1.67***	-0.01	−0.26**	**+0.05**	**-0.14**	**-0.08**
	(1.46)	(0.29)	(0.17)	(1.58)	(0.29)	(0.17)			
GSS_P	-0.75	-0.08	-0.10	−0.69*	-0.10	-0.09	**+0.06**	**-0.02**	**+0.01**
	(0.46)	(0.24)	(0.12)	(0.36)	(0.23)	(0.12)			
NS_{MP}	0.07	-0.01	−0.13*	-0.01	-0.07	−0.14**	**-0.08**	**-0.07**	**-0.01**
	(0.70)	(0.23)	(0.12)	(0.65)	(0.23)	(0.12)			
JK_{MP}	-0.09	-0.01	−0.17**	-0.31	-0.12	−0.22***	**-0.22**	**-0.11**	**-0.05**
	(0.81)	(0.29)	(0.13)	(0.61)	(0.25)	(0.11)			
JK_{CBI}	0.54	0.01	-0.02	0.60	0.01	0.02	**+0.07**	**+0.00**	**+0.04**
	(0.76)	(0.29)	(0.20)	(0.86)	(0.30)	(0.23)			

Table 3: Differences in Responses of Break-even Inflation to Shocks from Event Window Choice

Notes: Each estimate comes from a separate OLS regression. For each regression, the dependent variable is the one-day change in end-of-day break-even inflation, represented by columns. The independent variable is the change in monetary policy shock series over a given event window around the time of scheduled FOMC announcements. Starting from the left, columns 2–4 represent the OLS regressions performed on shock series constructed within 30-minute windows. Columns 5–7 are equivalent, but for shocks constructed within the median optimal window length of 50 minutes. The sample period consists of all scheduled FOMC statement release dates from May 1999 to October 2019, resulting in 165 observations. Newey and West (1987) standard errors are reported in parentheses. Columns 8–10 represent the differences between the coefficients of the regressions using the optimal windows v. those using 30-minute windows. * sig. at the 10% level. ** sig. at the 5% level. *** sig. at the 1% level.

	30-minute Window			Optimal Window			**Percentage Difference**		
	DP_{SPX}	DP_{ESc1}	DP_{ESc2}	DP_{SPX}	DP_{ESc1}	DP_{ESc2}	DP_{SPX}	DP_{ESc1}	DP_{ESc2}
GSS_T	−8.40***	−8.83***	−7.25***	−7.39***	−7.43***	−7.34***	**-11.99%**	**-15.92%**	**+1.25%**
	(2.71)	(2.68)	(2.78)	(3.10)	(3.15)	(3.12)			
GSS_P	−6.14***	−6.27***	−6.12***	−6.85***	−6.96***	−7.63***	**+11.51%**	**+11.00%**	**+24.61%**
	(1.81)	(1.83)	(1.76)	(2.88)	(2.91)	(2.81)			
NS_{MP}	−6.92***	−7.15***	−6.51***	−7.00***	−7.10***	−7.55***	**+1.23%**	**-1.00%**	**+16.00%**
	(1.32)	(1.37)	(1.31)	(1.85)	(1.89)	(1.84)			
JK_{MP}	−14.76***	−15.08***	−13.73***	−17.46***	−17.77***	−17.30***	**+18.25%**	**+17.88%**	**+26.00%**
	(0.81)	(0.91)	(0.94)	(1.04)	(1.08)	(1.06)			
JK_{CBI}	15.19***	13.84***	15.18***	14.08***	14.44***	12.12***	**-7.36%**	**-4.90%**	**-12.43%**
	(2.29)	(2.35)	(2.39)	(2.11)	(2.14)	(2.07)			

Table 4: Differences in Responses of Stock Prices to Shocks from Event Window Choice

Notes: Each estimate comes from a separate OLS regression. For each regression, the dependent variable is the price log-difference of the S&P 500 Index or E-mini futures, represented by columns. The independent variable is the change in monetary policy shock series over a given event window around the time of scheduled FOMC announcements. Starting from the left, columns 2–4 represent the OLS regressions performed on shock series constructed within 30-minute windows. Columns 5–7 are equivalent, but for shocks constructed within the median optimal window length of 50 minutes. The sample period consists of all scheduled FOMC statement release dates from May 1999 to October 2019, resulting in 165 observations. Newey and West (1987) standard errors are reported in parentheses. Columns 8–10 represent the percentage differences between the coefficients of the regressions using the optimal windows v. those using 30-minute windows, where positive (negative) values represent a stronger (weaker) effect in the same direction. *** sig. at the 1% level.

Metric	Simple	**Complicated**	**Different**	Similar	Unity	**Dissents**
Minimised MSE						
Average	1.25e-5	**1.06e-5**	**1.18e-5**	1.13e-5	9.57e-6	**1.43e-5**
Event Window Length (Minutes)						
Average	60	**71**	**62**	51	53	**83**

Table 5: Average MSEs and Event Window Lengths Using the "One Signal" Approach, Conditioned by FOMC Statement Complexity, Similarity, and Presence of Dissents
Notes: The complexity of FOMC statements is measured by the Flesch-Kincaid Grade Level, defined as: $0.39 \times$ average sentence length $+ 11.8 \times$ average number of syllables per word $-$ 15.59, and displayed in the first two columns. Changes in FOMC statements are measured using a pairwise-statement cosine similarity measure and displayed in the third and fourth columns from the left. The event window lengths are displayed in minutes. "Simple" statements have grade levels up to 16.5. "Complicated" statements have grade levels above 16.5. "Different" are sequential statements with a cosine similarity of less than to 0.885. "Similar" are sequential statements with a cosine similarity of more than 0.885. "Unity" statements are those without votes of dissent. "Dissents" are statements with recorded dissent votes. For all futures contracts, event window lengths considered as outliers under the "one signal" approach are set equal to the median of the sub-set window lengths in order to lessen their effects.

A Motivating Framework Simulations

For the motivating framework discussed in Section 2, I simulate the asset price process over time for multiple news announcements. The simulation illustrates the effect of cognitive noise and unrelated news on the optimal time horizon for price responses and demonstrates why a good signal is required for this time horizon estimation, both motivating the necessity of my methodology.

The simulation assumes that news is released at time $t = 0$. The fundamental price component, $P_{i,t}^f$, responds by jumping to and remaining at real value $P_i^f \in [-100, 100]$. Simultaneously, the cognitive noise component jumps to a random value within $[-100, 100]$ while unrelated news is equal to zero at $t = 0$. The noisy signal observed by the econometrician, s_i, is assumed to have its error component distributed normally with mean zero and finite variance. To analysis this on average, I simulate the price process and the noisy signal for $N = 10,000$ news announcements up to time $t = 100$, calculating both the true and estimated MSEs using P_i^f and signal s_i, respectively. Finally, the true and estimated times at which the market fully reacts on average, t^* and $\hat{t}$, can be determined according to Equation 4.

A.1 Simulation Results

Three scenarios are considered for the asset market price responding to the release of news: one that exhibits cognitive noise but little unrelated news, one with unrelated news but little cognitive noise, and one with both cognitive noise and unrelated news. The simulation parameters corresponding to these three scenarios can be found in the topmost eight rows of Appendix Table F1.

Reading the columns of the bottom two rows of Appendix Table F1 from left to right, one can see that an asset market possessing a lot of cognitive noise with little unrelated news will have a longer optimal time horizon. The middle column shows the opposite: when faced with a lot of unrelated news but little cognitive noise, t^* becomes very small. Both these results follow the dynamics predicted by Equation 3. Intuitively, when markets need more time to fully react to news, this means the optimal time horizon should be longer. In contrast, when markets are able to react quickly to the news and there is the presence of unrelated information coming out over time, using a short event window would be more appropriate. The rightmost column displays the most "realistic" scenario for an asset market: both cognitive noise and unrelated news exist. In such an event, the optimal time horizon lies somewhere between the two extreme scenarios.

Importantly, the simulations also present minimal differences between t^* and $\hat{t}$—the true and estimated time horizons, respectively—when using only the signal s_i. Therefore, a "good" signal allows for the possibility of estimating the time horizon that reflects when markets have fully reacted to news on average. However, recall that this the motivating framework and related simulations have strong linearity and orthogonality assumptions imposed to simply demonstrate the effects of noise on the estimation for the optimal time horizon. In reality, the compositions and processes of the fundamental price and noise components are unknown for equity markets, yet still influence the movement of the overall asset price.

Abstracting away from the simple framework, I argue in the rest of this paper that my text-analysis neural network approach is still able to extract a “good” signal from its approximation of the underlying mapping between monetary policy communications and asset price changes, regardless of what are the true noise processes at any point of time.

B The Primary Accuracy Metric for Judging XLNet

As mentioned in Subsection 4.4, the primary metric for neural network fine-tuning is a generalised R^2 statistic from Hawinkel, Waegeman, and Maere (2024):

$$R^2_{OOS} = 1 - \frac{\widehat{MSE}}{\widehat{MST}} = 1 - \frac{T^{-1}\sum_{i=1}^{T}(y_i - \hat{y}_i)^2}{\frac{T+1}{T(T-1)}\sum_{i=1}^{T}(y_i - \overline{y_{IS}})^2}, \tag{B1}$$

where $\widehat{MSE}$ is the mean squared error of the neural network over the out-of-sample observations; $\widehat{MST}$ is the mean squared error calculated using the in-sample mean $\overline{y_{IS}}$ as prediction over all out-of-sample observations; and T is the size of the testing sample. Although the explicit objective function of the neural network during fine-tuning is to minimise $\widehat{MSE}$, this simultaneously maximises R^2_{OOS}.

This metric is chosen for two main reasons. First, the conventional R^2 (i.e., the proportion of variance explained by a model) breaks down for non-linear methods like neural networks. For such models, the total variance no longer neatly decomposes into model and residual components, meaning the squared Pearson correlation coefficient does not equal R^2. Second, the R^2_{OOS} formula is specifically designed to assess out-of-sample performance, not in-sample fit.[47]

As is standard in machine learning, model performance is compared to a baseline. This baseline assumes *no* relationship between the FOMC statement text and futures price changes, naively predicting with the in-sample average. Therefore, the R^2_{OOS} measures the reduction in predictive error achieved by the neural network compared to this naive baseline.

During fine-tuning, XLNet-Base also tracks other metrics: the out-of-sample Pearson correlation between predicted and actual values, the out-of-sample mean absolute error, and the in-sample mean squared error. This last metric is important for verifying that the network is genuinely learning the underlying relationship.

The decisive criterion for systematic estimation is the out-of-sample R^2 *averaged across sample splits* for each event window length and futures contract:

$$\overline{R^2_{OOS}} = \frac{\sum_{i=1}^{K} R^2_{OOS}}{K}, \tag{B2}$$

where $K = 5$ is the number of sample splits. This $\overline{R^2_{OOS}}$ statistic measures the neural network's average performance improvement within a given window, relative to the naive baseline. The optimal event window is therefore the length that yields the largest $\overline{R^2_{OOS}}$. This is the point where the network's predictive power and generalisability is highest, which

[47] Using a conventional R^2 yields similar optimal window lengths and network quality, but R^2_{OOS} is a more appropriate metric.

only occurs when the impact of noise is minimised and the fundamental reaction is fully captured.

C Construction of High-Frequency Monetary Surprises

C.1 Financial Data Overview and Price Selection

The intraday data on interest-rate and equity futures is sourced from the Thomson Reuters Tick History database and provided by LSEG. The sample period for my analysis is from May 1999 to October 2019. Table F2 provides an overview of the financial outputs data. For each futures contract, I have a minutely series which includes the price of the first and last trades for a given minute. For the rest of this appendix section, I detail the construction of the interest-rate surprises from the contracts. Following previous papers in the literature, I use federal funds futures contracts to measure interest rate expectations within three months, Eurodollar futures to capture expectations between two and four quarters out, and Treasury futures for expectations out to fifteen years.

As discussed in Subsection 3.2, I collect price levels for all futures contracts at 10-minute intervals, starting from 10 minutes before to 18 hours after an FOMC statement release. Specifically, the last recorded traded price 10 minutes before while the first recorded traded price at 10-minute intervals after a statement release are collected. I make the core assumption that the futures price does not change in-between times of recorded trades. For example, if the recorded traded prices for $FF1$ were at 15 and 9 minutes before an FOMC statement release, I assume that the first recorded trading price 15 minutes before the release is the same at 10 minutes beforehand. Therefore, I collect the price 15 minutes beforehand. The equivalent logic applies to collecting futures prices for the 10-minute intervals after a monetary policy announcement.

C.2 Surprises from Federal Funds Futures

On the last day of a given expiry month, a federal funds rate futures (FFF) contract pays out 100 minutes the average (effective) federal funds rate over the expiry month. In other words, let $p_{\tau,t}^{FFj}$ be the price of the $(j-1)$ month-ahead FFF at time t for FOMC meeting τ. Then, the expected average federal funds rate of the $(j-1)$ month ahead FFF at time t for FOMC meeting τ is calculated as $100 - p_{\tau,t}^{FFj}$.

C.2.1 Surprises from Federal Funds Futures: Current Meeting

I calculate the federal funds rate surprise for the current meeting as follows:

1. For each time t, calculate the implied federal funds rate, $100 - p_{\tau,t}^{FFj}$.

2. Calculate $mp1_{\tau,t+n} = \frac{m}{m-d}\left(FF1_{\tau,t+n} - FF1_{\tau,t-10}\right)$, the federal funds rate surprise for the current meeting for each FOMC announcement τ, where d is the day of the announcement and m is the number of days in the month of the announcement.

3. If $m - d + 1 \leq 7$, i.e., the FOMC announcement occurs in the last seven days of the month, use the change in the price of the next month's FFF. In other words, the current-meeting surprise would be calculated as $FF2_{\tau,t+n} - FF2_{\tau,t-10}$. The reason behind this criterion is to avoid multiplying the difference in implied rates with large $\frac{m}{m-d}$.

C.2.2 Surprises from Federal Funds Futures: Next Meeting

I calculate the surprise in the expected federal funds rate, at current FOMC meeting τ, for the next meeting as follows:

1. For a given FOMC announcement τ, find the number of months out $(j-1)$ that has the next scheduled meeting. This determines the maturity of the FFF to use, i.e., j. It was found that $j \in \{2, 3, 4\}$.

2. Calculate $mp2_{\tau,t+n} = \frac{m_2}{m_2-d_2}\left\{[FFj_{\tau,t+n} - FFj_{\tau,t-10}] - \frac{d2}{m2}mp1_{\tau,t+n}\right\}$, where d_2 is the day of the next scheduled FOMC meeting, m_2 is the number of days in the month containing the next scheduled meeting, and $mp1_{\tau,t+n}$ is the current-meeting surprise calculated within an event window beginning 10 minutes before and ending $t+n$ minutes after announcement τ.

3. If $m_2 - d_2 + 1 \leq 7$, meaning the next scheduled FOMC announcement occurs in the last seven days of the month, use the change in the price of the $j+2$-month-ahead FFF contract. In other words, $mp2_{\tau,t+n} = FF(j+1)_{\tau,t+n} - FF(j+1)_{\tau,t-10}$.

C.3 Surprises from Eurodollar Futures

Eurodollar contracts are quarterly contracts that pay out 100 minus the US Dollar BBA LIBOR interest rate at the time of expiration. The last trading day of the contracts is the second London bank business day (usually a Monday) before the third Wednesday of the last month of the expiry quarter. In other words, let $p^{edj}_{\tau,t}$ be the price of the j^{th} closest quarterly Eurodollar futures contract (i.e., March, June, September, December) at time t on FOMC meeting date τ. Note that the expiration date of this contract is between $j-1$ and j quarters into the future at any point of time. We can write the implied interest rate to be $edj_{\tau,t} = 100 - p^{edj}_{\tau,t}$ at time t in FOMC meeting date τ. For a given FOMC announcement τ, I calculate the surprise in the implied rate of the contract maturity j as follows:

1. For each time t, calculate the implied rate, $edj_{\tau,t} = 100 - p^{edj}_{\tau,t}$.

2. Calculate the surprise in the implied rate of contract maturity $j, edj_{\tau,t+n} = edj_{\tau,t+n} - edj_{\tau,t-10}$.

C.4 Surprises from Treasury Futures

At the time of expiration, Treasury futures contracts oblige the seller to deliver Treasury bonds within a range of maturities to the buyer. Let $p^{tkj}_{\tau,t+n}$ be the price at time t on FOMC announcement τ of the j^{th} nearest quarterly k-year Treasury futures contract (i.e., March, June, September, December). To calculate the implied yield surprise around the FOMC meeting $\tau, tk_{\tau,t+n}$, I divide the price change by the approximate maturity and flip the sign:

$$tk_{\tau,t+n} = -\left(p^{tk^j}_{\tau,t+n} - p^{tk^j}_{\tau,t-10}\right)/l, \tag{C.1}$$

where approximate maturities $l \in \{2, 4, 7, 15\}$ come from Gürkaynak, Kisacikoğlu, and Wright (2020). If the FOMC meeting τ falls within the expiry month of the quarter (i.e., March, June, September, December), I use the $(j+1)^{th}$ nearest quarterly Treasury futures contract due to higher liquidity (Gorodnichenko and Ray, 2017).

D Robustness Checks

D.1 Original Surprise Instruments for MP Surprises & Shocks

I perform the same exercises in Subsection 6.2, but for monetary policy shocks constructed with the original instrument set of monetary policy surprises used by the authors.[48]

Visual Shock Differences I plot the shock series constructed in both window lengths using the original instrument set, displayed in Figure E28. For each sub-figure, the horizontal axis depicts the FOMC statement release dates. The vertical axis depicts the principal components re-scaled according to the original specifications of the authors. Interestingly, the differences in all four monetary policy shocks containing forward guidance have extremely similar visual characteristics as those shown in the main results of the paper, which use the full instrument set. The similar differences from event window choice also apply to JK_{CBI}, which still observe more "shifts of importance" when using the optimal event window length due to the sign restriction method of Jarociński and Karadi (2020).

Responses of Interest Rates Regression results following the specification of Equation 8 for the responses of nominal and real interest rates to monetary policy shocks within different event windows are presented in Appendix Table F11 and Appendix Table F12, respectively. Overall, we can see that when changing only the event window length, monetary policy shocks containing forward guidance have stronger point estimates in units of basis points. Interestingly, using the median optimal event window length of 50 minutes also increases the statistical significance of several coefficient estimates, such as regressing the daily change in the two-year TIPS yield on GSS_P or various nominal Treasury yields on the central bank information shock of Jarociński and Karadi (2020).

Responses of Break-even Inflation Regression results for break-even inflation using the original surprise instruments can be found in Appendix Table F13. Similar to the main results of the paper, monetary policy shocks constructed within the optimal event window length results in more negative impacts on break-even inflation for all maturities. Furthermore, moving away from the 30-minute window results in increased statistical significance for several estimates. These results show that the main takeaways of the paper are robust

[48] Technically, Jarociński and Karadi (2020) use $FF4$ instead of the next-meeting surprise, $mp2_\tau$. Because both instruments broadly represent the same underlying maturity and the other authors use this instrument over $FF4$, I construct monetary policy shocks with the method of Jarociński and Karadi (2020) using $mp2_\tau$ for convenience.

to the choice of surprise instruments.

Responses of Equities Using the original instrument set to construct all considered monetary policy shocks, I also investigate the effect of event window size on the impact of monetary policy shocks on equity prices. Appendix Table F14 presents analogous results for the S&P 500 Index and its E-mini futures following the regression form of Equation 8. The sole difference is that the dependent variable is now 100 times the price log-difference of equities. Compared to the main results in Subsection 6.2, the effects of event window choice on the impact of monetary policy shocks are more mixed. For example, while GSS_P results in a decline in the S&P 500 Index by an additional 2.52% when using the optimal window length, other shocks like NS_{MP} see a weaker effect by 3.32%. while all monetary policy shocks still have statistically significant effects on equity prices which go in the expected direction, the effects of window size choice are overall smaller. In contrast, the percentage changes to the responses of the second-month S&P 500 E-mini futures contract are similar to those seen in the main results.

The overall results from constructing monetary policy shocks with the original instrument set suggests that the shorter underlying maturity of the surprises could be understating the shock effects. Because my sample period includes the effective lower bound, the effects of forward guidance and other non-conventional monetary policy tools might not be fully captured within the implied surprises of interest-rate surprises whose expected path only goes out one year into the future. Indeed, recent studies (e.g., Brennan et al., 2024; An, Stedman, and Lusompa, 2025) have shown that incorporating instruments with longer underlying maturities can help prevent this understatement in constructed monetary policy shocks.

Responses of the Macroeconomy via Local Projections I use the lag-augmented local projection method from Olea and Plagborg-Møller (2021) to estimate the responses of log CPI, log IP, the nominal two-year Treasury yield, and the EBP to the considered monetary policy shock series. The separate regressions of each outcome variable on each shock follow the form specified in Equation 9. Figures E29–E32 visually display the impulse responses for all outcome variables to all shock series. Similar to the main results of the paper, the effects of event window length don't really change most point estimates of the responses, although there are some exceptions to this trend (e.g., the EBP increases positively with statistical significance to the NS_{MP} shock starting at 36 months after initial impact). Also similar to the main results is the fact that the confidence intervals of the impulse responses using the optimal window length appear smaller than those using 30-minute windows. When calculating the average and median confidence interval width ratios discussed in Subsection 6.2, using the original instrument set yields ratios of 0.9033 and 0.9506, respectively. In other words, using the optimal event window length results in estimated responses of macroeconomic variables to be more precise.

D.2 Optimal Event Window Lengths Beyond 70 Minutes?

An additional advantage with the "one signal" approach is that I am able to crudely check if event windows beyond the optimal window length from the "joint" estimation have a larger $\overline{R^2_{OOS}}$. Appendix Figures E21–E27 visually display the results of this exercise for maturities of federal funds futures, Eurodollar futures, Treasury futures of two, five, ten, and thirty years, and both the S&P 500 Index and its E-mini futures. The horizontal axis depicts the event window lengths, starting from the systematically estimated event window length and ending at an event window starting 10 minutes before and ending 18 hours after FOMC statement release. The cross points represent the event window length associated with the highest $\overline{R^2_{OOS}}$. One can see that for most maturities and assets, the event window length with the largest $\overline{R^2_{OOS}}$ is the length systematically estimated by the "joint" approach. In other words, the results from this exercise can be viewed as a robustness check in support of XLNet-Base providing a "good" signal overall for systematic event window length estimation. However, there are few futures contracts that show a difference between the "joint" and "one signal" estimated event window lengths, such as how the "one signal" approach estimates the optimal window length for the 1-month-ahead federal fund futures ($FF2$), 3-month-ahead federal fund futures $FF4$), front-month 2-year Treasury futures ($TUc1$), second-month 10-year Treasury futures ($TYc2$), and front-month 30-year Treasury futures ($USc1$) are longer than those systematically estimated. To confirm the validity of these estimated window lengths, I calculate $\overline{R^2_{OOS}}$ under the "joint" approach for these futures. The results are depicted as the rightmost box-and-whisker plot in the respective sub-figures of systematic estimation for these futures contracts. Two insights can be taken away from this robustness exercise. First, the optimal event window for these futures contracts is still within 60 minutes in length. However, "joint" estimation for all these markets showed that the considered window lengths beyond 60 minutes after statement released yielded $\overline{R^2_{OOS}}$ greater than that obtained within a 30-minute window. This leads to the second insight, which is the interesting possibility that the global optimal event window length might possibly be between 60 minutes and the amount of time predicted by the "one signal" approach after FOMC statement release. Regardless, a 30-minute window is never predicted to be the optimal event window length.

Appendix Figures

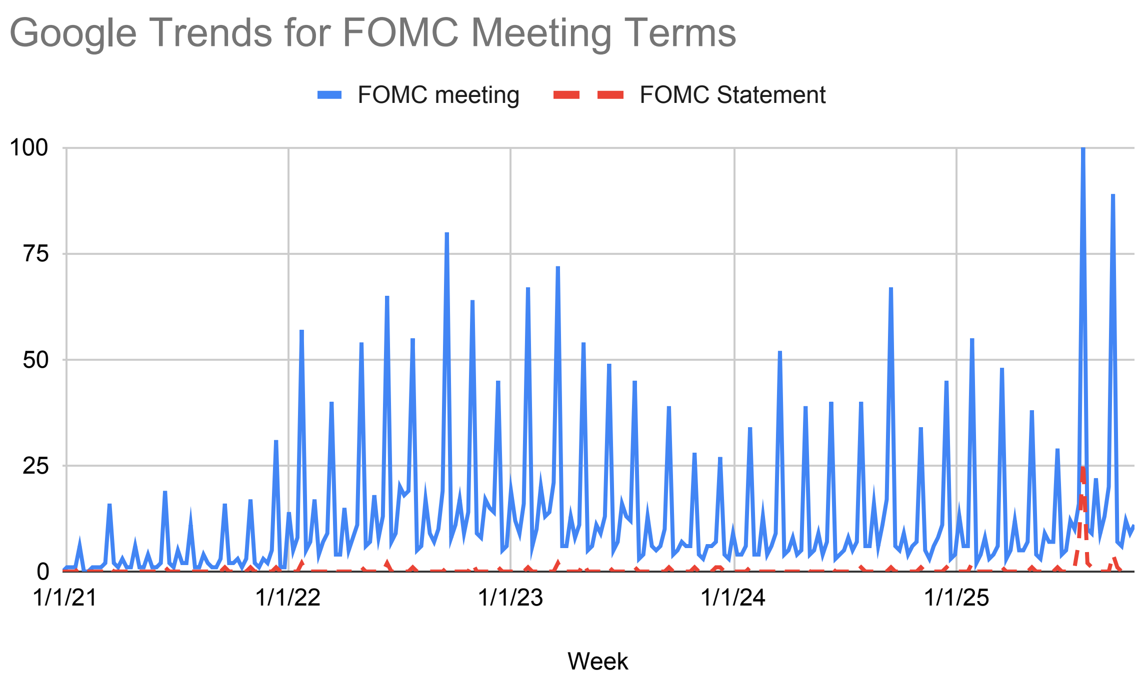

Figure E1: Google Trends Interest About FOMC Statements over Time, January 2021–October 2025

Notes: The blue and solid (red and dashed) line represents the search interest for the phrase "FOMC meeting" ("FOMC statement") in the U.S. The interest gauge is substantially greater than

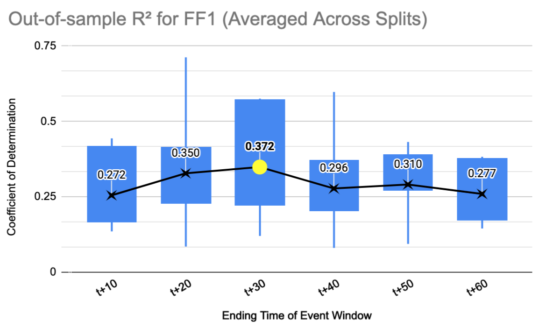

(a) $FF1$

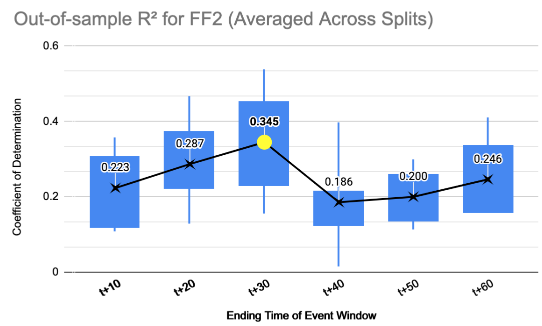

(b) $FF2$

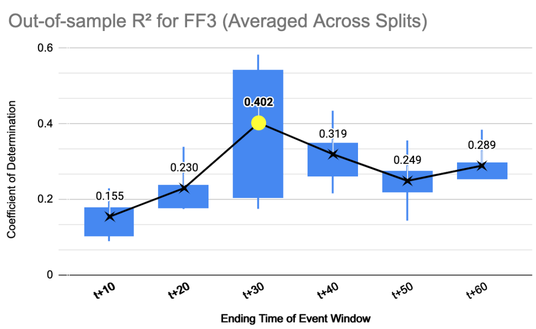

(c) $FF3$

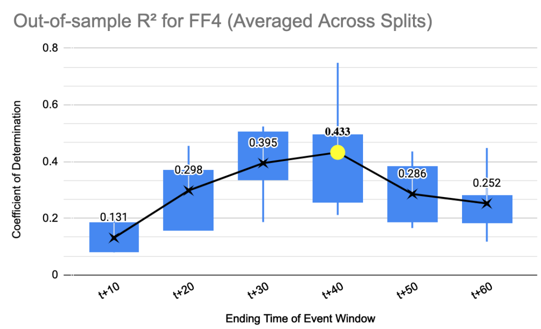

(d) $FF4$

Figure E2: Optimal Event Window Lengths for Federal Funds Futures

Notes: The horizontal axis of each figure depicts the end time of the considered event window lengths and the vertical axis represents the percentage points of price changes by the market directly and only due to the FOMC statement text. The cross points represent $\overline{R^2_{OOS}}$ for each event window size, where the solid yellow point represents the event window with the largest $\overline{R^2_{OOS}}$. For each event window, box-and-whisker plots are shown surrounding the corresponding averages. $0 \leq \overline{R^2_{OOS}} \leq 1$ represents the proportion of the mean squared error from predicting with the in-sample average explained by the superior out-of-sample performance of XLNet-Base.

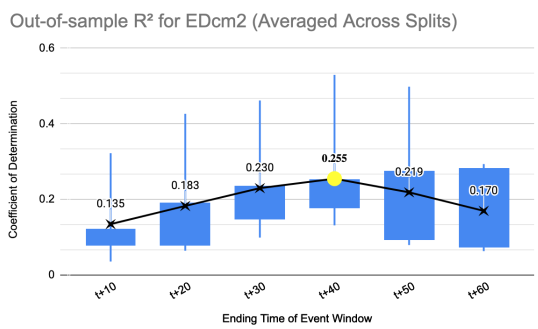

(a) $EDcm2$

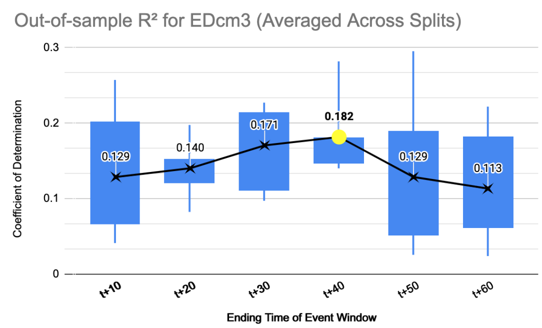

(b) $EDcm3$

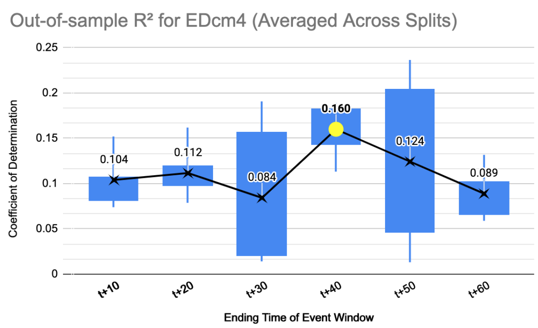

(c) $EDcm4$

Figure E3: Optimal Event Window Lengths for Eurodollar Futures

Notes: The horizontal axis of each figure depicts the end time of the considered event window lengths and the vertical axis represents the percentage points of price changes by the market directly and only due to the FOMC statement text. The cross points represent $\overline{R^2_{OOS}}$ for each event window size, where the solid yellow point represents the event window with the largest $\overline{R^2_{OOS}}$. For each event window, box-and-whisker plots are shown surrounding the corresponding averages. $0 \leq \overline{R^2_{OOS}} \leq 1$ represents the proportion of the mean squared error from predicting with the in-sample average explained by the superior out-of-sample performance of XLNet-Base.

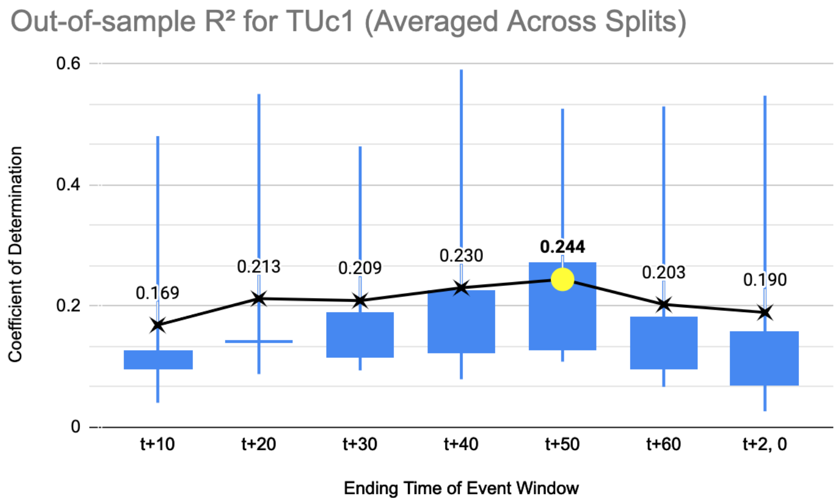

(a) $TUc1$

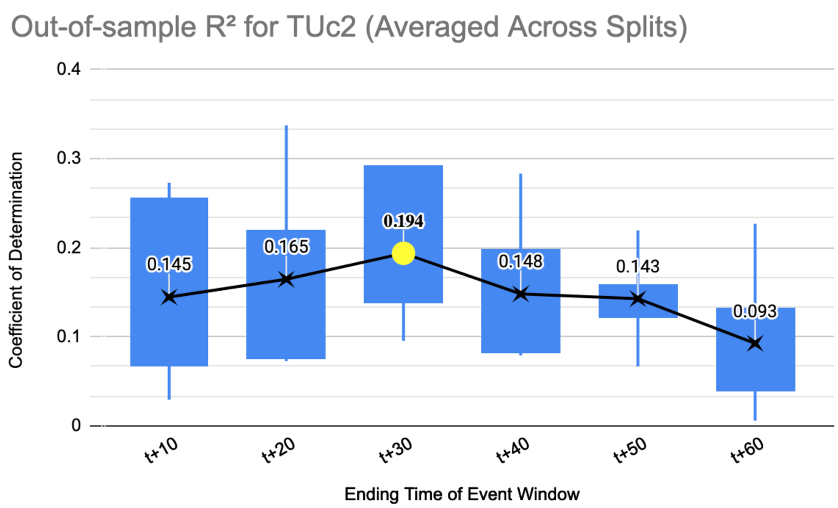

(b) $TUc2$

Figure E4: Optimal Event Window Lengths for 2-Year Treasury Futures

Notes: The horizontal axis of each figure depicts the end time of the considered event window lengths and the vertical axis represents the percentage points of price changes by the market directly and only due to the FOMC statement text. The cross points represent $\overline{R^2_{OOS}}$ for each event window size, where the solid yellow point represents the event window with the largest $\overline{R^2_{OOS}}$. For each event window, box-and-whisker plots are shown surrounding the corresponding averages. $0 \leq \overline{R^2_{OOS}} \leq 1$ represents the proportion of the mean squared error from predicting with the in-sample average explained by the superior out-of-sample performance of XLNet-Base.

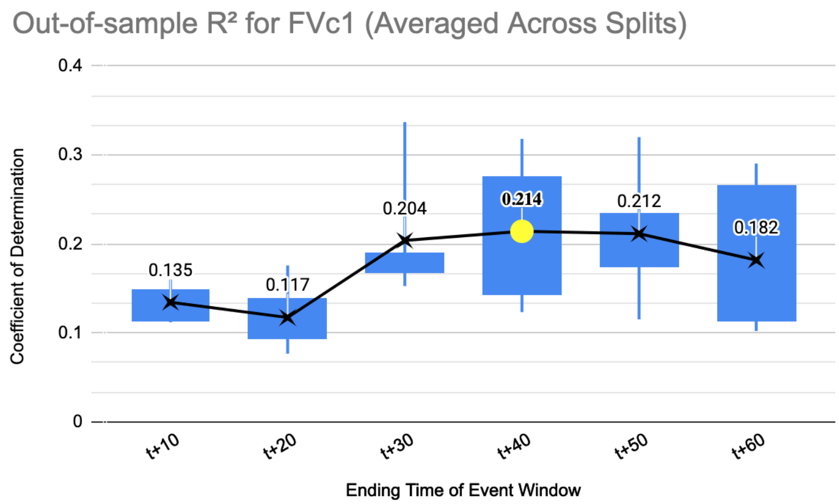

(a) $FVc1$

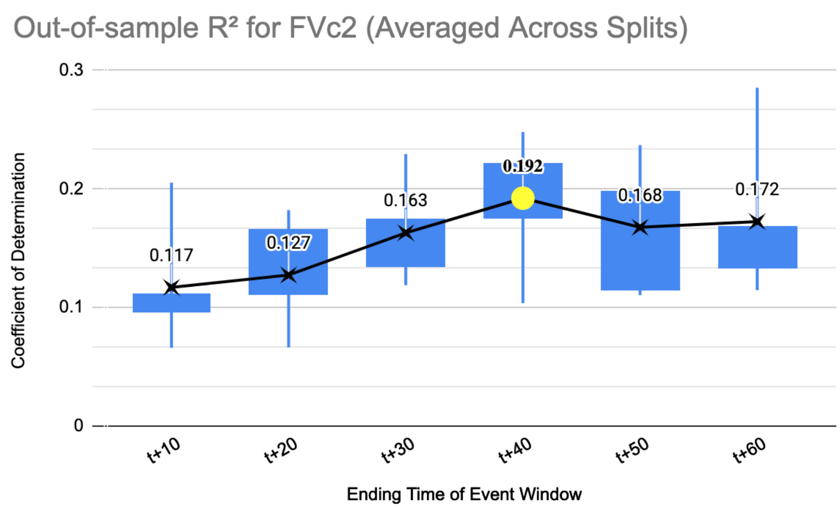

(b) $FVc2$

Figure E5: Optimal Event Window Lengths for 5-Year Treasury Futures

Notes: The horizontal axis of each figure depicts the end time of the considered event window lengths and the vertical axis represents the percentage points of price changes by the market directly and only due to the FOMC statement text. The cross points represent $\overline{R^2_{OOS}}$ for each event window size, where the solid yellow point represents the event window with the largest $\overline{R^2_{OOS}}$. For each event window, box-and-whisker plots are shown surrounding the corresponding averages. $0 \leq \overline{R^2_{OOS}} \leq 1$ represents the proportion of the mean squared error from predicting with the in-sample average explained by the superior out-of-sample performance of XLNet-Base.

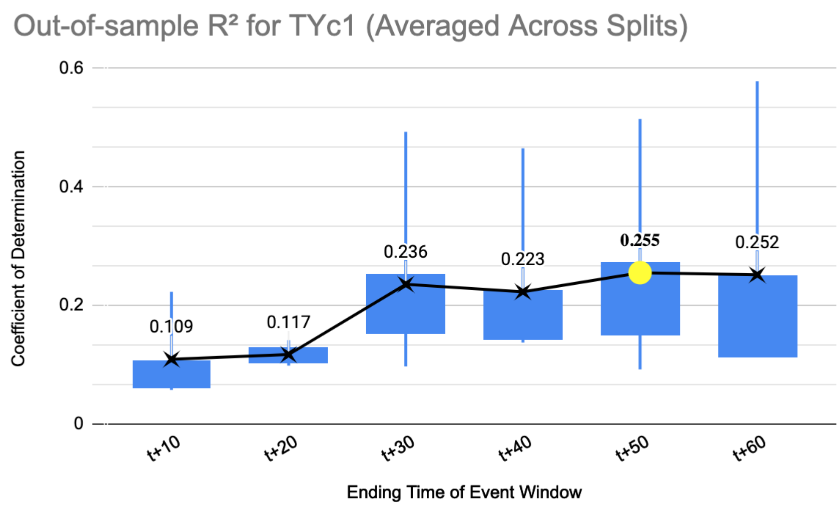

(a) $TYc1$

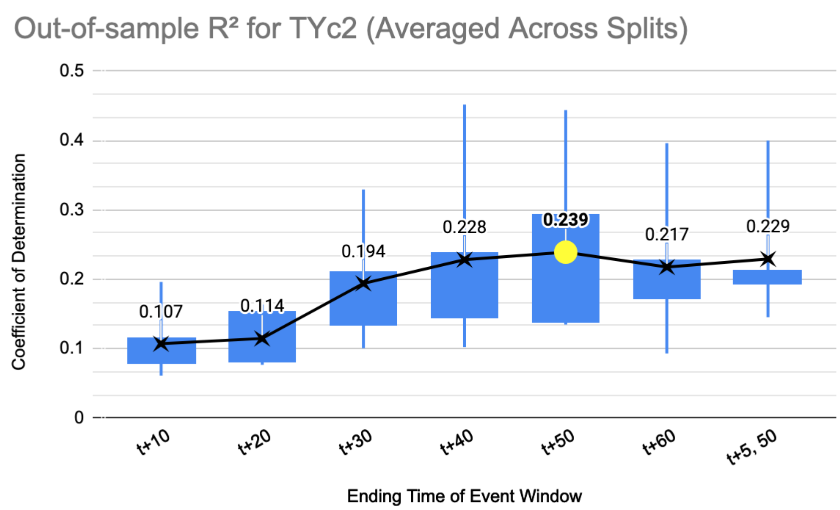

(b) $TYc2$

Figure E6: Optimal Event Window Lengths for 10-Year Treasury Futures

Notes: The horizontal axis of each figure depicts the end time of the considered event window lengths and the vertical axis represents the percentage points of price changes by the market directly and only due to the FOMC statement text. The cross points represent $\overline{R^2_{OOS}}$ for each event window size, where the solid yellow point represents the event window with the largest $\overline{R^2_{OOS}}$. For each event window, box-and-whisker plots are shown surrounding the corresponding averages. $0 \leq \overline{R^2_{OOS}} \leq 1$ represents the proportion of the mean squared error from predicting with the in-sample average explained by the superior out-of-sample performance of XLNet-Base.

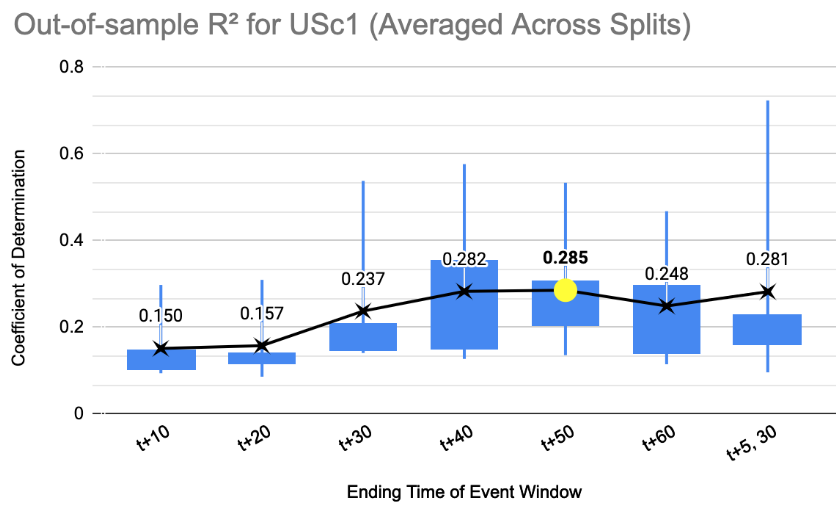

(a) $USc1$

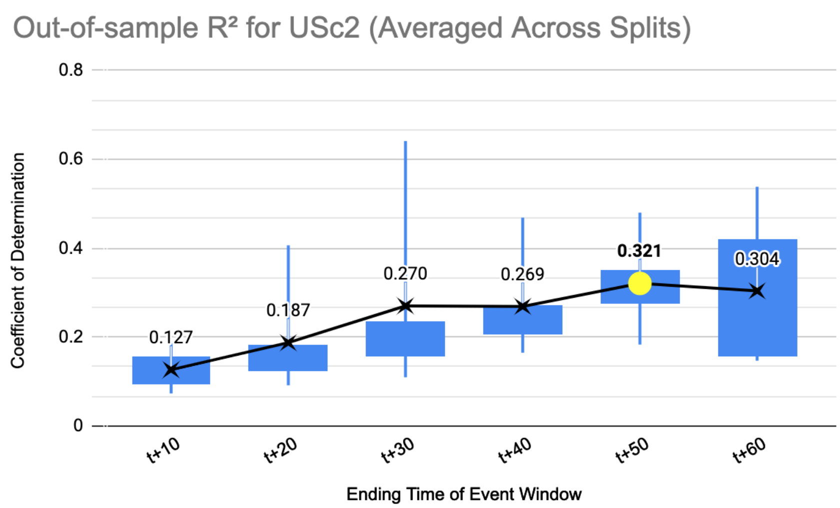

(b) $USc2$

Figure E7: Optimal Event Window Lengths for 30-Year Treasury Futures

Notes: The horizontal axis of each figure depicts the end time of the considered event window lengths and the vertical axis represents the percentage points of price changes by the market directly and only due to the FOMC statement text. The cross points represent $\overline{R^2_{OOS}}$ for each event window size, where the solid yellow point represents the event window with the largest $\overline{R^2_{OOS}}$. For each event window, box-and-whisker plots are shown surrounding the corresponding averages. $0 \leq \overline{R^2_{OOS}} \leq 1$ represents the proportion of the mean squared error from predicting with the in-sample average explained by the superior out-of-sample performance of XLNet-Base.

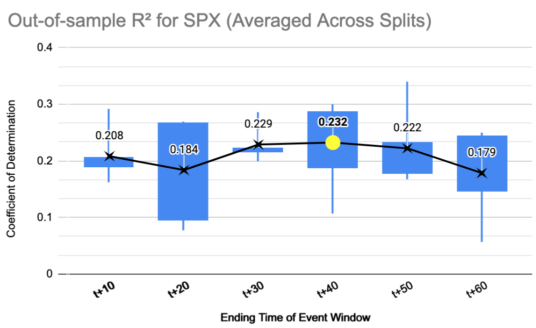

(a) SPX

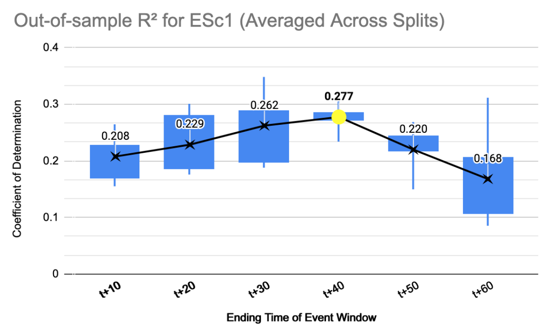

(b) $ESc1$

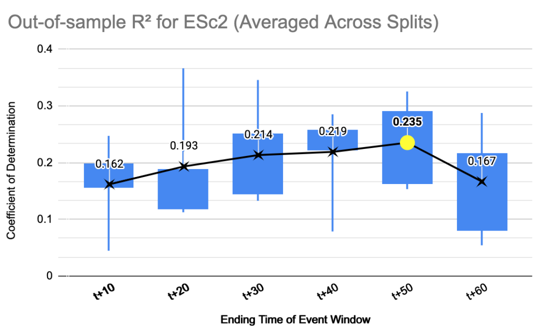

(c) $ESc2$

Figure E8: Optimal Event Window Lengths for S&P 500 and E-mini Futures

Notes: The horizontal axis of each figure depicts the end time of the considered event window lengths and the vertical axis represents the percentage points of price changes by the market directly and only due to the FOMC statement text. The cross points represent $\overline{R^2_{OOS}}$ for each event window size, where the solid yellow point represents the event window with the largest $\overline{R^2_{OOS}}$. For each event window, box-and-whisker plots are shown surrounding the corresponding averages. $0 \leq \overline{R^2_{OOS}} \leq 1$ represents the proportion of the mean squared error from predicting with the in-sample average explained by the superior out-of-sample performance of XLNet-Base.

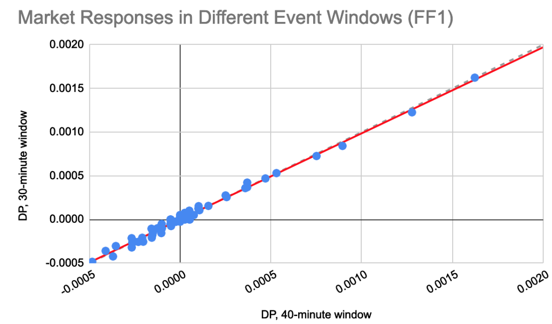

(a) $FF1^{**}$

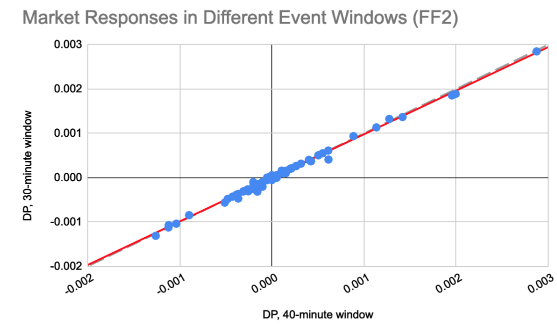

(b) $FF2^{***}$

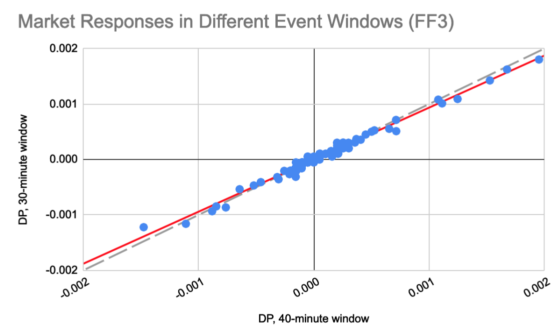

(c) $FF3^{***}$

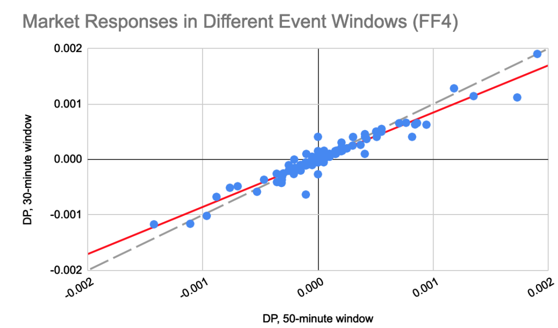

(d) $FF4^{***}$

Figure E9: Comparing Market Responses in Different Event Windows for Federal Funds Futures

Notes: The horizontal axis depicts the price log-difference within the systematically estimated event window length. The vertical axis represents DP_{t+20}. The 45-degree line is depicted as grey and dashed. The blue dots are market reactions on scheduled FOMC meetings. The price-log differences calculated within the optimal window lengths are regressed on DP_{t+20} through OLS. If the slope coefficient is greater (less) than one and its difference with one is statistically significant for at most $\alpha = 0.10$, then the financial market under-reacts to the information content of FOMC statements on release, ex post, and is depicted in red. ** sig. at the 5% level, *** sig. at the 1% level.

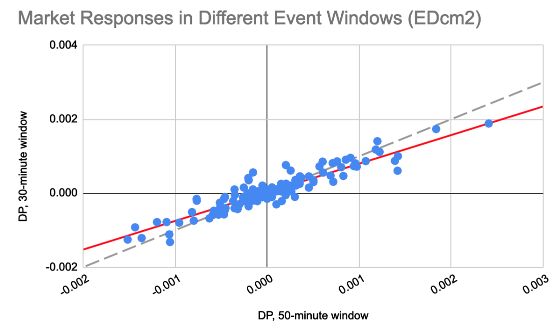

(a) $EDcm2^{***}$

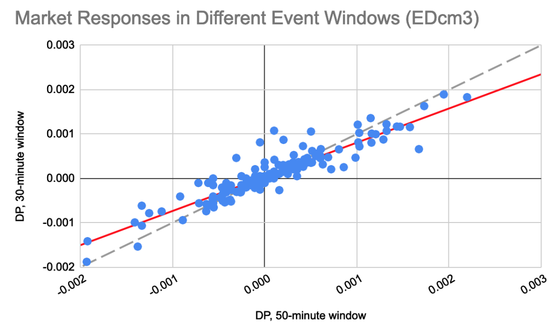

(b) $EDcm3^{***}$

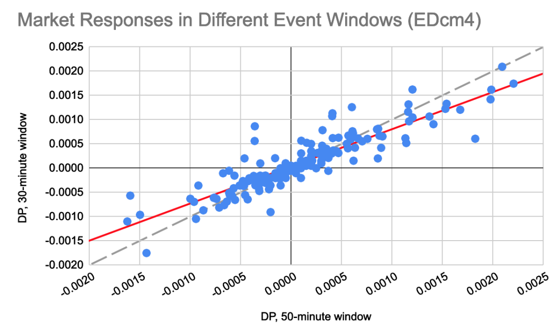

(c) $EDcm4^{***}$

Figure E10: Comparing Market Responses in Different Event Windows for Eurodollar Futures

Notes: The horizontal axis depicts the price log-difference within the systematically estimated event window length. The vertical axis represents DP_{t+20}. The 45-degree line is depicted as grey and dashed. The blue dots are market reactions on scheduled FOMC meetings. The price-log differences calculated within the optimal window lengths are regressed on DP_{t+20} through OLS. If the slope coefficient is greater (less) than one and its difference with one is statistically significant for at most $\alpha = 0.10$, then the financial market under-reacts to the information content of FOMC statements on release, ex post, and is depicted in red. ** sig. at the 5% level, *** sig. at the 1% level.

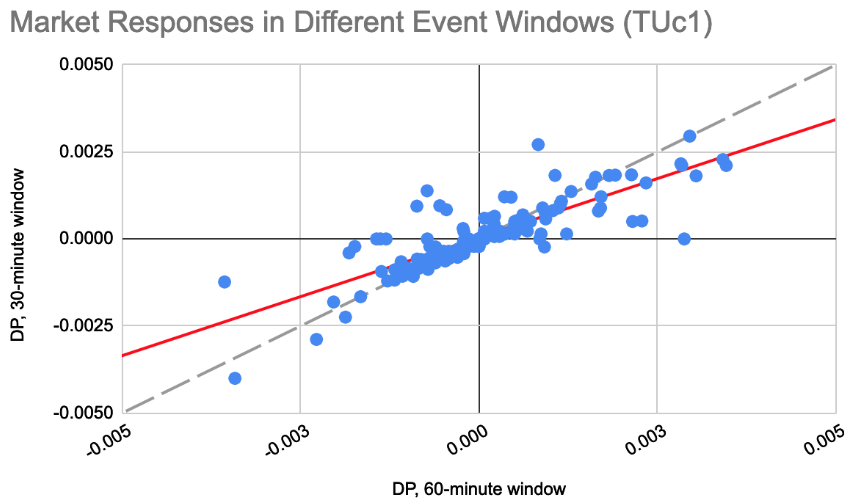

(a) $TUc1^{***}$

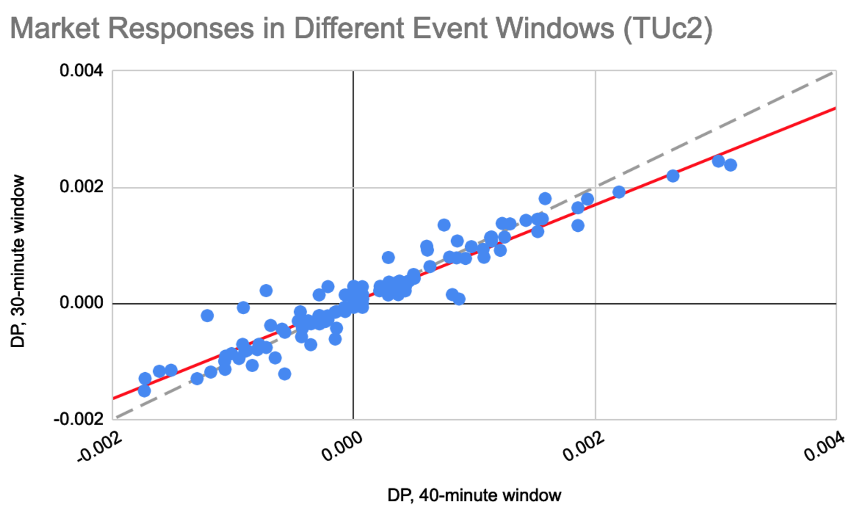

(b) $TUc2^{***}$

Figure E11: Comparing Market Responses in Different Event Windows for 2-Year Treasury Futures
Notes: The horizontal axis depicts the price log-difference within the systematically estimated event window length. The vertical axis represents DP_{t+20}. The 45-degree line is depicted as grey and dashed. The blue dots are market reactions on scheduled FOMC meetings. The price-log differences calculated within the optimal window lengths are regressed on DP_{t+20} through OLS. If the slope coefficient is greater (less) than one and its difference with one is statistically significant for at most $\alpha = 0.10$, then the financial market under-reacts to the information content of FOMC statements on release, ex post, and is depicted in red. ** sig. at the 5% level, *** sig. at the 1% level.

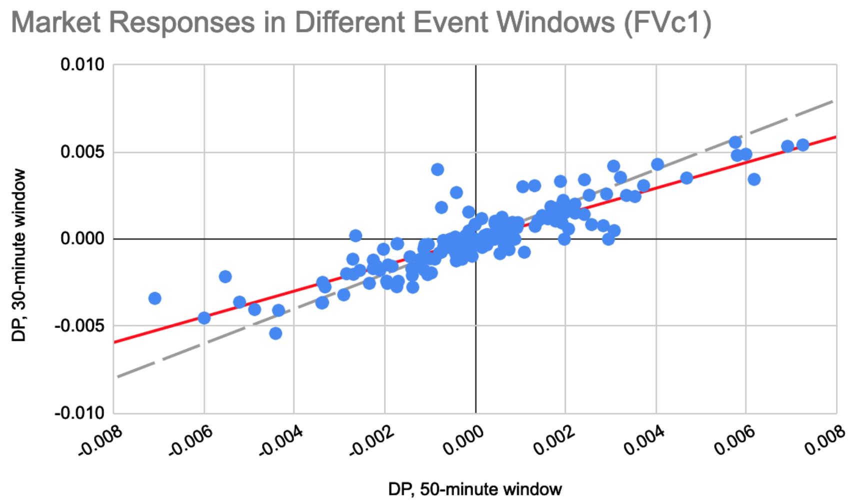

(a) $FVc1^{***}$

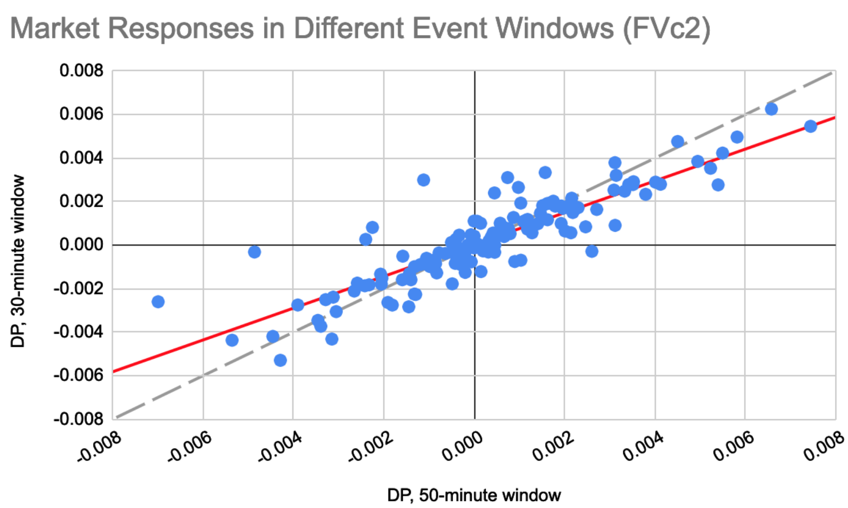

(b) $FVc2^{***}$

Figure E12: Comparing Market Responses in Different Event Windows for 5-Year Treasury Futures

Notes: The horizontal axis depicts the price log-difference within the systematically estimated event window length. The vertical axis represents DP_{t+20}. The 45-degree line is depicted as grey and dashed. The blue dots are market reactions on scheduled FOMC meetings. The price-log differences calculated within the optimal window lengths are regressed on DP_{t+20} through OLS. If the slope coefficient is greater (less) than one and its difference with one is statistically significant for at most $\alpha = 0.10$, then the financial market under-reacts to the information content of FOMC statements on release, ex post, and is depicted in red. ** sig. at the 5% level, *** sig. at the 1% level.

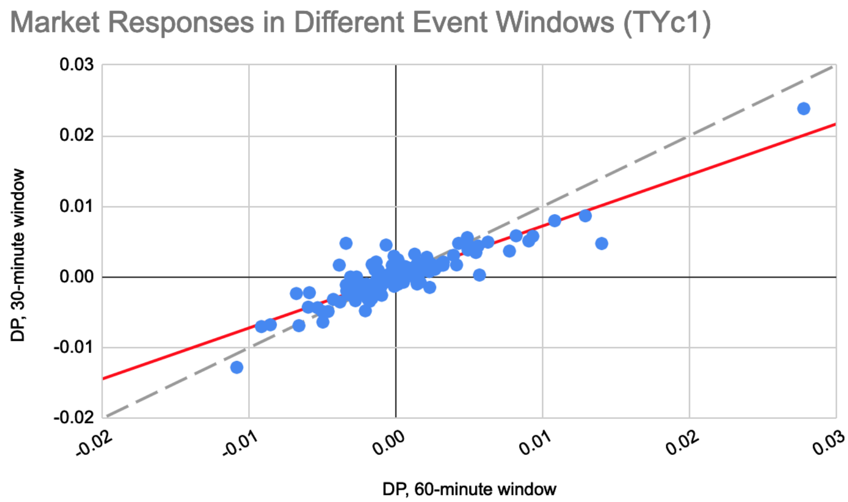

(a) $TYc1^{***}$

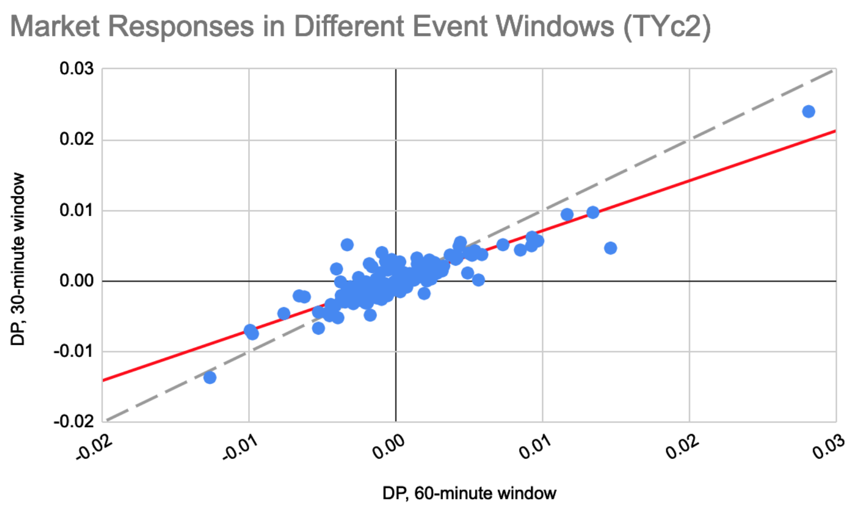

(b) $TYc2^{***}$

Figure E13: Comparing Market Responses in Different Event Windows for 10-Year Treasury Futures

Notes: The horizontal axis depicts the price log-difference within the systematically estimated event window length. The vertical axis represents DP_{t+20}. The 45-degree line is depicted as grey and dashed. The blue dots are market reactions on scheduled FOMC meetings. The price-log differences calculated within the optimal window lengths are regressed on DP_{t+20} through OLS. If the slope coefficient is greater (less) than one and its difference with one is statistically significant for at most $\alpha = 0.10$, then the financial market under-reacts to the information content of FOMC statements on release, ex post, and is depicted in red. ** sig. at the 5% level, *** sig. at the 1% level.

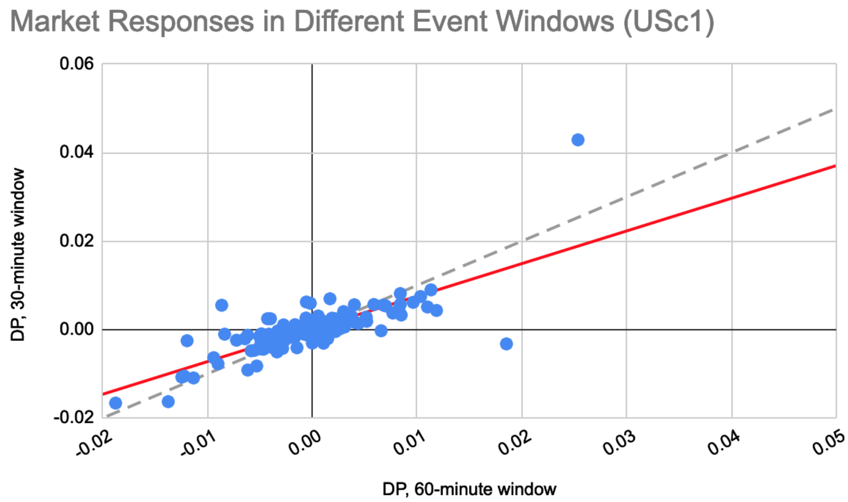

(a) $USc1^{***}$

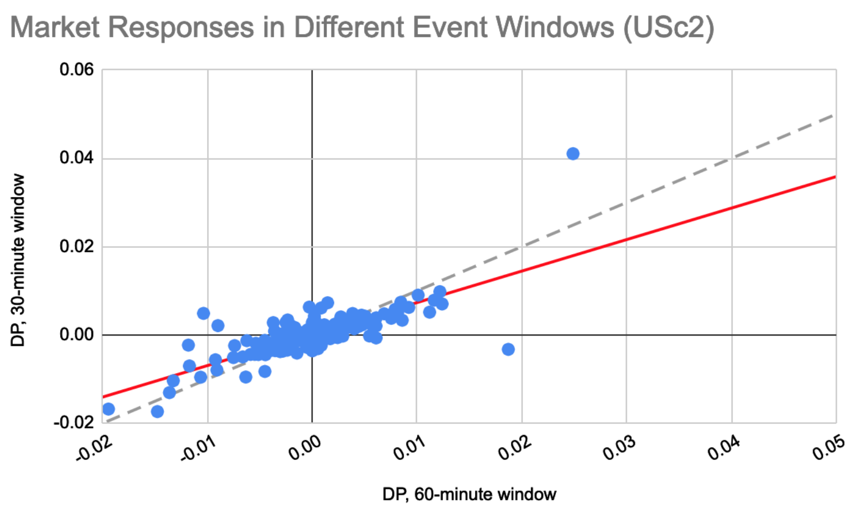

(b) $USc2^{***}$

Figure E14: Comparing Market Responses in Different Event Windows for 30-Year Treasury Futures
Notes: The horizontal axis depicts the price log-difference within the systematically estimated event window length. The vertical axis represents DP_{t+20}. The 45-degree line is depicted as grey and dashed. The blue dots are market reactions on scheduled FOMC meetings. The price-log differences calculated within the optimal window lengths are regressed on DP_{t+20} through OLS. If the slope coefficient is greater (less) than one and its difference with one is statistically significant for at most $\alpha = 0.10$, then the financial market under-reacts to the information content of FOMC statements on release, ex post, and is depicted in red. ** sig. at the 5% level, *** sig. at the 1% level.

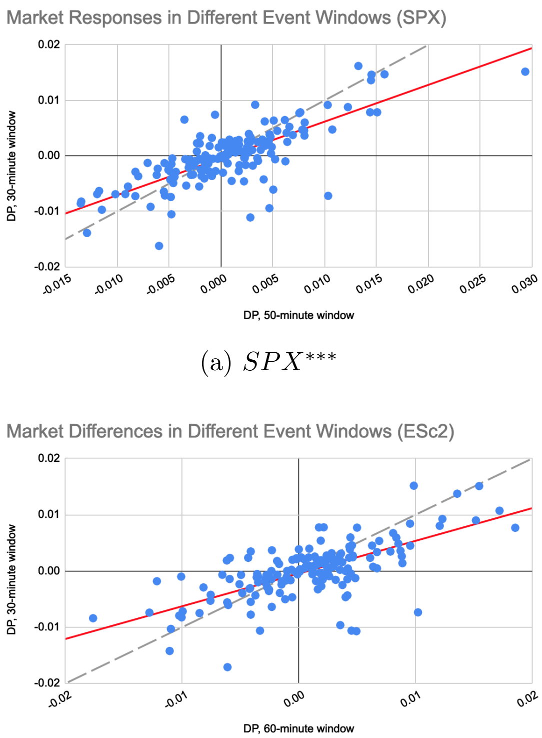

(a) SPX^{***} (b) $ESc1^{***}$

(c) $ESc2^{***}$

Figure E15: Comparing Market Responses in Different Event Windows for S&P 500 and E-mini Futures

Notes: The horizontal axis depicts the price log-difference within the systematically estimated event window length. The vertical axis represents DP_{t+20}. The 45-degree line is depicted as grey and dashed. The blue dots are market reactions on scheduled FOMC meetings. The price-log differences calculated within the optimal window lengths are regressed on DP_{t+20} through OLS. If the slope coefficient is greater (less) than one and its difference with one is statistically significant for at most $\alpha = 0.10$, then the financial market under-reacts to the information content of FOMC statements on release, ex post, and is depicted in red. ** sig. at the 5% level, *** sig. at the 1% level.

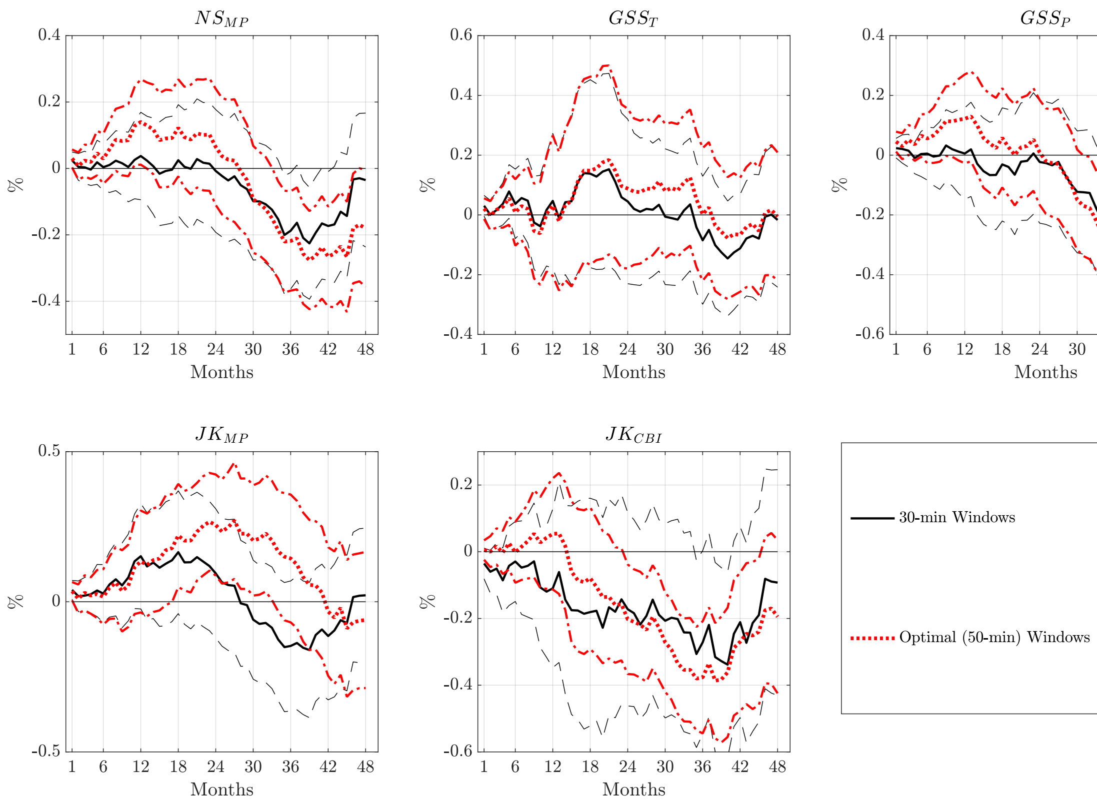

Figure E16: Effects of Event Window Choice on IP Impulse Responses

Notes: Impulse responses are calculated using the lag-augmented local projection from Olea and Plagborg-Møller (2021). Confidence bands are at the 90% level. Standard errors are Eicker-Huber-White heteroscedasticity-robust standard errors. The responses are to a 100 basis point increase in the monetary policy shock series. The monetary policy shocks are constructed using the full set of monetary policy surprise instruments.

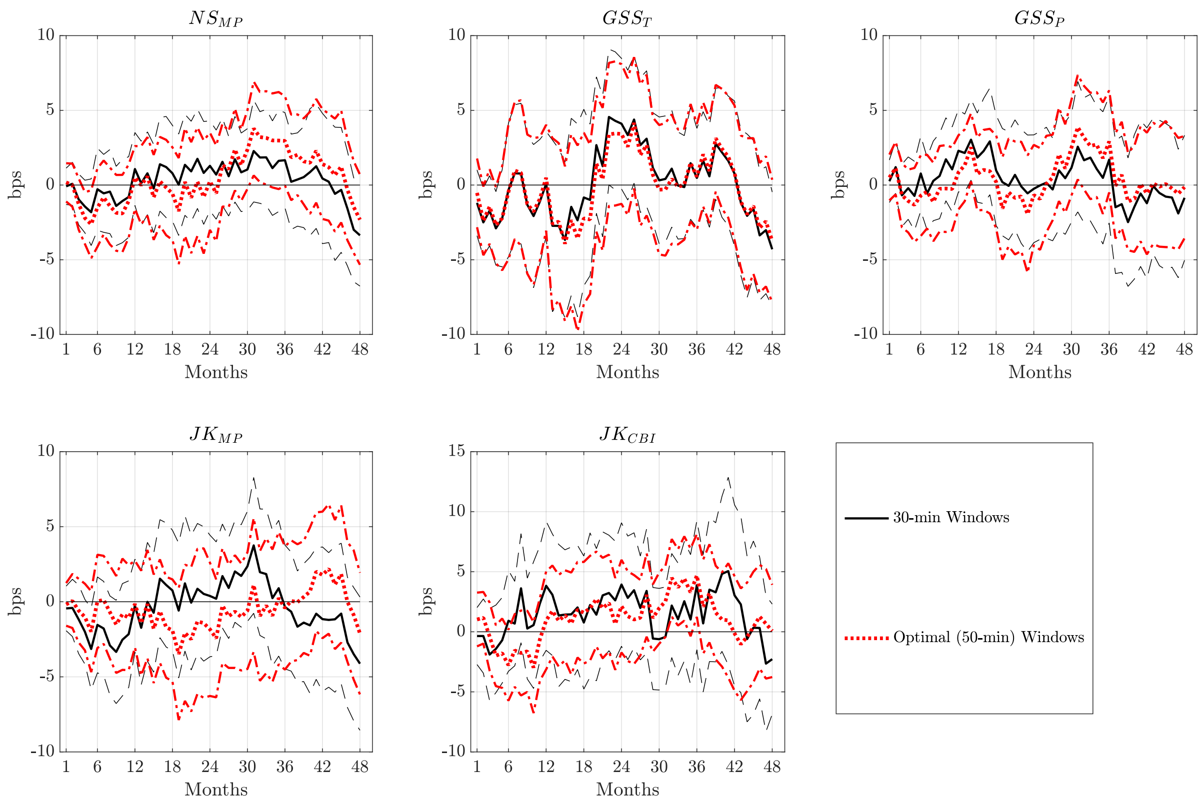

Figure E17: Effects of Event Window Choice on Excess Bond Premium Impulse Responses

Notes: Impulse responses are calculated using the lag-augmented local projection from Olea and Plagborg-Møller (2021). Confidence bands are at the 90% level. Standard errors are Eicker-Huber-White heteroscedasticity-robust standard errors. The responses are to a 100 basis point increase in the monetary policy shock series. The monetary policy shocks are constructed using the full set of monetary policy surprise instruments.

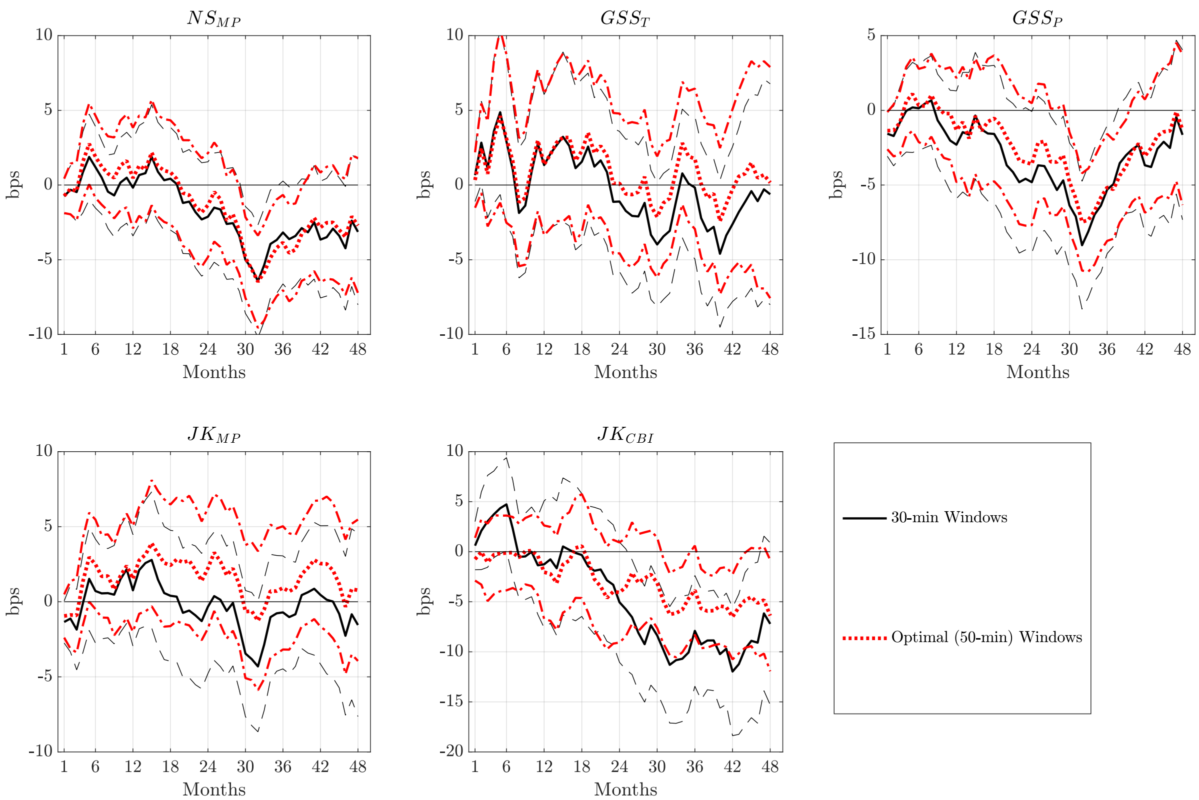

Figure E18: Effects of Event Window Choice on Two-year Treasury Yield Impulse Responses

Notes: Impulse responses are calculated using the lag-augmented local projection from Olea and Plagborg-Møller (2021). Confidence bands are at the 90% level. Standard errors are Eicker-Huber-White heteroscedasticity-robust standard errors. The responses are to a 100 basis point increase in the monetary policy shock series. The monetary policy shocks are constructed using the full set of monetary policy surprise instruments.

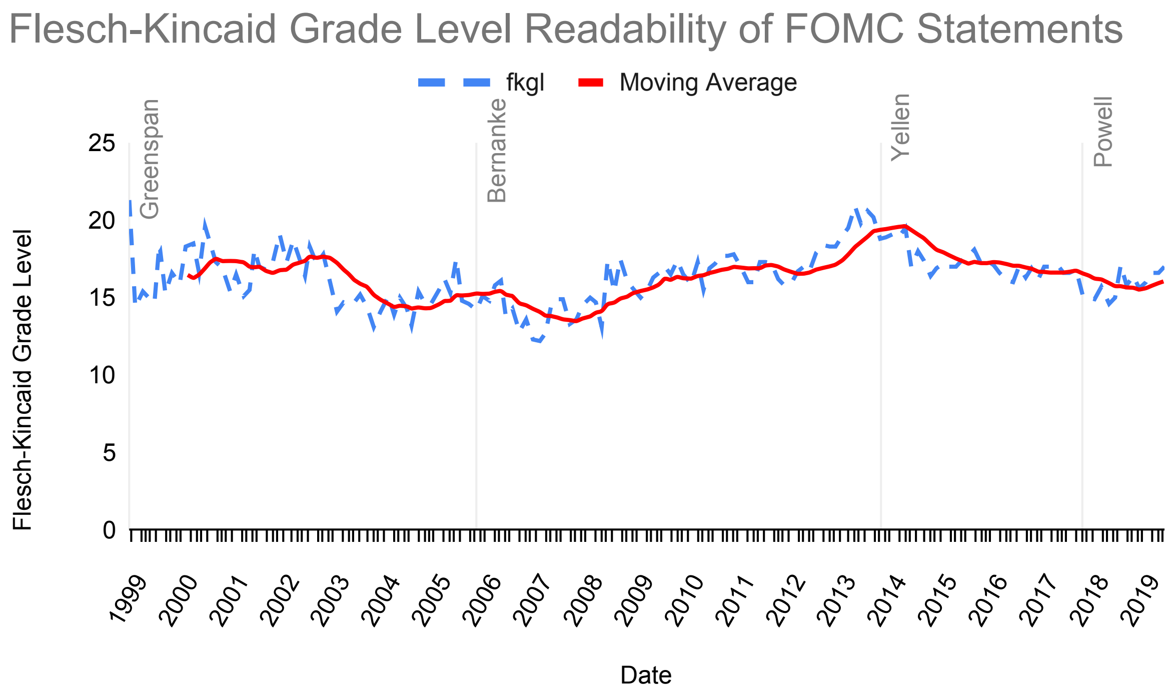

Figure E19: Flesch-Kincaid Grade Level Readability of FOMC Statements

Notes: The complexity of FOMC statements is measured by the Flesch-Kincaid Grade Level, defined as: $0.39 \times$ average sentence length $+ 11.8 \times$ average number of syllables per word $-$ 15.59. From left to right, the vertical grey lines indicate the first FOMC meeting with Greenspan, Bernanke, Yellen, and Powell as Fed Chair. The moving average in solid red is calculated with a period of 10. For a description of U.S. education grade levels, see https://en.wikipedia.org/wiki/Education_in_the_United_States.

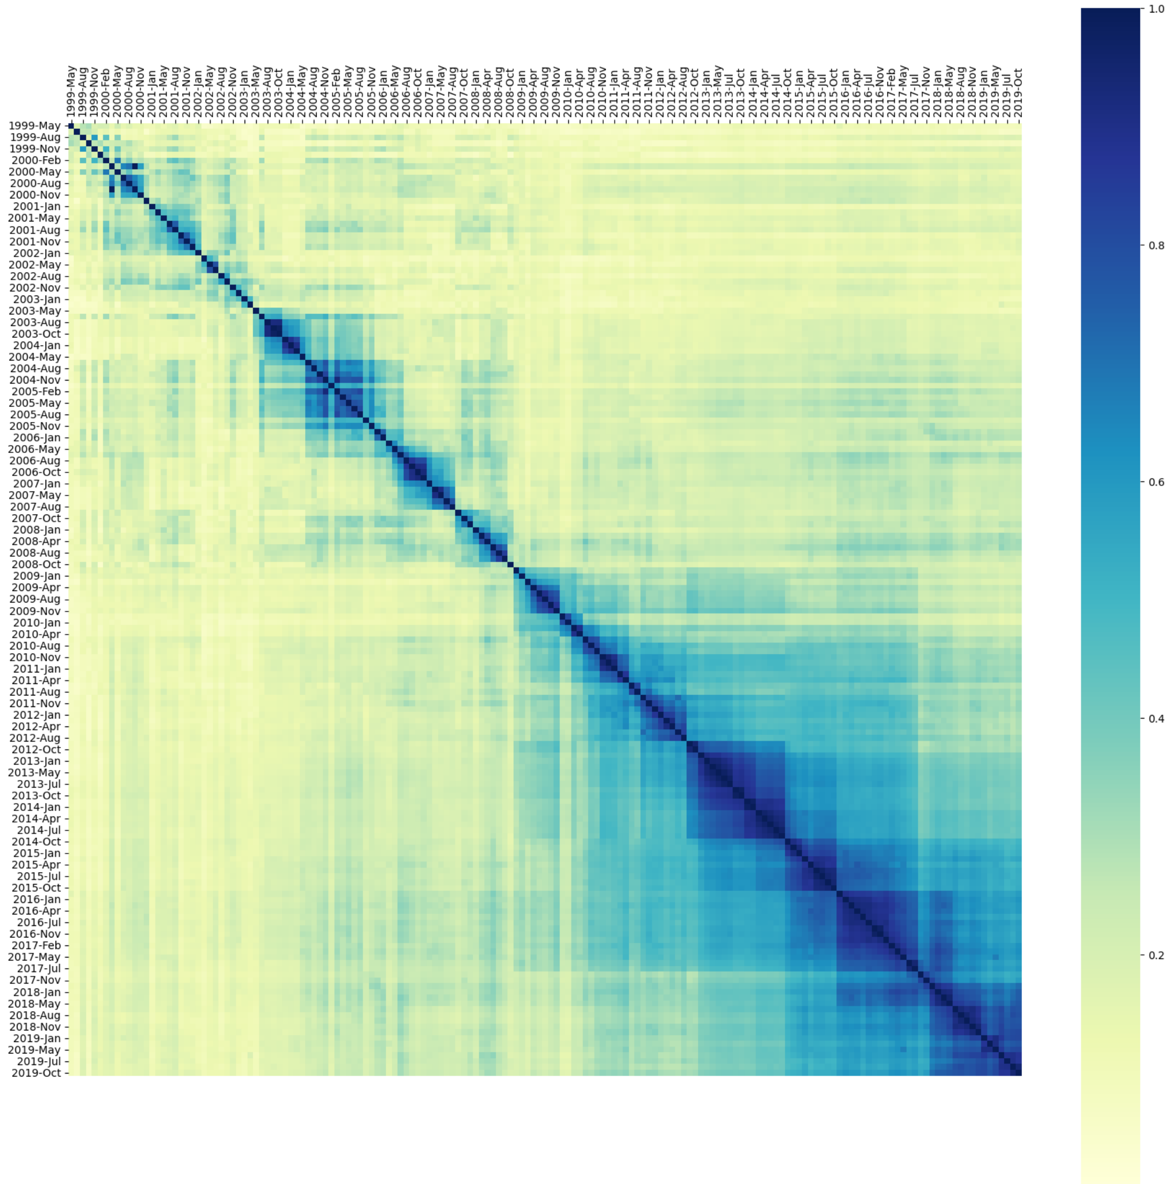

Figure E20: Document Similarity Matrix for FOMC Statements

Notes: Each element is the cosine similarity between two FOMC statements of that row and column. The cosine similarity value measures how similar the terms of both statements are. The darker shade of blue represents pairs of documents that have higher similarity measures closer to one, while the light shade of green represents pairs of documents that do not have terms in common and similarity values closer to zero. The main diagonal has all ones because the cosine similarity value is calculated between an FOMC statement with itself. All FOMC statements have been preprocessed such that all words are lowercase, all words are reduced to base form, and all stop-words are removed.

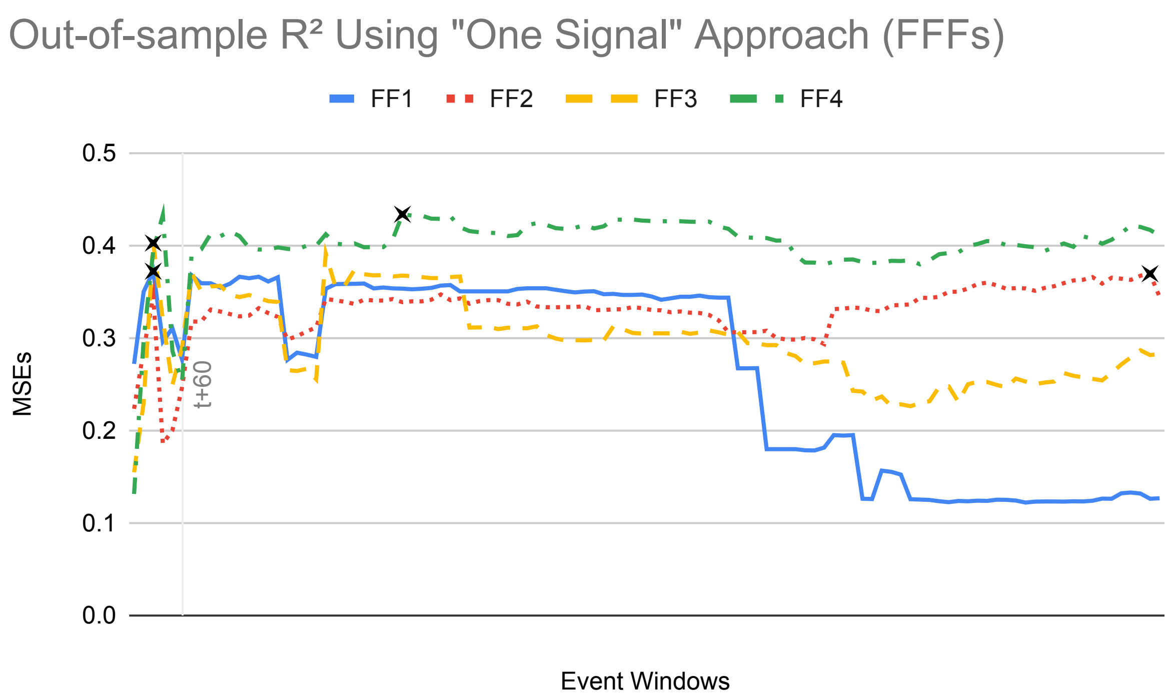

Figure E21: $\overline{R^2_{OOS}}$ Calculated Using the "One Signal" Approach for Federal Funds Futures

Notes: The horizontal axis depicts the event window lengths, starting from a 20-minute event window and ending at an event window starting 10 minutes before and ending 18 hours after FOMC statement release. The cross points represent the event window length associated with the largest $\overline{R^2_{OOS}}$. Estimates to the left of the vertical grey line labelled "$t + 60$" are calculated using the "joint' approach. The 1-month and 3-month-ahead federal funds futures contract are the two exceptions that sees the largest $\overline{R^2_{OOS}}$ within a window length than the systematically estimated length. As a robustness check, $\overline{R^2_{OOS}}$ was calculated using systematic estimation for both futures and considered event window lengths, which can be found in the rightmost box-and-whisker plot of Appendix Sub-figures E2b and E2d, respectively.

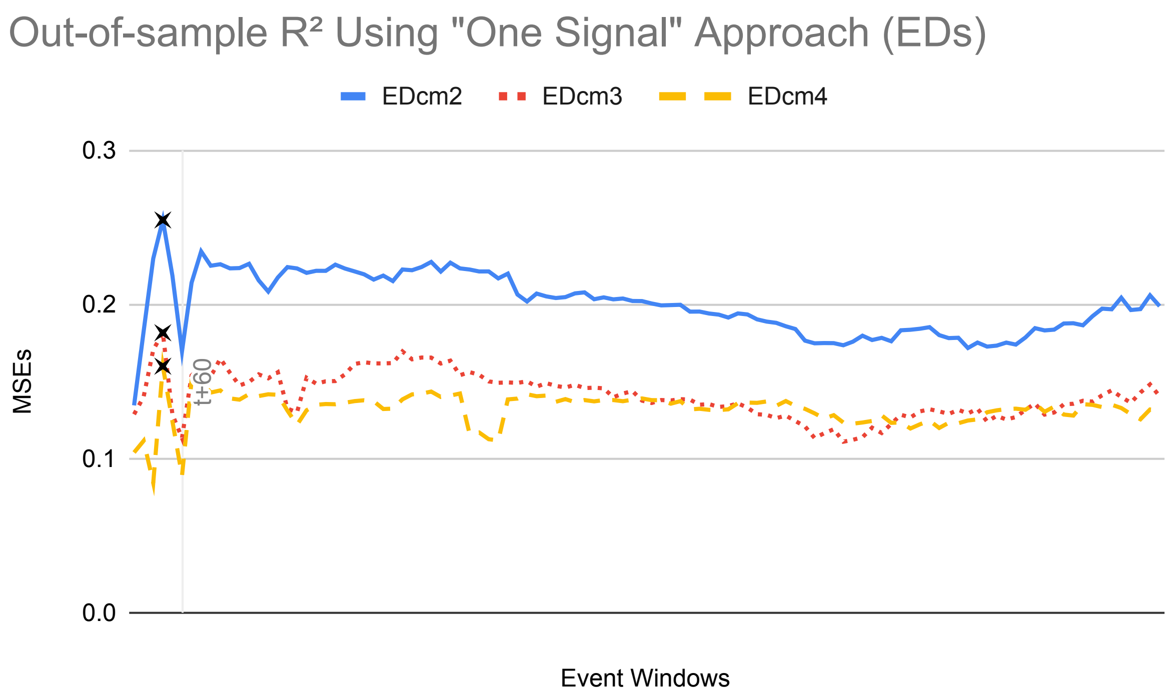

Figure E22: $\overline{R^2_{OOS}}$ Calculated Using the "One Signal" Approach for Eurodollar Futures

Notes: The horizontal axis depicts the event window lengths, starting from a 20-minute event window and ending at an event window starting 10 minutes before and ending 18 hours after FOMC statement release. The cross points represent the event window length associated with the largest $\overline{R^2_{OOS}}$. Estimates to the left of the vertical grey line labelled "$t + 60$" are calculated using the "joint" approach.

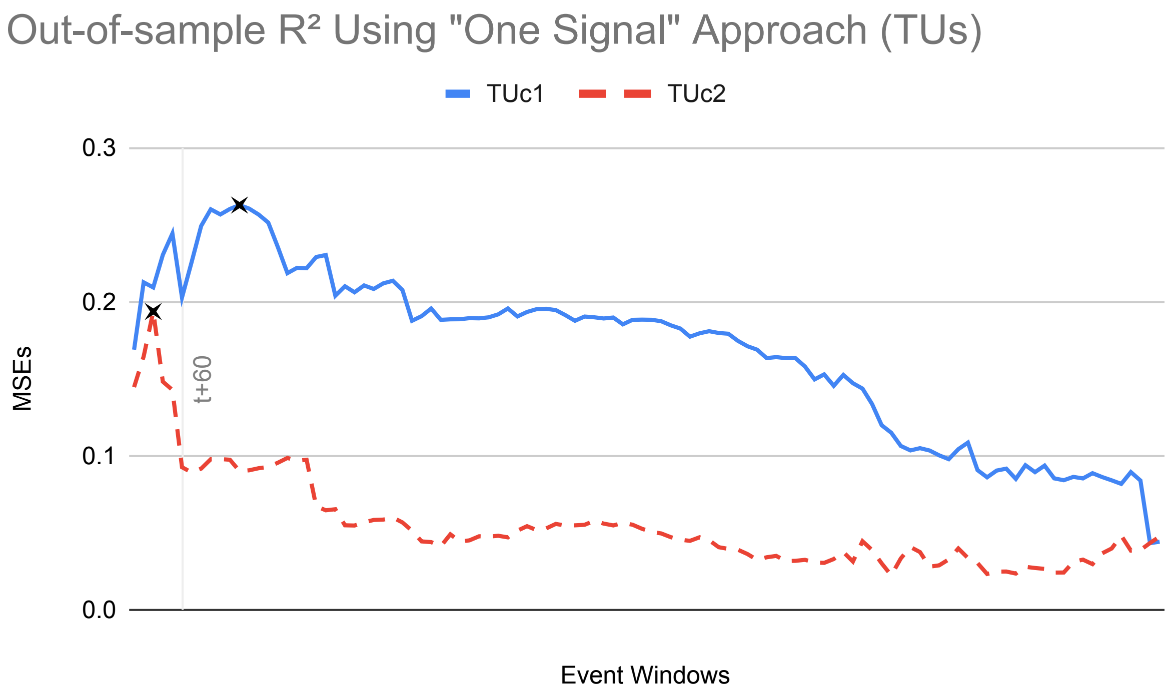

Figure E23: $\overline{R^2_{OOS}}$ Calculated Using the "One Signal" Approach for 2-year Treasury Futures

Notes: The horizontal axis depicts the event window lengths, starting from the systematically estimated event window length and ending at an event window starting 10 minutes before and ending 18 hours after FOMC statement release. The cross points represent the event window length associated with the largest $\overline{R^2_{OOS}}$. Estimates to the left of the vertical grey line labelled "$t + 60$" are calculated using the "joint" approach. The front-month 2-year Treasury futures contract is the one exception that sees the largest $\overline{R^2_{OOS}}$ within a window length than the systematically estimated length. As a robustness check, $\overline{R^2_{OOS}}$ was calculated with systematic estimation for the futures contract and considered event window length, which can be found in the rightmost box-and-whisker plot of Appendix Sub-figure E4a.

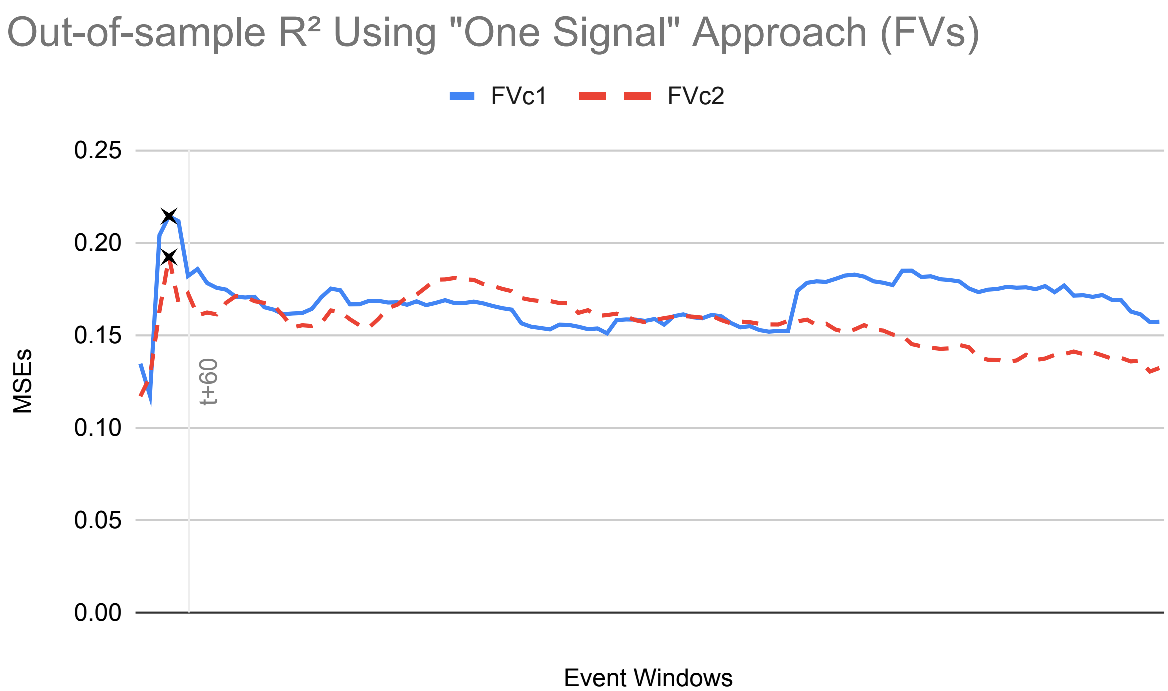

Figure E24: $\overline{R^2_{OOS}}$ Calculated Using the "One Signal" Approach for 5-year Treasury Futures

Notes: The horizontal axis depicts the event window lengths, starting from the systematically estimated event window length and ending at an event window starting 10 minutes before and ending 18 hours after FOMC statement release. The cross points represent the event window length associated with the largest $\overline{R^2_{OOS}}$. Estimates to the left of the vertical grey line labelled "$t + 60$" are calculated using the "joint" approach.

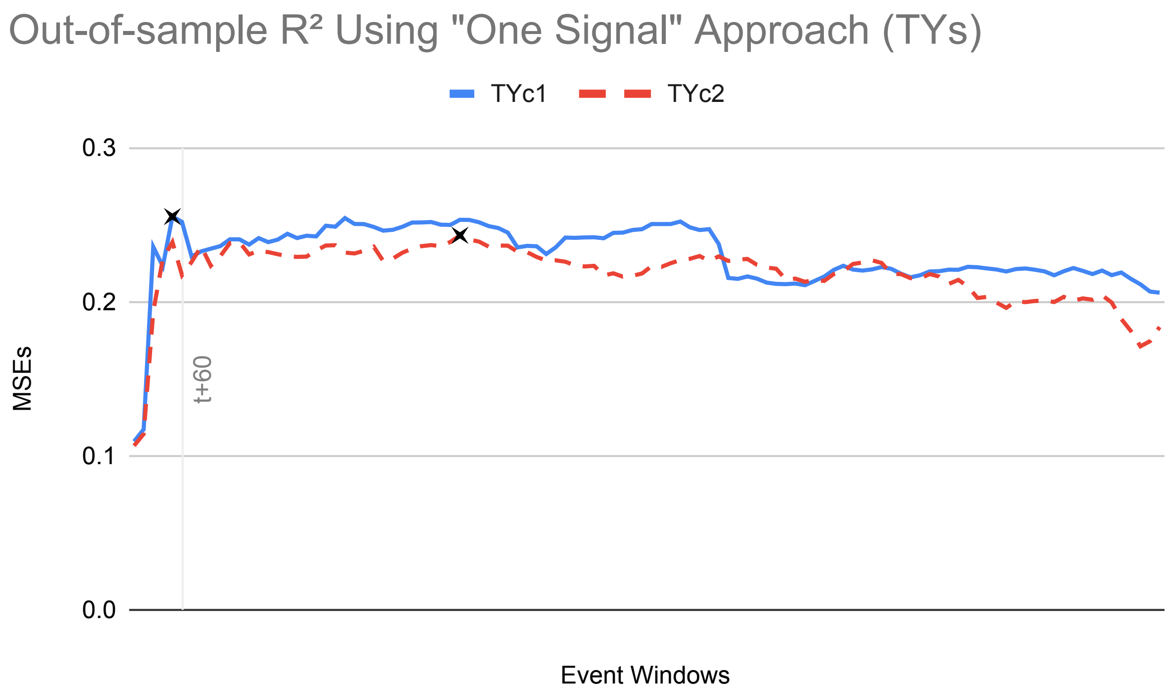

Figure E25: $\overline{R^2_{OOS}}$ Calculated Using the "One Signal" Approach for 10-year Treasury Futures

Notes: The horizontal axis depicts the event window lengths, starting from the systematically estimated event window length and ending at an event window starting 10 minutes before and ending 18 hours after FOMC statement release. The cross points represent the event window length associated with the largest $\overline{R^2_{OOS}}$. Estimates to the left of the vertical grey line labelled "$t + 60$" are calculated using the "joint" approach. The second-month 10-year Treasury futures contract is the one exception that sees the largest $\overline{R^2_{OOS}}$ within a window length than the systematically estimated length. As a robustness check, $\overline{R^2_{OOS}}$ was calculated with systematic estimation for the futures contract and considered event window length, which can be found in the rightmost box-and-whisker plot of Appendix Sub-figure E6b.

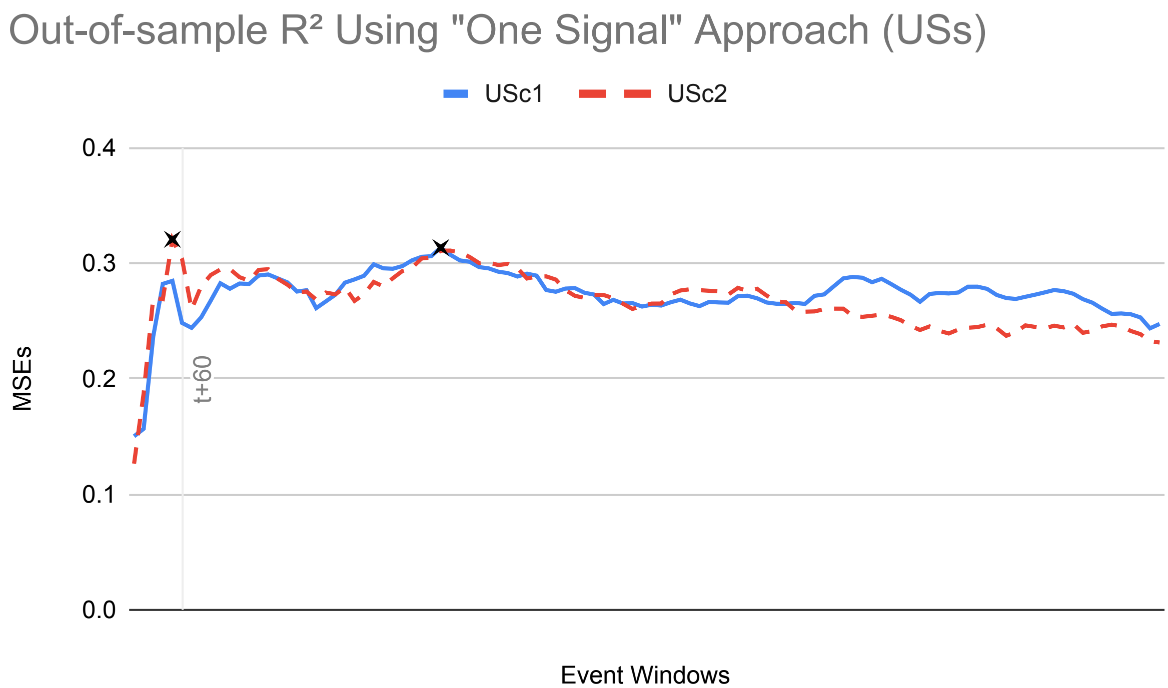

Figure E26: $\overline{R^2_{OOS}}$ Calculated Using the "One Signal" Approach for 30-year Treasury Futures

Notes: The horizontal axis depicts the event window lengths, starting from the systematically estimated event window length and ending at an event window starting 10 minutes before and ending 18 hours after FOMC statement release. The cross points represent the event window length associated with the largest $\overline{R^2_{OOS}}$. Estimates to the left of the vertical grey line labelled "$t + 60$" are calculated using the "joint" approach. The front-month 30-year Treasury futures contract is the one exception that sees the largest $\overline{R^2_{OOS}}$ within a window length than the systematically estimated length. As a robustness check, $\overline{R^2_{OOS}}$ was calculated with systematic estimation for the futures contract and considered event window length, which can be found in the rightmost box-and-whisker plot of Appendix Sub-figure E7a.

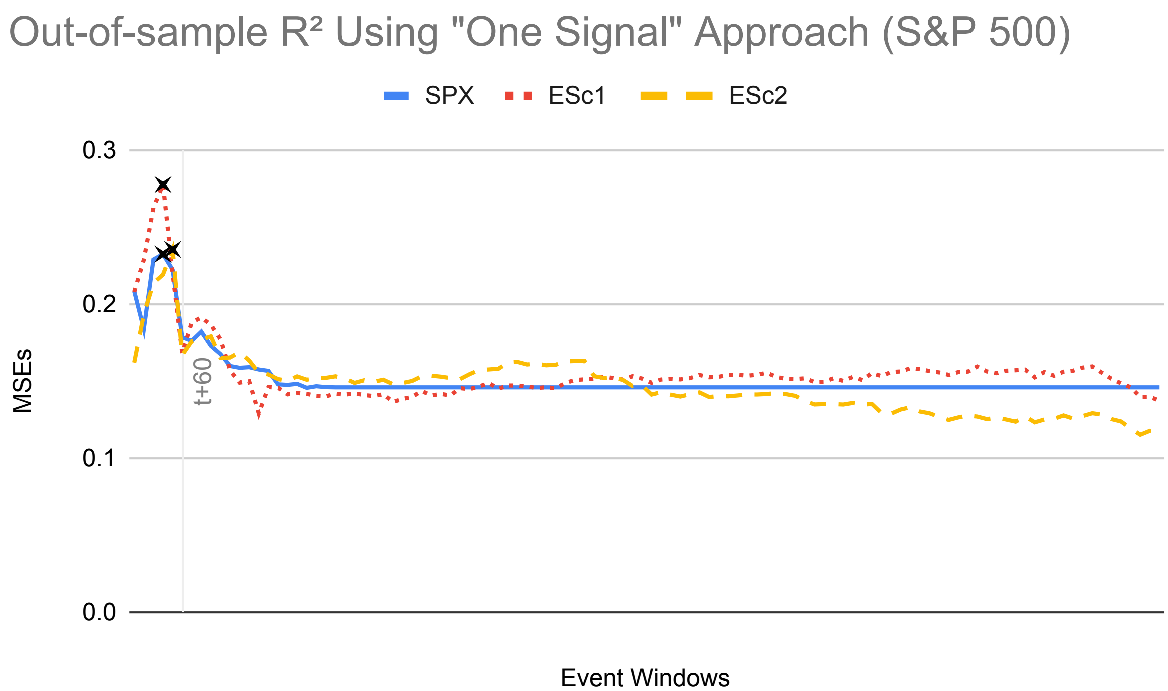

Figure E27: $\overline{R^2_{OOS}}$ Calculated Using the "One Signal" Approach for S&P 500 and E-mini Futures

Notes: The horizontal axis depicts the event window lengths, starting from the systematically estimated event window length and ending at an event window starting 10 minutes before and ending 18 hours after FOMC statement release. The cross points represent the event window length associated with the largest $\overline{R^2_{OOS}}$. Estimates to the left of the vertical grey line labelled "$t + 60$" are calculated using the "joint" approach.

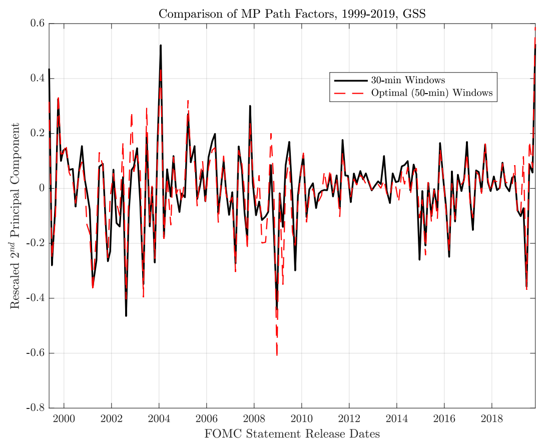

(a) GSS_P

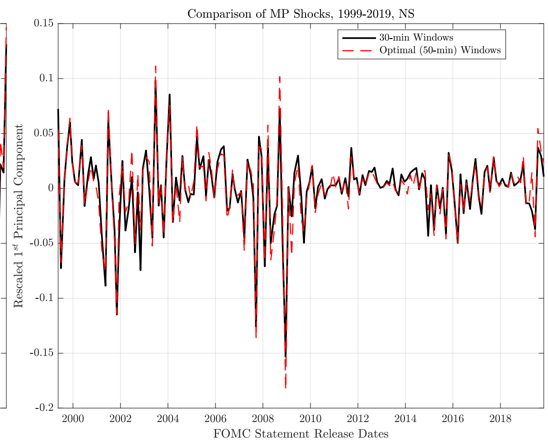

(b) NS_{MP}

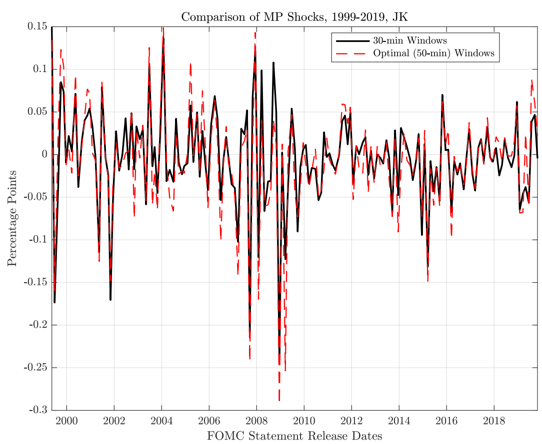

(c) JK_{MP}

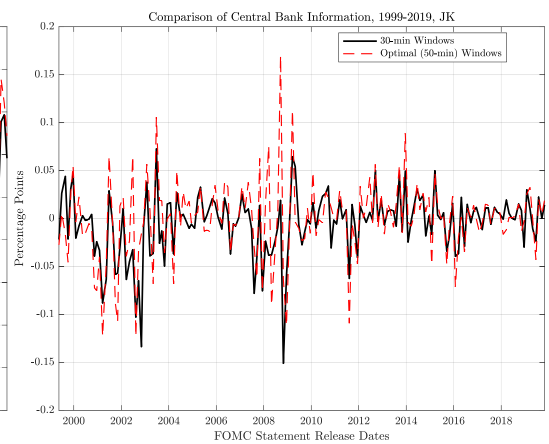

(d) JK_{CBI}

Figure E28: Comparing Monetary Policy Shock Series Derived from Optimal Window Length v. 30 Minutes

Notes: For each sub-figure, the horizontal axis represents FOMC statement release dates. The vertical axis depicts the principal components scaled according to the original specification of the authors. For all construction methods, the black-solid and red-dotted lines represent the shocks derived from surprises measured within 30 minutes and the median optimal event window length of 50 minutes, respectively.

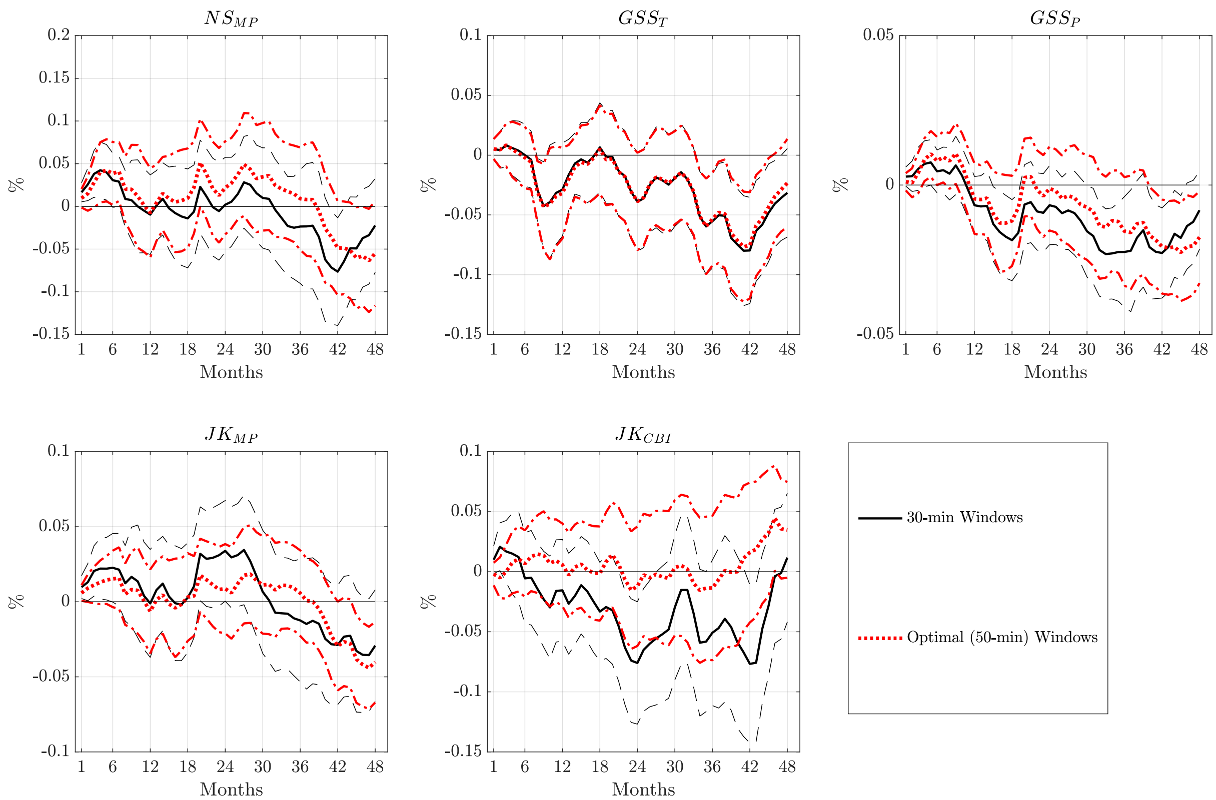

Figure E29: Effects of Event Window Choice on CPI Impulse Responses

Notes: Impulse responses are calculated using the lag-augmented local projection from Olea and Plagborg-Møller (2021). Confidence bands are at the 90% level. Standard errors are Eicker-Huber-White heteroscedasticity-robust standard errors. The responses are to a 100 basis point increase in the monetary policy shock series. The monetary policy shocks are constructed using the full set of monetary policy surprise instruments.

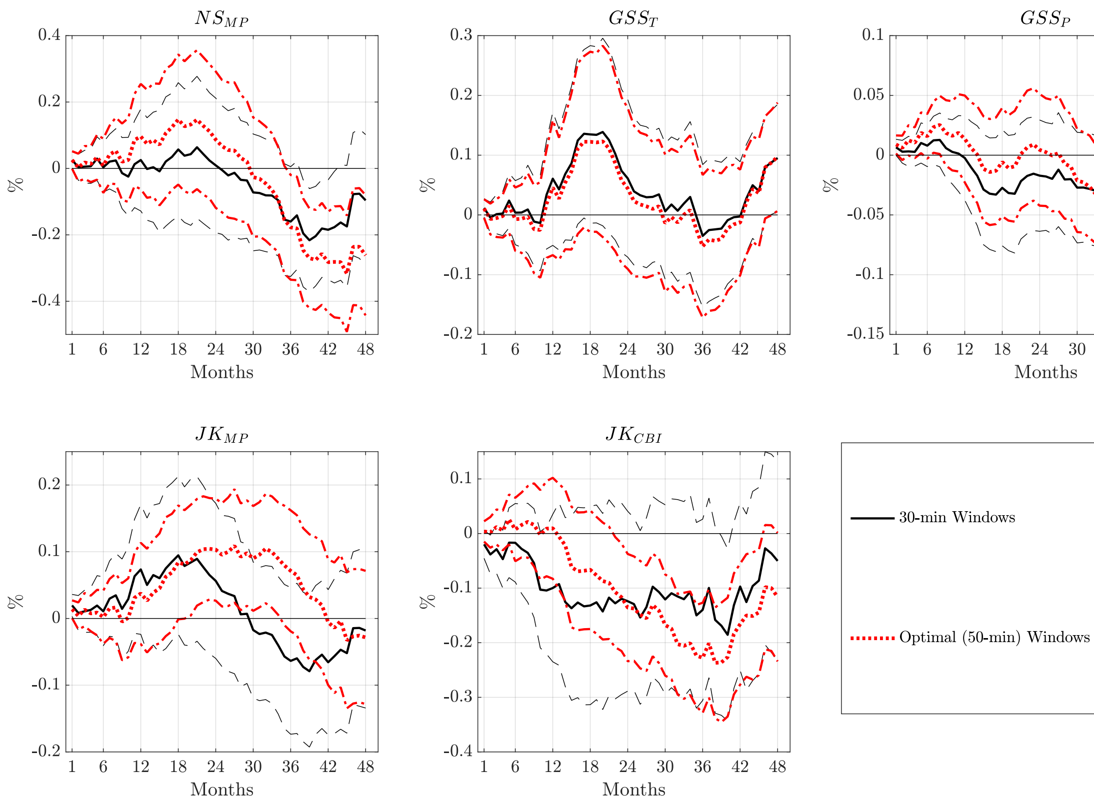

Figure E30: Effects of Event Window Choice on IP Impulse Responses

Notes: Impulse responses are calculated using the lag-augmented local projection from Olea and Plagborg-Møller (2021). Confidence bands are at the 90% level. Standard errors are Eicker-Huber-White heteroscedasticity-robust standard errors. The responses are to a 100 basis point increase in the monetary policy shock series. The monetary policy shocks are constructed using the full set of monetary policy surprise instruments.

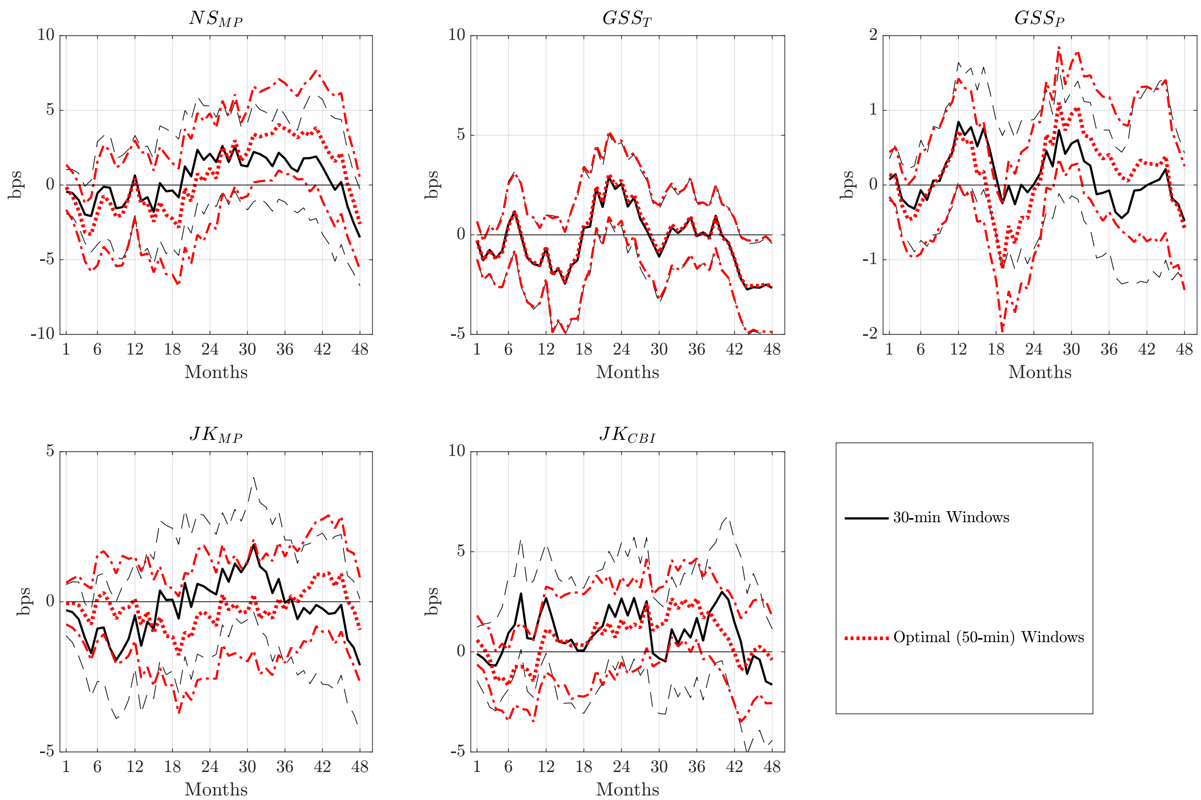

Figure E31: Effects of Event Window Choice on Excess Bond Premium Impulse Responses

Notes: Impulse responses are calculated using the lag-augmented local projection from Olea and Plagborg-Møller (2021). Confidence bands are at the 90% level. Standard errors are Eicker-Huber-White heteroscedasticity-robust standard errors. The responses are to a 100 basis point increase in the monetary policy shock series. The monetary policy shocks are constructed using the full set of monetary policy surprise instruments.

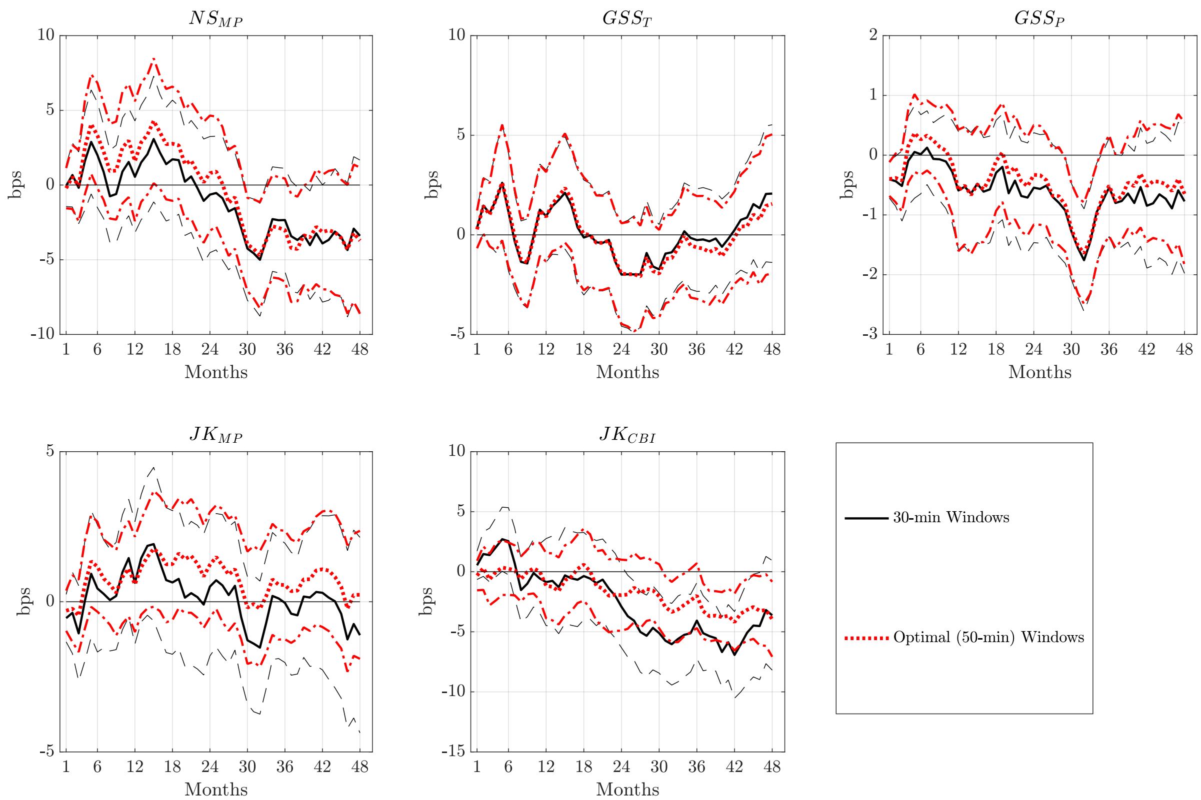

Figure E32: Effects of Event Window Choice on Two-year Treasury Yield Impulse Responses

Notes: Impulse responses are calculated using the lag-augmented local projection from Olea and Plagborg-Møller (2021). Confidence bands are at the 90% level. Standard errors are Eicker-Huber-White heteroscedasticity-robust standard errors. The responses are to a 100 basis point increase in the monetary policy shock series. The monetary policy shocks are constructed using the full set of monetary policy surprise instruments.

Appendix Tables

	Scenario 1	Scenario 2	Scenario 3
Framework Simulation Parameters			
P_i^f	$\in [-100, 100]$	$\in [-100, 100]$	$\in [-100, 100]$
$\varepsilon_{i,0}^c$	$\in [-100, 100]$	$\in [-100, 100]$	$\in [-100, 100]$
$\varepsilon_{i,0}^n$	0	0	0
σ_c	100	0.1	50
$\mathcal{D}$	0.5	1	0.75
σ_n	0.1	10	1
ρ_c	0.47	0.47	0.47
σ_s	$\in \mathbb{R}$	$\in \mathbb{R}$	$\in \mathbb{R}$
Simulation Results			
t^*	16	2	10
$\hat{t}$	15	2	10

Table F1: Framework Simulation Parameters and Results for Different Asset Market Scenarios
Notes: The results are from 10,000 simulations of the asset price framework for the three asset market scenarios considered. t^* and $\hat{t}$ are defined as the time at which the asset market fully reacts to news on average, where the former is calculated using the fundamental price component and the latter with the observed signal according to the motivating framework.

Name	Maturity	Ticker	Sample	Observations
Inputs				
FOMC Statements	N/A	N/A	1999–2019	165
Outputs				
Federal Funds Rate Futures	Front-month	$FFc1$	1999–2019	165
Federal Funds Rate Futures	1-month-ahead	$FFc2$	1999–2019	165
Federal Funds Rate Futures	2-month-ahead	$FFc3$	1999–2019	165
Federal Funds Rate Futures	3-month-ahead	$FFc4$	1999–2019	165
Eurodollar Futures	2-quarter	$EDcm2$	1999–2019	165
Eurodollar Futures	3-quarter	$EDcm3$	1999–2019	165
Eurodollar Futures	4-quarter	$EDcm4$	1999–2019	165
2-year Treasury Futures	Front-month	$TUc1$	1999–2019	165
2-year Treasury Futures	Second-month	$TUc2$	1999–2019	165
5-year Treasury Futures	Front-month	$FVc1$	1999–2019	165
5-year Treasury Futures	Second-month	$FVc2$	1999–2019	165
10-year Treasury Futures	Front-month	$TYc1$	1999–2019	165
10-year Treasury Futures	Second-month	$TYc2$	1999–2019	165
30-year Treasury Futures	Front-month	$USc1$	1999–2019	165
30-year Treasury Futures	Second-month	$USc2$	1999–2019	165
S&P 500 Index	N/A	SPX	1999–2019	165
S&P 500 E-mini Futures	Front-month	$ESc1$	1999–2019	165
S&P 500 E-mini Futures	Second-month	$ESc2$	1999–2019	165

Table F2: Independent and Dependent Variables for Systematic Estimation of Appropriate Event Window Lengths by Neural Networks
Notes: The table shows the FOMC statements and financial market asset prices considered as independent and dependent variables, respectively, in my analysis. The statements are collected from the Board of Governors of the Federal Reserve System website and the asset prices are from the Thomson Reuters Tick History database. All series begin in May 1999 and end in October 2019. *Ticker* refers to the Reuters Instrument Code, which uniquely identifies the financial instrument. The letters *c* and *cm* standard for continuous futures contracts.

Hyperparameter	Value
Number of Layers	12
Hidden Size	768
Number of Attention Heads	12
Attention Head Size	64
Hidden Dropout	0.1
Attention Dropout	0.1
Max Sequence Length	512
Vocab Size	32,000
Training Batch Size	8
Evaluation Batch Size	8
Max Num of Steps	2040
Max Num of Epochs	120
Manual Random Seed	47
Learning Rate	*
Warmup Ratio	0.06
Adam Epsilon	1e-8
Learning Rate Decay	Linear

Table F3: XLNet-Base Hyperparameters for Fine-tuning

Notes: * denotes the hyperparameter that undergoes tuning during the fine-tuning process. The random seed is set as a constant to ensure replicability by always initialising random components (e.g., weights and biases) to the same initial values.

Asset	$\overline{R^2_{OOS}}$, 30-min	$\overline{R^2_{OOS}}$, Optimal	**Difference**
$FF1$	35.0%	37.2%	$+$**2.2 p.p.**
$FF2$	28.7%	34.5%	$+$**5.8 p.p.**
$FF3$	23.0%	40.2%	$+$**17.2 p.p.**
$FF4$	29.8%	43.3%	$+$**13.5 p.p.**
$EDcm2$	18.3%	23.3%	$+$**5 p.p.**
$EDcm3$	14.0%	18.2%	$+$**4.2 p.p.**
$EDcm4$	11.2%	16.0%	$+$**4.8 p.p.**
$TUc1$	21.3%	24.4%	$+$**3.1 p.p.**
$TUc2$	16.5%	19.4%	$+$**2.9 p.p.**
$FVc1$	11.7%	21.4%	$+$**9.7 p.p.**
$FVc2$	12.7%	19.2%	$+$**6.5 p.p.**
$TYc1$	11.7%	25.5%	$+$**13.8 p.p.**
$TYc2$	11.4%	23.9%	$+$**12.5 p.p.**
$USc1$	15.7%	28.5%	$+$**12.8 p.p.**
$USc2$	18.7%	32.1%	$+$**13.4 p.p.**
SPX	18.4%	23.2%	$+$**4.8 p.p.**
$ESc1$	22.9%	27.7%	$+$**4.8 p.p.**
$ESc2$	19.3%	23.5%	$+$**4.2 p.p.**

Table F4: Differences of $\overline{R^2_{OOS}}$ between 30-minute and Optimal Event Windows

Series Name	Notation	Description
Gürkaynak, Sack, and Swanson (2005) Target Shocks	GSS_T	First principal component of monetary policy surprises that is rotated such that it drives unexpected changes in the surprise in the current federal funds rate. The authors re-scale the shock originally to be one-for-one with the surprise in the current federal funds rate.
Gürkaynak, Sack, and Swanson (2005) Path Shocks	GSS_P	Second principal component of monetary policy surprises that is rotated such that on average, it has no effect on the surprise in the current federal funds rate. The authors re-scale the shock originally to have equal effect with GSS_T on the surprise in the four-quarter Eurodollar futures.
Nakamura and Steinsson (2018) Shocks	NS_{MP}	First principal component of monetary policy surprises. The authors re-scale the shock originally to be one-for-one with the daily change in the zero-coupon, nominal one-year Treasury yield.
Jarociński and Karadi (2020) Shocks	JK_{MP}	First principal component of monetary policy surprises that have negative comovement with stock market changes. The authors make the first principal component to have equal standard deviation with that of the surprise in the four-quarter Eurodollar futures.
Jarociński and Karadi (2020) Central Bank Information	JK_{CBI}	First principal component of monetary policy surprises that have positive comovement with stock market changes. The authors make the first principal component to have equal standard deviation with that of the surprise in the four-quarter Eurodollar futures.

Table F5: Monetary Policy Shock Series

Notes: All monetary policy shocks are re-scaled to be one-for-one with the daily change in the zero-coupon, nominal one-year Treasury yield for easy interpretation and comparison. Before applying sign restrictions to obtain JK_{MP} and JK_{CBI}, I follow Jarociński and Karadi (2020) by first re-scaling the first principal component of monetary policy surprises in terms of the standard deviation of the zero-coupon, nominal one-year Treasury yield.

Metric	GSS_T	GSS_P	NS_{MP}	JK_{MP}	JK_{CBI}
Count	165	165	165	165	165
	(165)	(165)	(165)	(165)	(165)
Mean	0	0	0	-0.0023	-0.0028
	(0)	(0)	(0)	(-0.0035)	(-0.0004)
SD	0.0219	0.0282	0.0356	0.0305	0.0180
	(0.0216)	(0.0309)	(0.0375)	(0.0297)	(0.0233)
Max	0.0756	0.1223	0.1333	0.0968	0.0388
	(0.0652)	(0.1079)	(0.1135)	(0.0840)	(0.0984)
75^{th}	0.0076	0.0134	0.0193	0.0175	0.0069
	(0.0084)	(0.0167)	(0.0203)	(0.0124)	(0.0112)
Median	0.0011	-0.0010	0.0033	-0.0017	-0.0003
	(0.0023)	(0.0019)	(0.0003)	(-0.0009)	(0.0022)
25^{th}	-0.0069	-0.0124	-0.0187	-0.0153	-0.0086
	(-0.0077)	(-0.0127)	(-0.0164)	(-0.0138)	(-0.0103)
Min	-0.1038	-0.1611	-0.1403	-0.1222	-0.0812
	(-0.1025)	(-0.1618)	(-0.1722)	(-0.1491)	(-0.0828)

Table F6: Descriptive Statistics for Monetary Policy Shock Series, FOMC Statement Frequency for 1999–2019
Notes: Numbers in parentheses are summary statistics for each shock series derived from monetary policy surprises calculated within the median optimal event window length of 50 minutes. All shocks series have been re-scaled to be one-for-one with the daily change in the zero-coupon, nominal one-year Treasury yield.

Metric	$ln(CPI)$	$ln(IP)$	EBP	TY_2
Count	246	246	246	246
Mean	5.356	4.597	0.110	2.210
SD	0.126	0.057	0.715	1.794
Max	5.551	4.706	3.283	6.650
75^{th}	5.466	4.642	0.341	3.569
Median	5.381	4.605	-0.108	1.673
25^{th}	5.241	4.548	-0.334	0.662
Min	5.112	4.467	-1.140	0.188

Table F7: Descriptive Statistics for Impulse Response Variables, Monthly for 1999–2019
Notes: All logs are natural logarithms. Consumer Price Index (CPI) and Industrial Production (IP) are sourced from the Federal Reserve Economic Data. The Excess Bond Premium (EBP) is from Gilchrist and Zakrajšek (2012) and is expressed in percentage points. The two-year Treasury yield (TY_2) is from Gürkaynak, Sack, and Swanson (2005).

Metric	GSS_T	GSS_P	NS_{MP}	JK_{MP}	JK_{CBI}
Count	165	165	165	165	165
	(165)	(165)	(165)	(165)	(165)
Mean	0	0	0	-0.0015	-0.0019
	(0)	(0)	(0)	(-0.0024)	(-0.0002)
SD	0.0180	0.0230	0.0291	0.0250	0.0148
	(0.0177)	(0.0253)	(0.0307)	(0.0243)	(0.0190)
Max	0.0756	0.1223	0.1333	0.0968	0.0388
	(0.0652)	(0.1079)	(0.1135)	(0.0840)	(0.0984)
75^{th}	0.0042	0.0078	0.0114	0.0055	0.0031
	(0.0042)	(0.0089)	(0.0110)	(0.0058)	(0.0075)
Median	0	0	0	0	0
	(0)	(0)	(0)	(0)	(0)
25^{th}	-0.0021	-0.0064	-0.0036	-0.0073	-0.0037
	(-0.0011)	(-0.0052)	(-0.0039)	(-0.0070)	(-0.0035)
Min	-0.1038	-0.1611	-0.1403	-0.1222	-0.0812
	(-0.1025)	(-0.1618)	(-0.1722)	(-0.1491)	(-0.0828)

Table F8: Descriptive Statistics for Monetary Policy Shock Series, Monthly for 1999–2019
Notes: Numbers in parentheses are summary statistics for each shock series derived from monetary policy surprises calculated within the median optimal event window length of 50 minutes. All shocks series have been re-scaled to be one-for-one with the daily change in the zero-coupon, nominal one-year Treasury yield. All shocks are zero for any month that does not have an FOMC meeting.

Metric	FKGL	S^1
Count	165	164
Mean	16.361	0.751
SD	1.715	0.212
Max	21.3	0.984
75^{th}	17.3	0.920
Median	16.5	0.826
25^{th}	15.1	0.622
Min	12.2	0.200

Table F9: Descriptive Statistics for Heterogeneity Analyses
Notes: The complexity of FOMC statements is measured by the Flesch-Kincaid Grade Level (FKGL), defined as: $0.39 \times \text{average sentence length} + 11.8 \times \text{average number of syllables per word} - 15.59$. S^1 is the cosine similarity measure between sequential FOMC statements.

domestic	alreadytight	recov	tilt	buildup
alert	alter	foreign	imbal	undermin
direct	excess	quit	favor	perform
strength	fall	trend	eas	concern
background	firm	gener	demand	potenti
core	subdu	cost	longer	gain

Table F10: FOMC Statement Base Terms with Top 30 TFIDF Scores

Notes: TFIDF is a weighted frequency combining word and document counts such that greater weight is given to words that are more informative about the information content of the FOMC statements relative to other statements where the words are not found. Specifically, the weighted term frequency gives higher weight to terms that occur more frequently in a given document. The term frequency is then divided by the number of documents that has this term appear. The more documents that have the word, the less importance and weight will be given to the word as it is less informative for distinguishing documents from one another.

	30-minute Window				Optimal Window				**Difference**		
	ΔTY_1	ΔTY_2	ΔTY_5	ΔTY_{10}	ΔTY_1	ΔTY_2	ΔTY_5	ΔTY_{10}	ΔTY_2	ΔTY_5	ΔTY_{10}
GSS_T	0.36***	0.28**	0.15	0.05	0.35***	0.27**	0.13	0.03	**-0.01**	**-0.01**	**-0.02**
	(0.11)	(0.13)	(0.12)	(0.11)	(0.11)	(0.13)	(0.11)	(0.10)			
GSS_P	0.21***	0.25***	0.24***	0.15***	0.21***	0.28***	0.27***	0.19***	**+0.01**	**+0.02**	**+0.03**
	(0.02)	(0.03)	(0.04)	(0.04)	(0.03)	(0.04)	(0.05)	(0.06)			
NS_{MP}	1.00***	1.11***	0.95***	0.57***	1.00***	1.16***	1.04***	0.67***	**+0.05**	**+0.09**	**+0.10**
	(0.09)	(0.14)	(0.18)	(0.19)	(0.11)	(0.15)	(0.20)	(0.22)			
JK_{MP}	0.53***	0.64***	0.61***	0.39***	0.44***	0.55***	0.56***	0.39***	**-0.09**	**-0.09**	**-0.06**
	(0.08)	(0.11)	(0.16)	(0.18)	(0.08)	(0.11)	(0.18)	(0.21)			
JK_{CBI}	0.51***	0.39**	0.15	0.03	0.52***	0.50***	0.29**	0.12	**+0.01**	**+0.11**	**+0.15**
	(0.16)	(0.19)	(0.23)	(0.24)	(0.15)	(0.19)	(0.25)	(0.24)			

Table F11: Differences in Responses of Nominal Interest Rates to Shocks from Event Window Choice

Notes: Each estimate comes from a separate OLS regression. For each regression, the dependent variable is the one-day change in end-of-day Treasury yields, represented by columns. The independent variable is the change in monetary policy shock series over a given event window around the time of scheduled FOMC announcements. Starting from the left, columns 2–5 represent the OLS regressions performed on shock series constructed within 30-minute windows. Columns 6–9 are equivalent, but for shocks constructed within the median optimal window length of 50 minutes. The sample period consists of all scheduled FOMC statement release dates from May 1999 through October 2019, resulting in 165 observations. Newey and West (1987) standard errors are reported in parentheses. Columns 10–12 represent the differences between the coefficients of the regressions using the optimal windows v. those using 30-minute windows. ** sig. at the 5% level. *** sig. at the 1% level.

	30-minute Window			Optimal Window			**Difference**		
	$\Delta TIPS_2$	$\Delta TIPS_5$	$\Delta TIPS_{10}$	$\Delta TIPS_2$	$\Delta TIPS_5$	$\Delta TIPS_{10}$	$\Delta TIPS_2$	$\Delta TIPS_5$	$\Delta TIPS_{10}$
GSS_T	-0.43	0.21	0.15	-0.40	0.21	0.14	**+0.03**	**-0.00**	**-0.01**
	(0.75)	(0.21)	(0.15)	(0.71)	(0.20)	(0.14)			
GSS_P	0.25**	0.21***	0.16***	0.28***	0.26***	0.21***	**+0.02**	**+0.06**	**+0.04**
	(0.12)	(0.06)	(0.05)	(0.13)	(0.08)	(0.07)			
NS_{MP}	0.56***	0.89***	0.69***	0.68***	1.07***	0.83***	**+0.12**	**+0.18**	**+0.13**
	(0.93)	(0.30)	(0.23)	(0.93)	(0.32)	(0.27)			
JK_{MP}	0.53**	0.60***	0.48***	0.54**	0.60***	0.49***	**+0.01**	**-0.00**	**+0.01**
	(0.52)	(0.24)	(0.20)	(0.38)	(0.25)	(0.22)			
JK_{CBI}	-0.39	0.07	0.03	-0.23	0.25	0.11	**+0.16**	**+0.17**	**+0.08**
	(0.56)	(0.25)	(0.23)	(0.71)	(0.29)	(0.25)			

Table F12: Differences in Responses of Real Interest Rates to Shocks from Event Window Choice

Notes: Each estimate comes from a separate OLS regression. For each regression, the dependent variable is the one-day change in end-of-day TIPS yields, represented by columns. The independent variable is the change in monetary policy shock series over a given event window around the time of scheduled FOMC announcements. Starting from the left, columns 2–4 represent the OLS regressions performed on shock series constructed within 30-minute windows. Columns 5–7 are equivalent, but for shocks constructed within the median optimal window length of 50 minutes. The sample period consists of all scheduled FOMC statement release dates from May 1999 to October 2019, resulting in 165 observations. Newey and West (1987) standard errors are reported in parentheses. Columns 8–10 represent the differences between the coefficients of the regressions using the optimal windows v. those using 30-minute windows. ** sig. at the 5% level. *** sig. at the 1% level.

	30-minute Window			Optimal Window			**Difference**		
	ΔBEI_2	ΔBEI_5	ΔBEI_{10}	ΔBEI_2	ΔBEI_5	ΔBEI_{10}	ΔBEI_2	ΔBEI_5	ΔBEI_{10}
GSS_T	0.71**	-0.06	−0.10*	0.67**	-0.07	−0.11**	**-0.04**	**-0.02**	**-0.01**
	(0.75)	(0.13)	(0.09)	(0.71)	(0.13)	(0.09)			
GSS_P	0.00	0.03	-0.01	-0.00	0.01	-0.02	**-0.00**	**-0.02**	**-0.01**
	(0.11)	(0.05)	(0.03)	(0.12)	(0.05)	(0.03)			
NS_{MP}	0.55	0.06	−0.13*	0.48	-0.03	−0.15**	**-0.07**	**-0.09**	**-0.03**
	(0.90)	(0.24)	(0.13)	(0.92)	(0.24)	(0.12)			
JK_{MP}	0.11	0.02	−0.09**	0.01	-0.04	−0.10***	**-0.10**	**-0.05**	**-0.01**
	(0.50)	(0.16)	(0.07)	(0.36)	(0.12)	(0.13)			
JK_{CBI}	0.78**	0.07	-0.00	0.73**	0.04	0.01	**-0.05**	**-0.03**	+**0.01**
	(0.55)	(0.17)	(0.11)	(0.67)	(0.18)	(0.13)			

Table F13: Differences in Responses of Break-even Inflation to Shocks from Event Window Choice

Notes: Each estimate comes from a separate OLS regression. For each regression, the dependent variable is the one-day change in end-of-day break-even inflation, represented by columns. The independent variable is the change in monetary policy shock series over a given event window around the time of scheduled FOMC announcements. Starting from the left, columns 2–4 represent the OLS regressions performed on shock series constructed within 30-minute windows. Columns 5–7 are equivalent, but for shocks constructed within the median optimal window length of 50 minutes. The sample period consists of all scheduled FOMC statement release dates from May 1999 to October 2019, resulting in 165 observations. Newey and West (1987) standard errors are reported in parentheses. Columns 8–10 represent the differences between the coefficients of the regressions using the optimal windows v. those using 30-minute windows. * sig. at the 10% level. ** sig. at the 5% level. *** sig. at the 1% level.

	30-minute Window			Optimal Window			**Percentage Difference**		
	DP_{SPX}	DP_{ESc1}	DP_{ESc2}	DP_{SPX}	DP_{ESc1}	DP_{ESc2}	DP_{SPX}	DP_{ESc1}	DP_{ESc2}
GSS_T	-3.46^{***}	-3.54^{***}	-2.68^{***}	-3.24^{***}	-3.25^{***}	-2.67^{**}	**-6.52%**	**-8.15%**	**-0.57%**
	(1.36)	(1.39)	(1.40)	(1.23)	(1.25)	(1.20)			
GSS_P	-1.39^{***}	-1.47^{***}	-1.39^{***}	-1.43^{***}	-1.45^{***}	-1.69^{***}	**+2.52%**	**-1.12%**	**+21.24%**
	(0.39)	(0.40)	(0.36)	(0.51)	(0.53)	(0.48)			
NS_{MP}	-7.57^{***}	-7.88^{***}	-6.96^{***}	-7.33^{***}	-7.41^{***}	-7.82^{***}	**-3.32%**	**-6.00%**	**+12.30%**
	(1.52)	(1.58)	(1.57)	(1.81)	(1.87)	(1.82)			
JK_{MP}	-8.16^{***}	-8.35^{***}	-7.53^{***}	-8.01^{***}	-8.15^{***}	-7.91^{***}	**-1.75%**	**-2.35%**	**+5.07%**
	(0.43)	(0.47)	(0.57)	(0.58)	(0.60)	(0.62)			
JK_{CBI}	8.27^{***}	8.20^{***}	7.68^{***}	8.04^{***}	8.25^{***}	7.00^{***}	**-2.84%**	**+0.66%**	**-8.97%**
	(1.25)	(1.29)	(1.24)	(1.44)	(1.46)	(1.47)			

Table F14: Differences in Responses of Stock Prices to Shocks from Event Window Choice

Notes: Each estimate comes from a separate OLS regression. For each regression, the dependent variable is the price log-difference of the S&P 500 Index or E-mini futures, represented by columns. The independent variable is the change in monetary policy shock series over a given event window around the time of scheduled FOMC announcements. Starting from the left, columns 2–4 represent the OLS regressions performed on shock series constructed within 30-minute windows. Columns 5–7 are equivalent, but for shocks constructed within the median optimal window length of 50 minutes. The sample period consists of all scheduled FOMC statement release dates from May 1999 to October 2019, resulting in 165 observations. Newey and West (1987) standard errors are reported in parentheses. Columns 8–10 represent the percentage differences between the coefficients of the regressions using the optimal windows v. those using 30-minute windows, where positive (negative) values represent a stronger (weaker) effect in the same direction. *** sig. at the 1% level.